

CS 350S: Privacy-Preserving Systems

Oblivious RAM

Outline

1. Overview
2. Square-root construction
3. Hierarchical construction
4. Limitations
5. Student presentation: PathORAM

Access patterns can reveal sensitive information



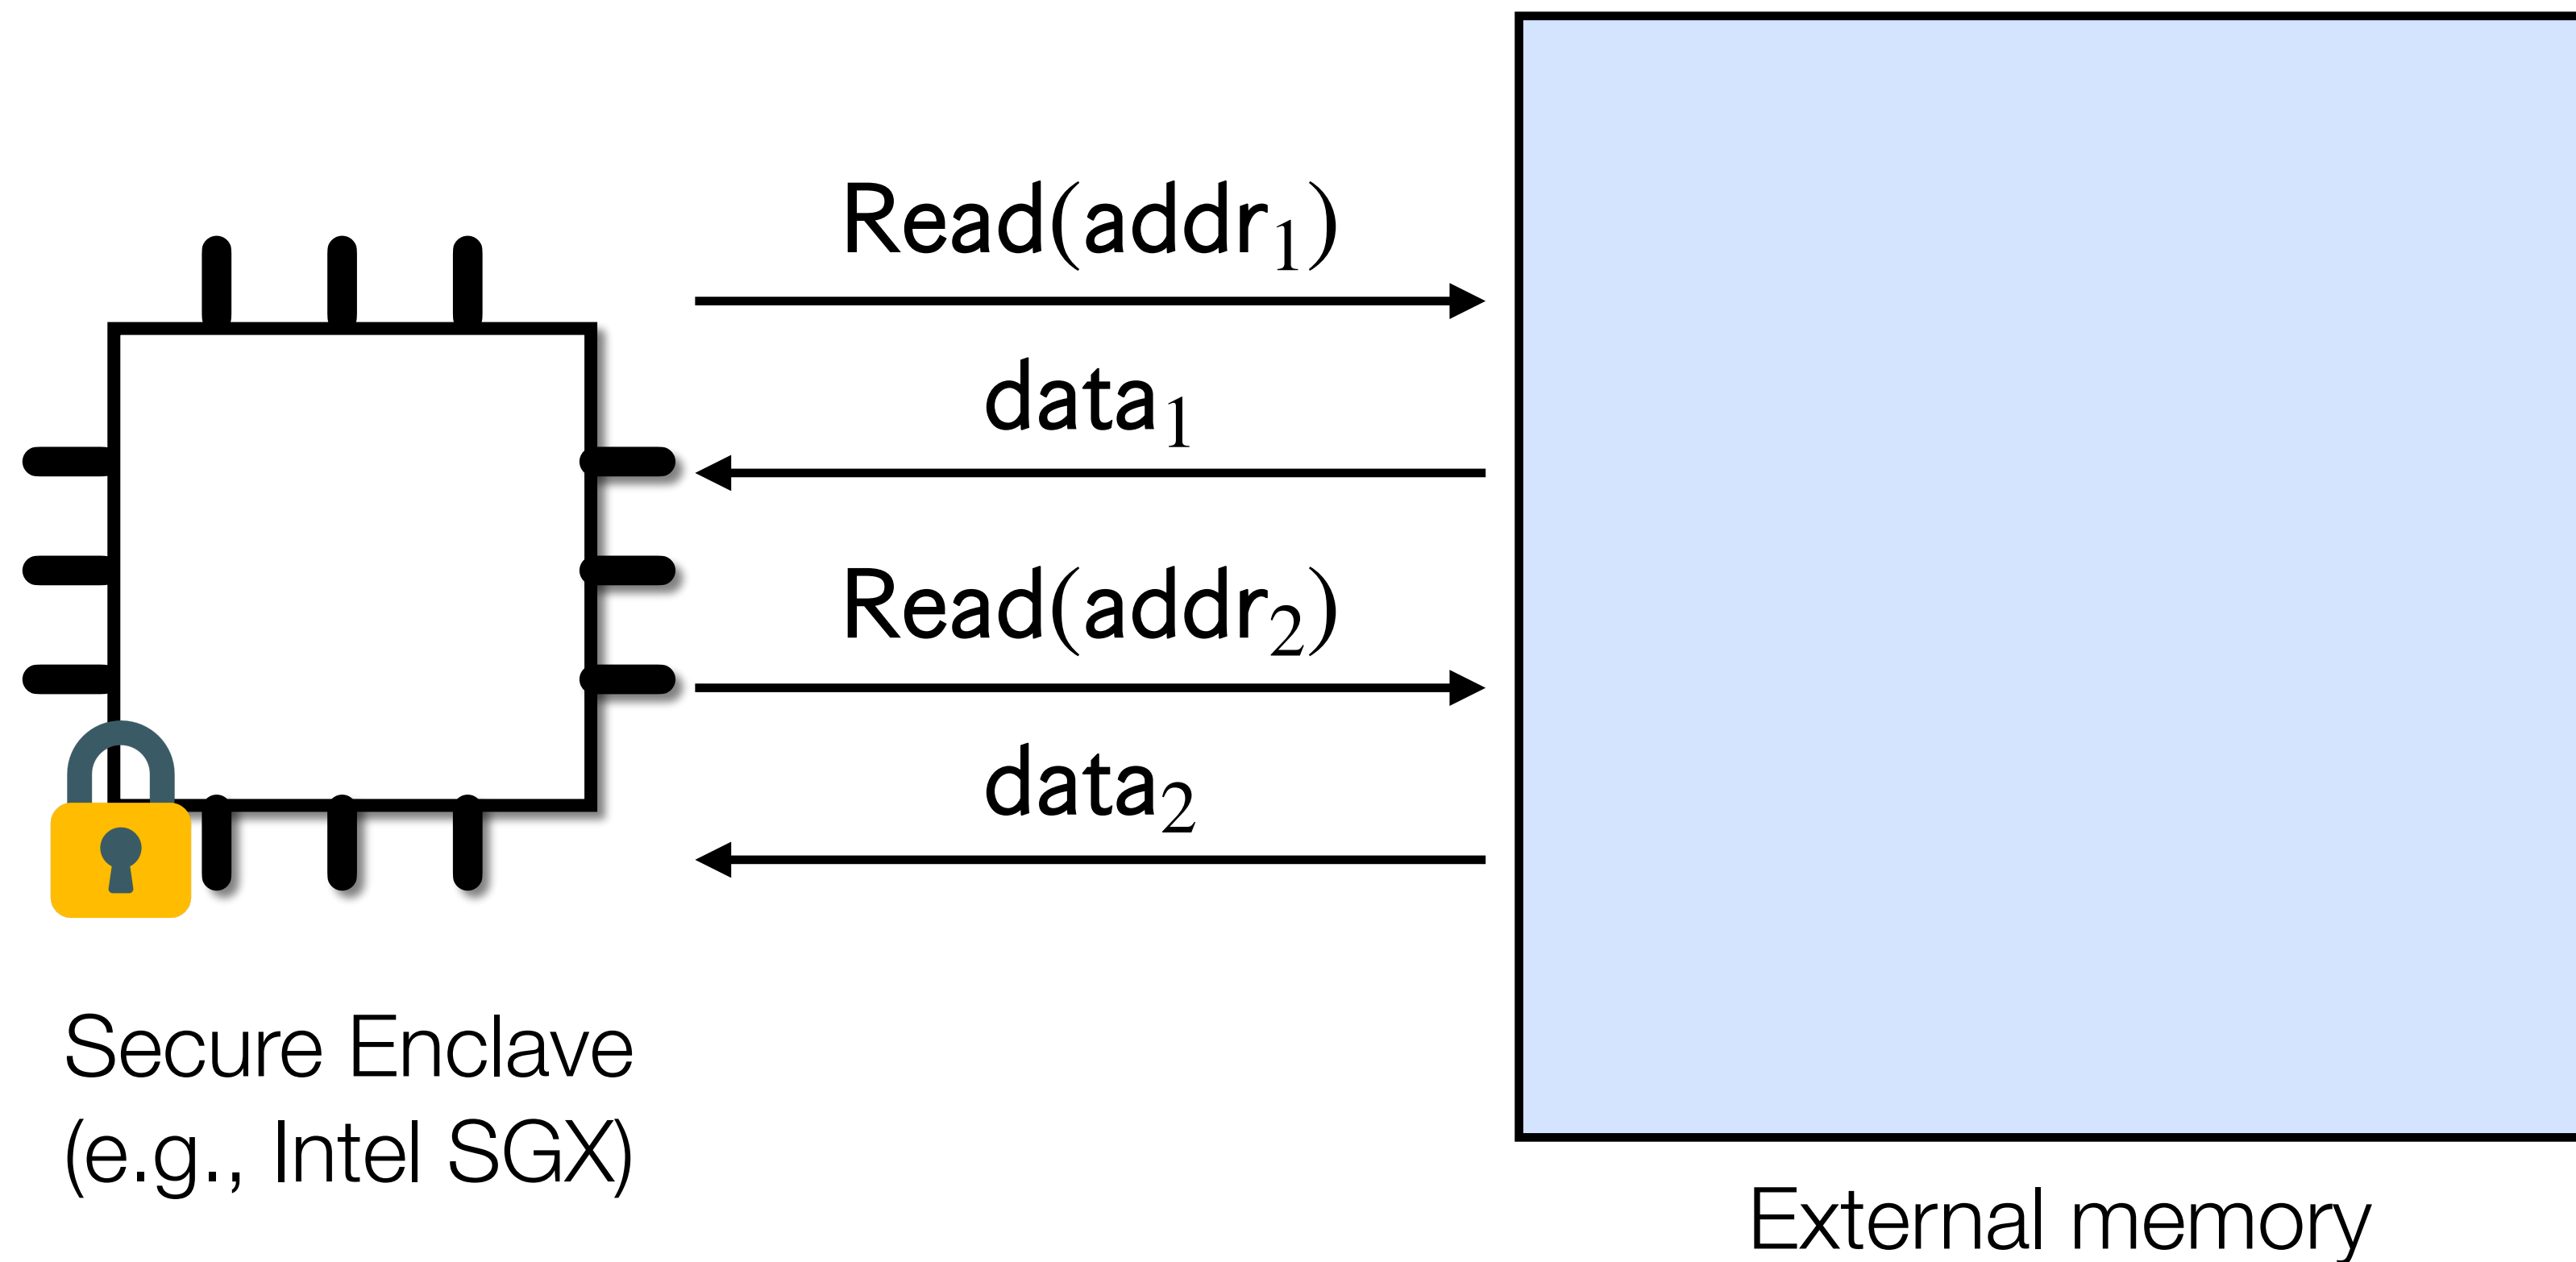
$\text{Enc}_k(\text{heart disease})$	$\text{Enc}_k(\text{info}_1)$
$\text{Enc}_k(\text{cancer})$	$\text{Enc}_k(\text{info}_2)$
$\text{Enc}_k(\text{diabetes})$	$\text{Enc}_k(\text{info}_3)$
...	...



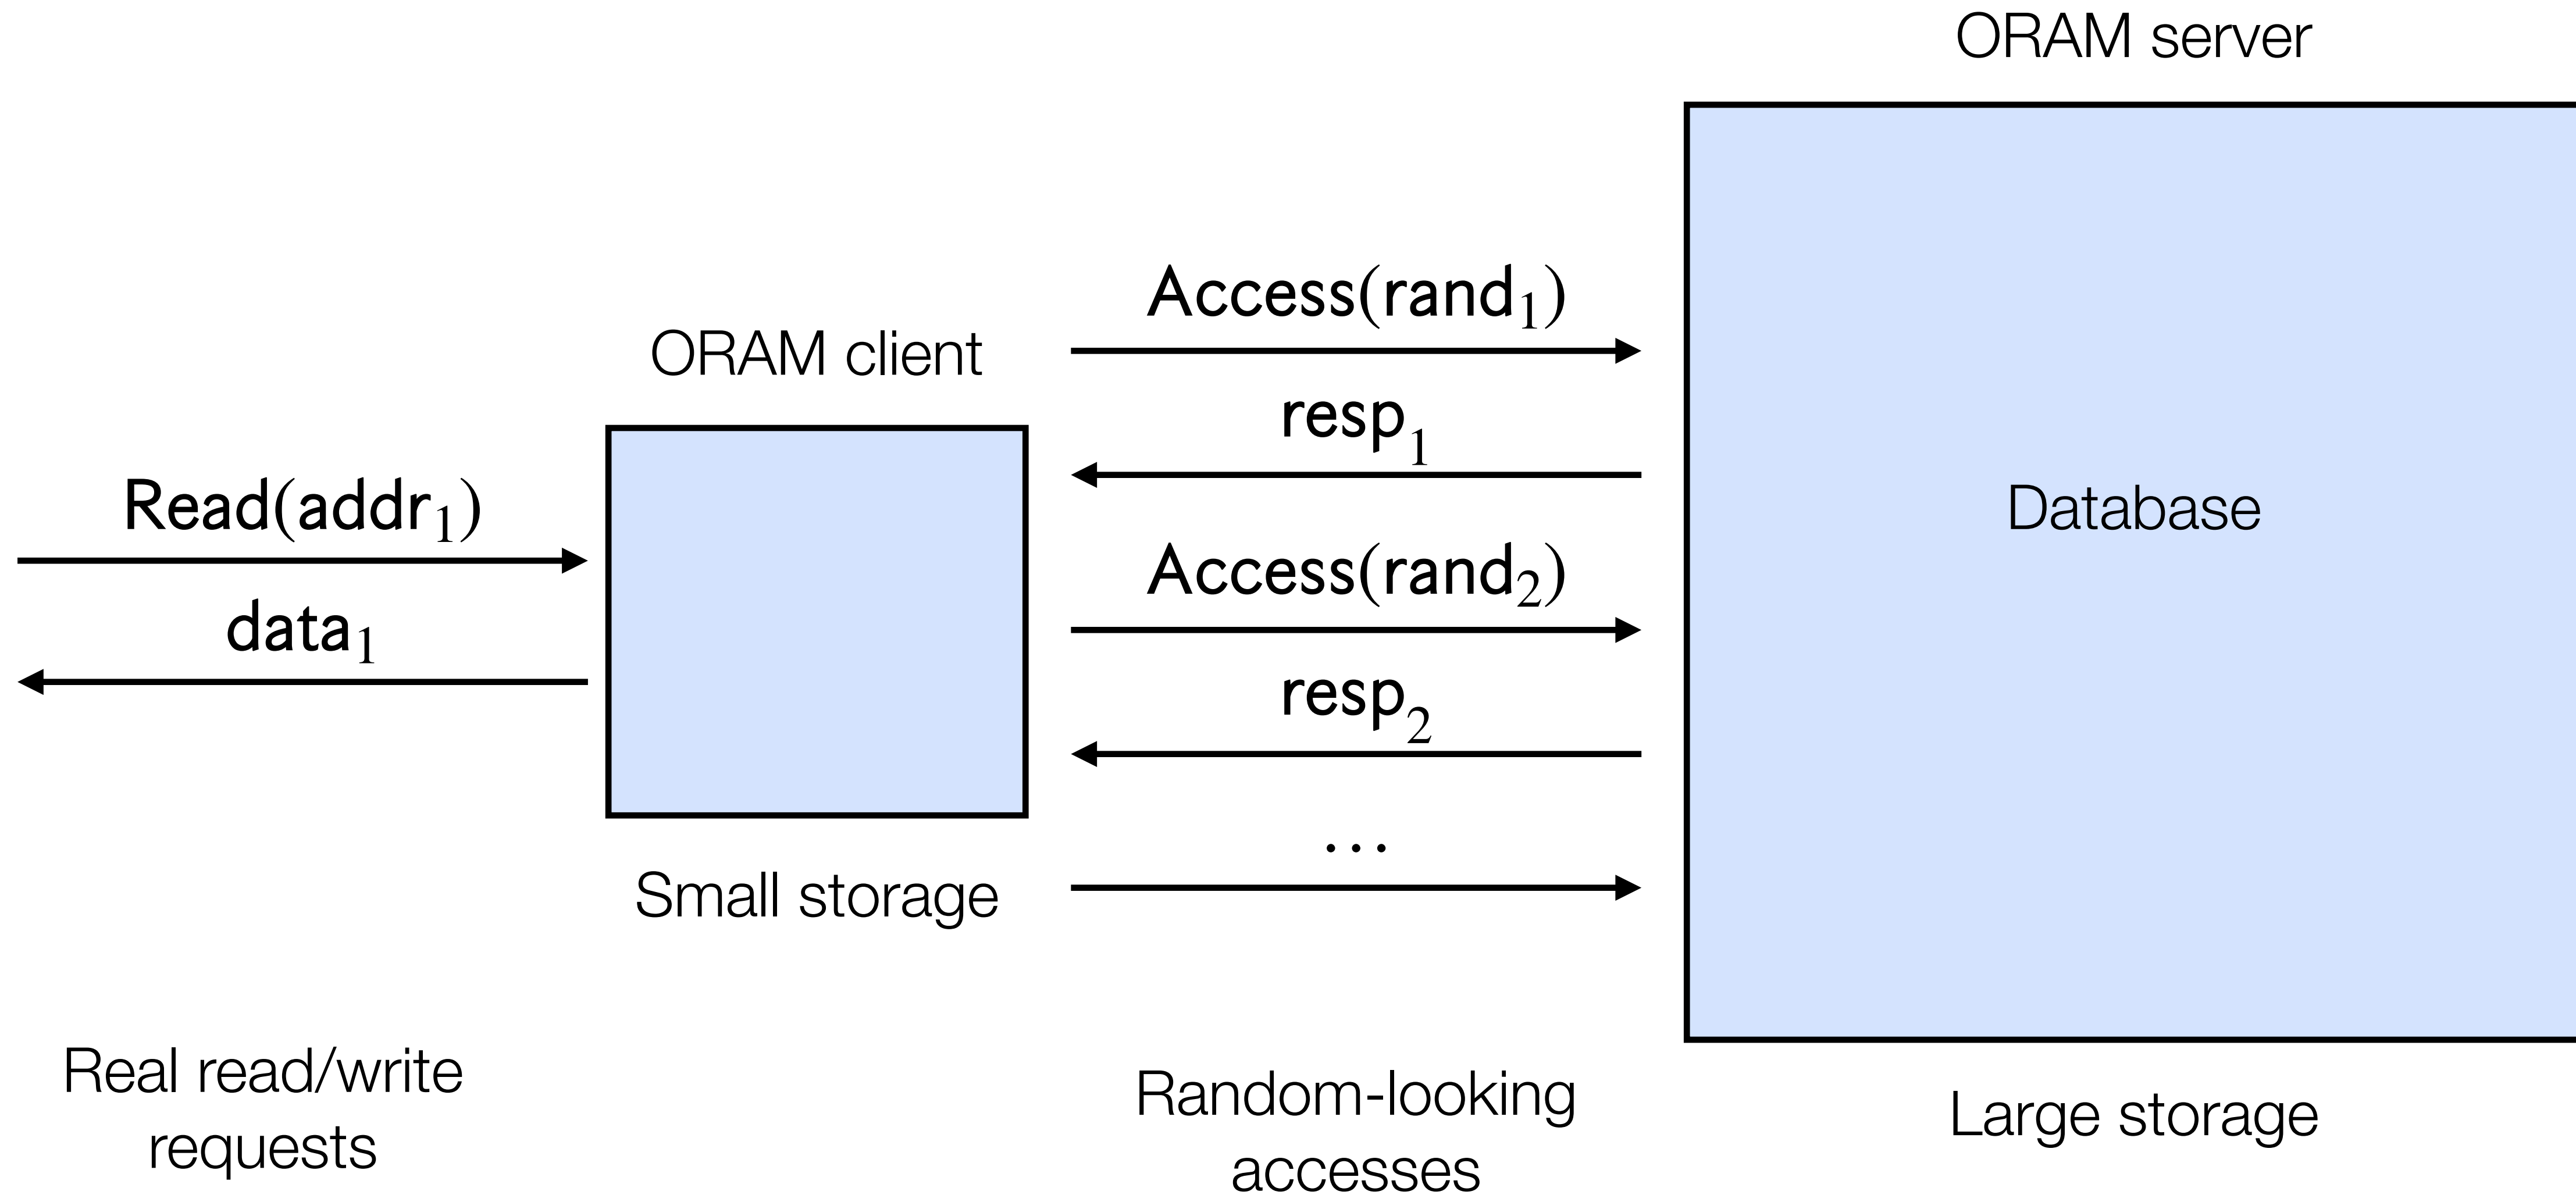
Attacker can infer that Alice and Bob
have the same medical condition



Access patterns can reveal sensitive information



Oblivious RAM

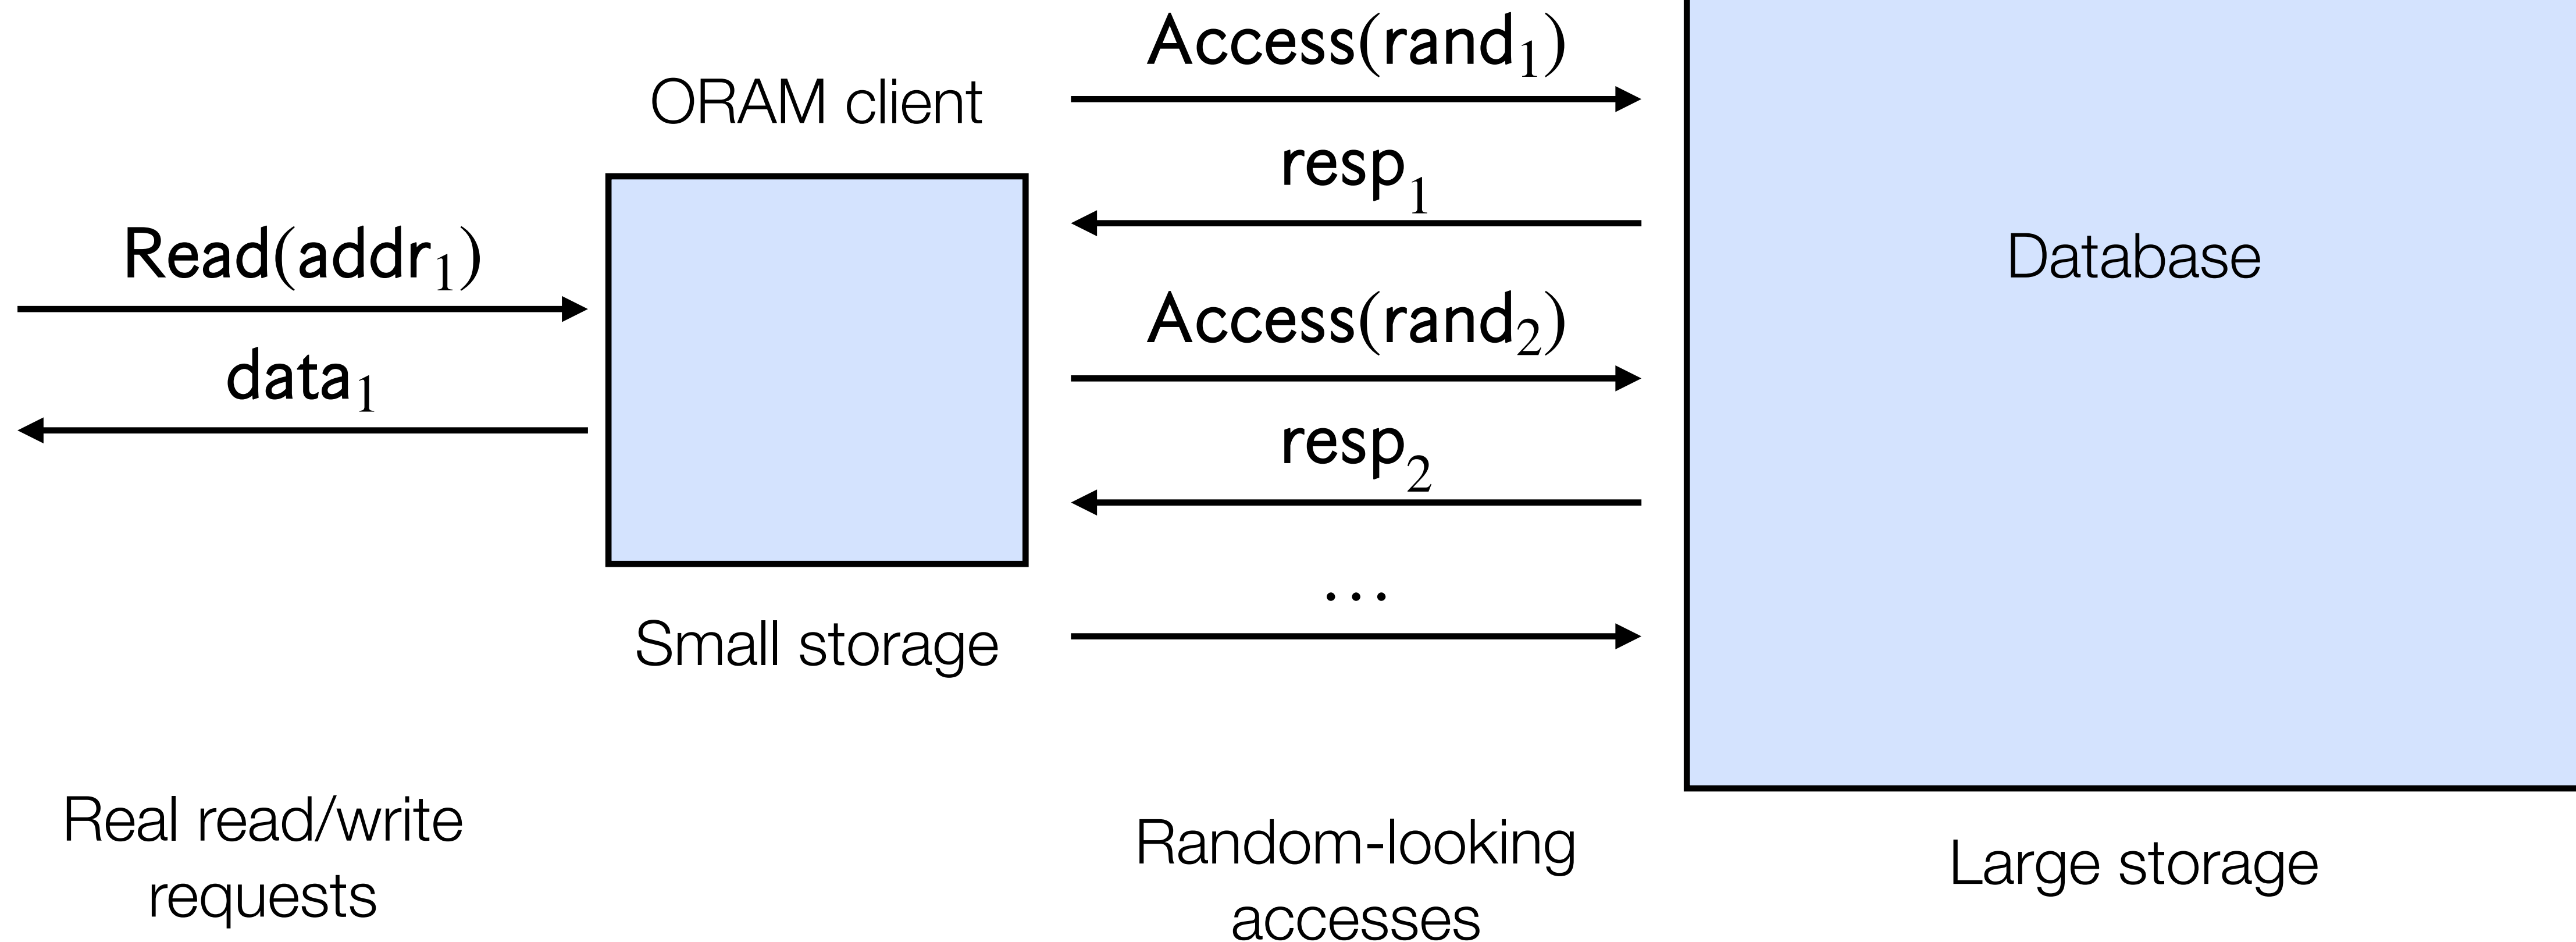


Oblivious RAM

Attacker that observes memory accesses “learns nothing” about the real read/write requests



ORAM server



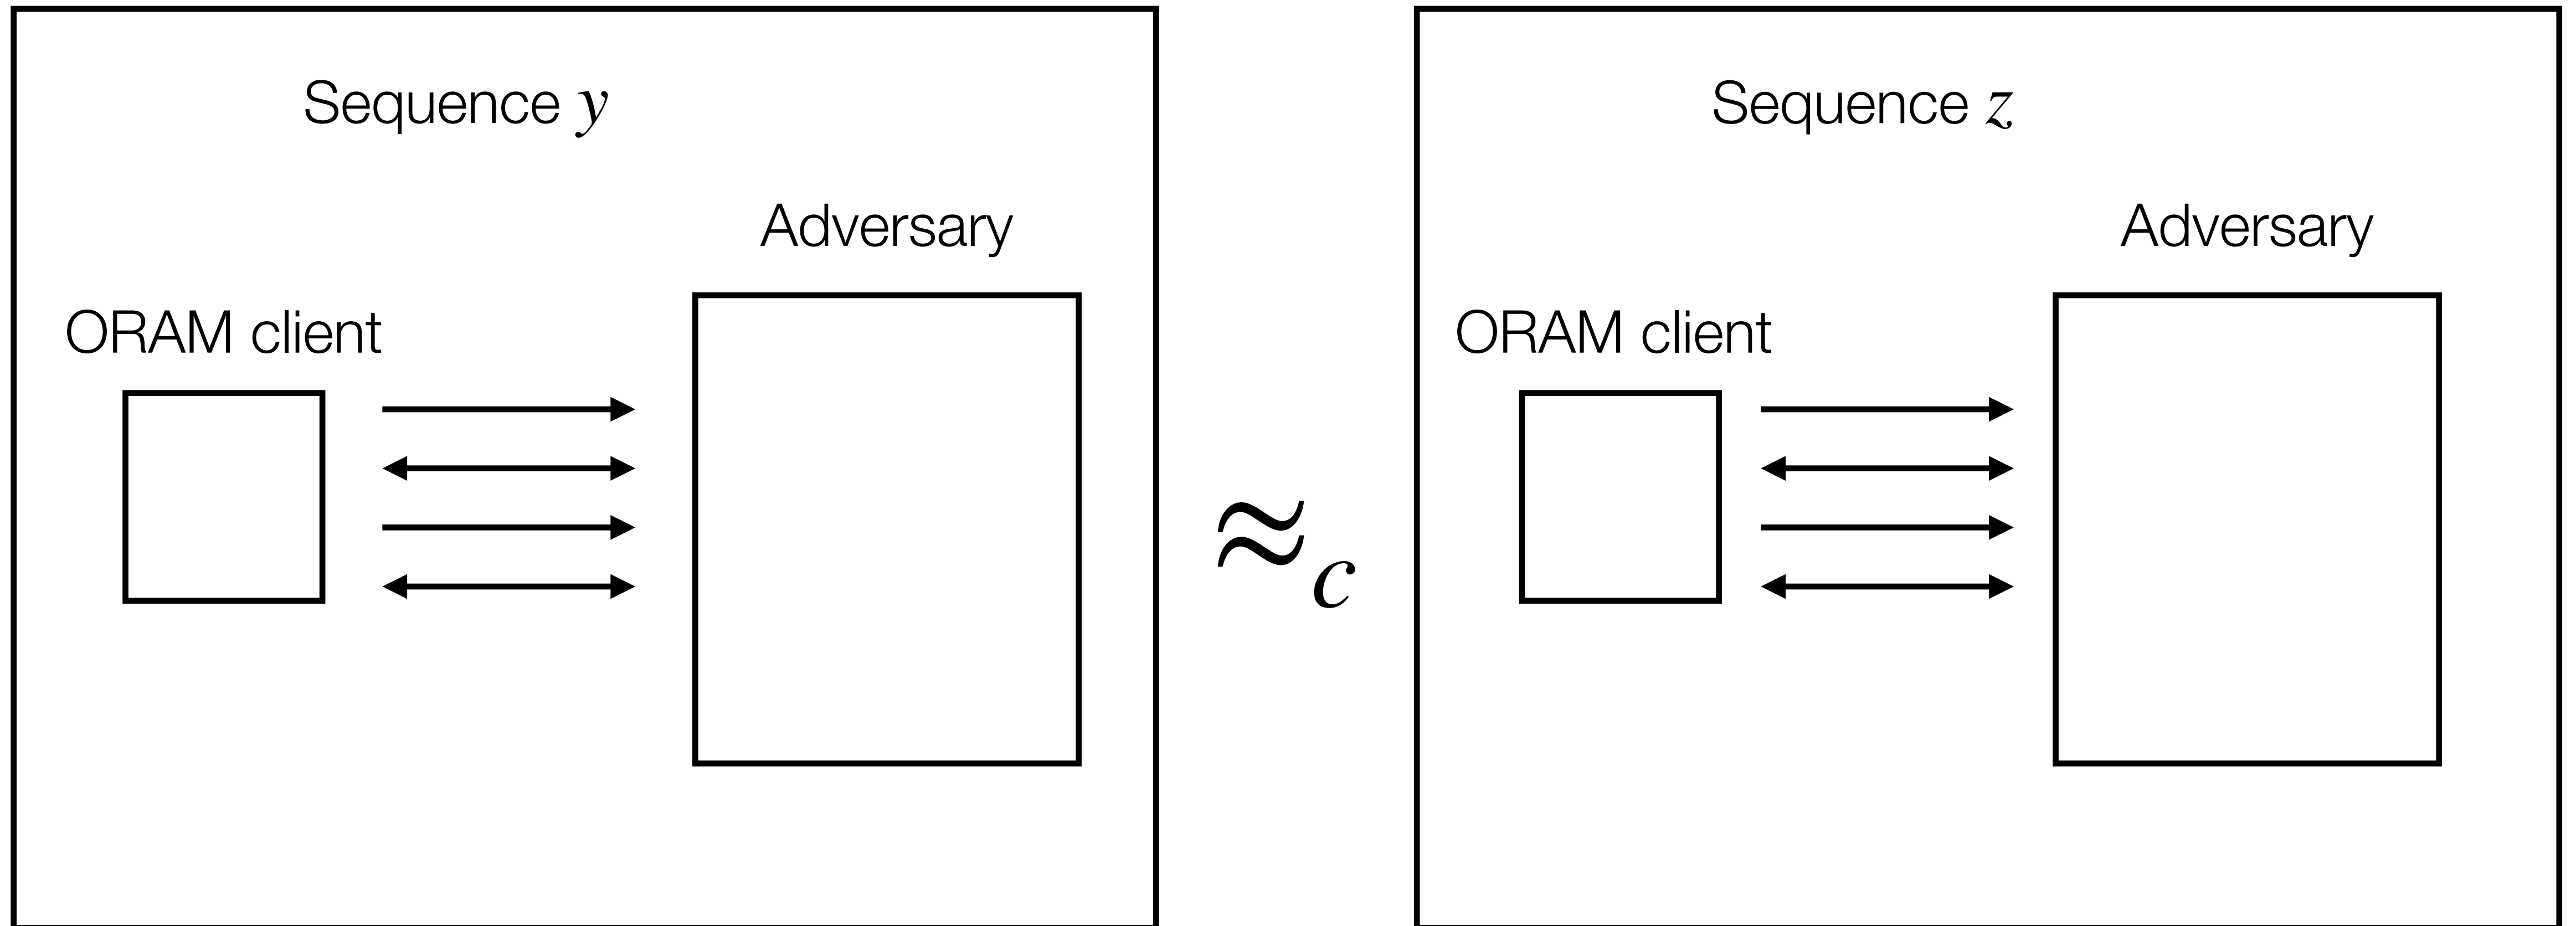
Definitions

Sequence of operations $y = ((op_1, addr_1, data_1), (op_2, addr_2, data_2), \dots)$
where $op \in \{\text{read}, \text{write}\}$

Correctness: An ORAM construction is correct if the responses from the ORAM for a sequence of operations y matches the responses from a standard RAM (with overwhelming probability).

Security: For any two request sequences y, z of the same length, their access patterns (i.e., interactions between ORAM client and server) are indistinguishable.


Security definition



Outline

1. Overview
- 2. Square-root construction**
3. Hierarchical construction
4. Limitations
5. Student presentation: PathORAM

A very simple, very expensive construction



Why private?

Server touched every element, and so the client could be accessing any element

Square-root ORAM construction

Goldreich and Ostrovsky

- Can we reduce the costs?
- Take advantage of:
 - Randomness
 - Large, mutable server state
 - Small, mutable client state

On the way to square-root ORAM construction

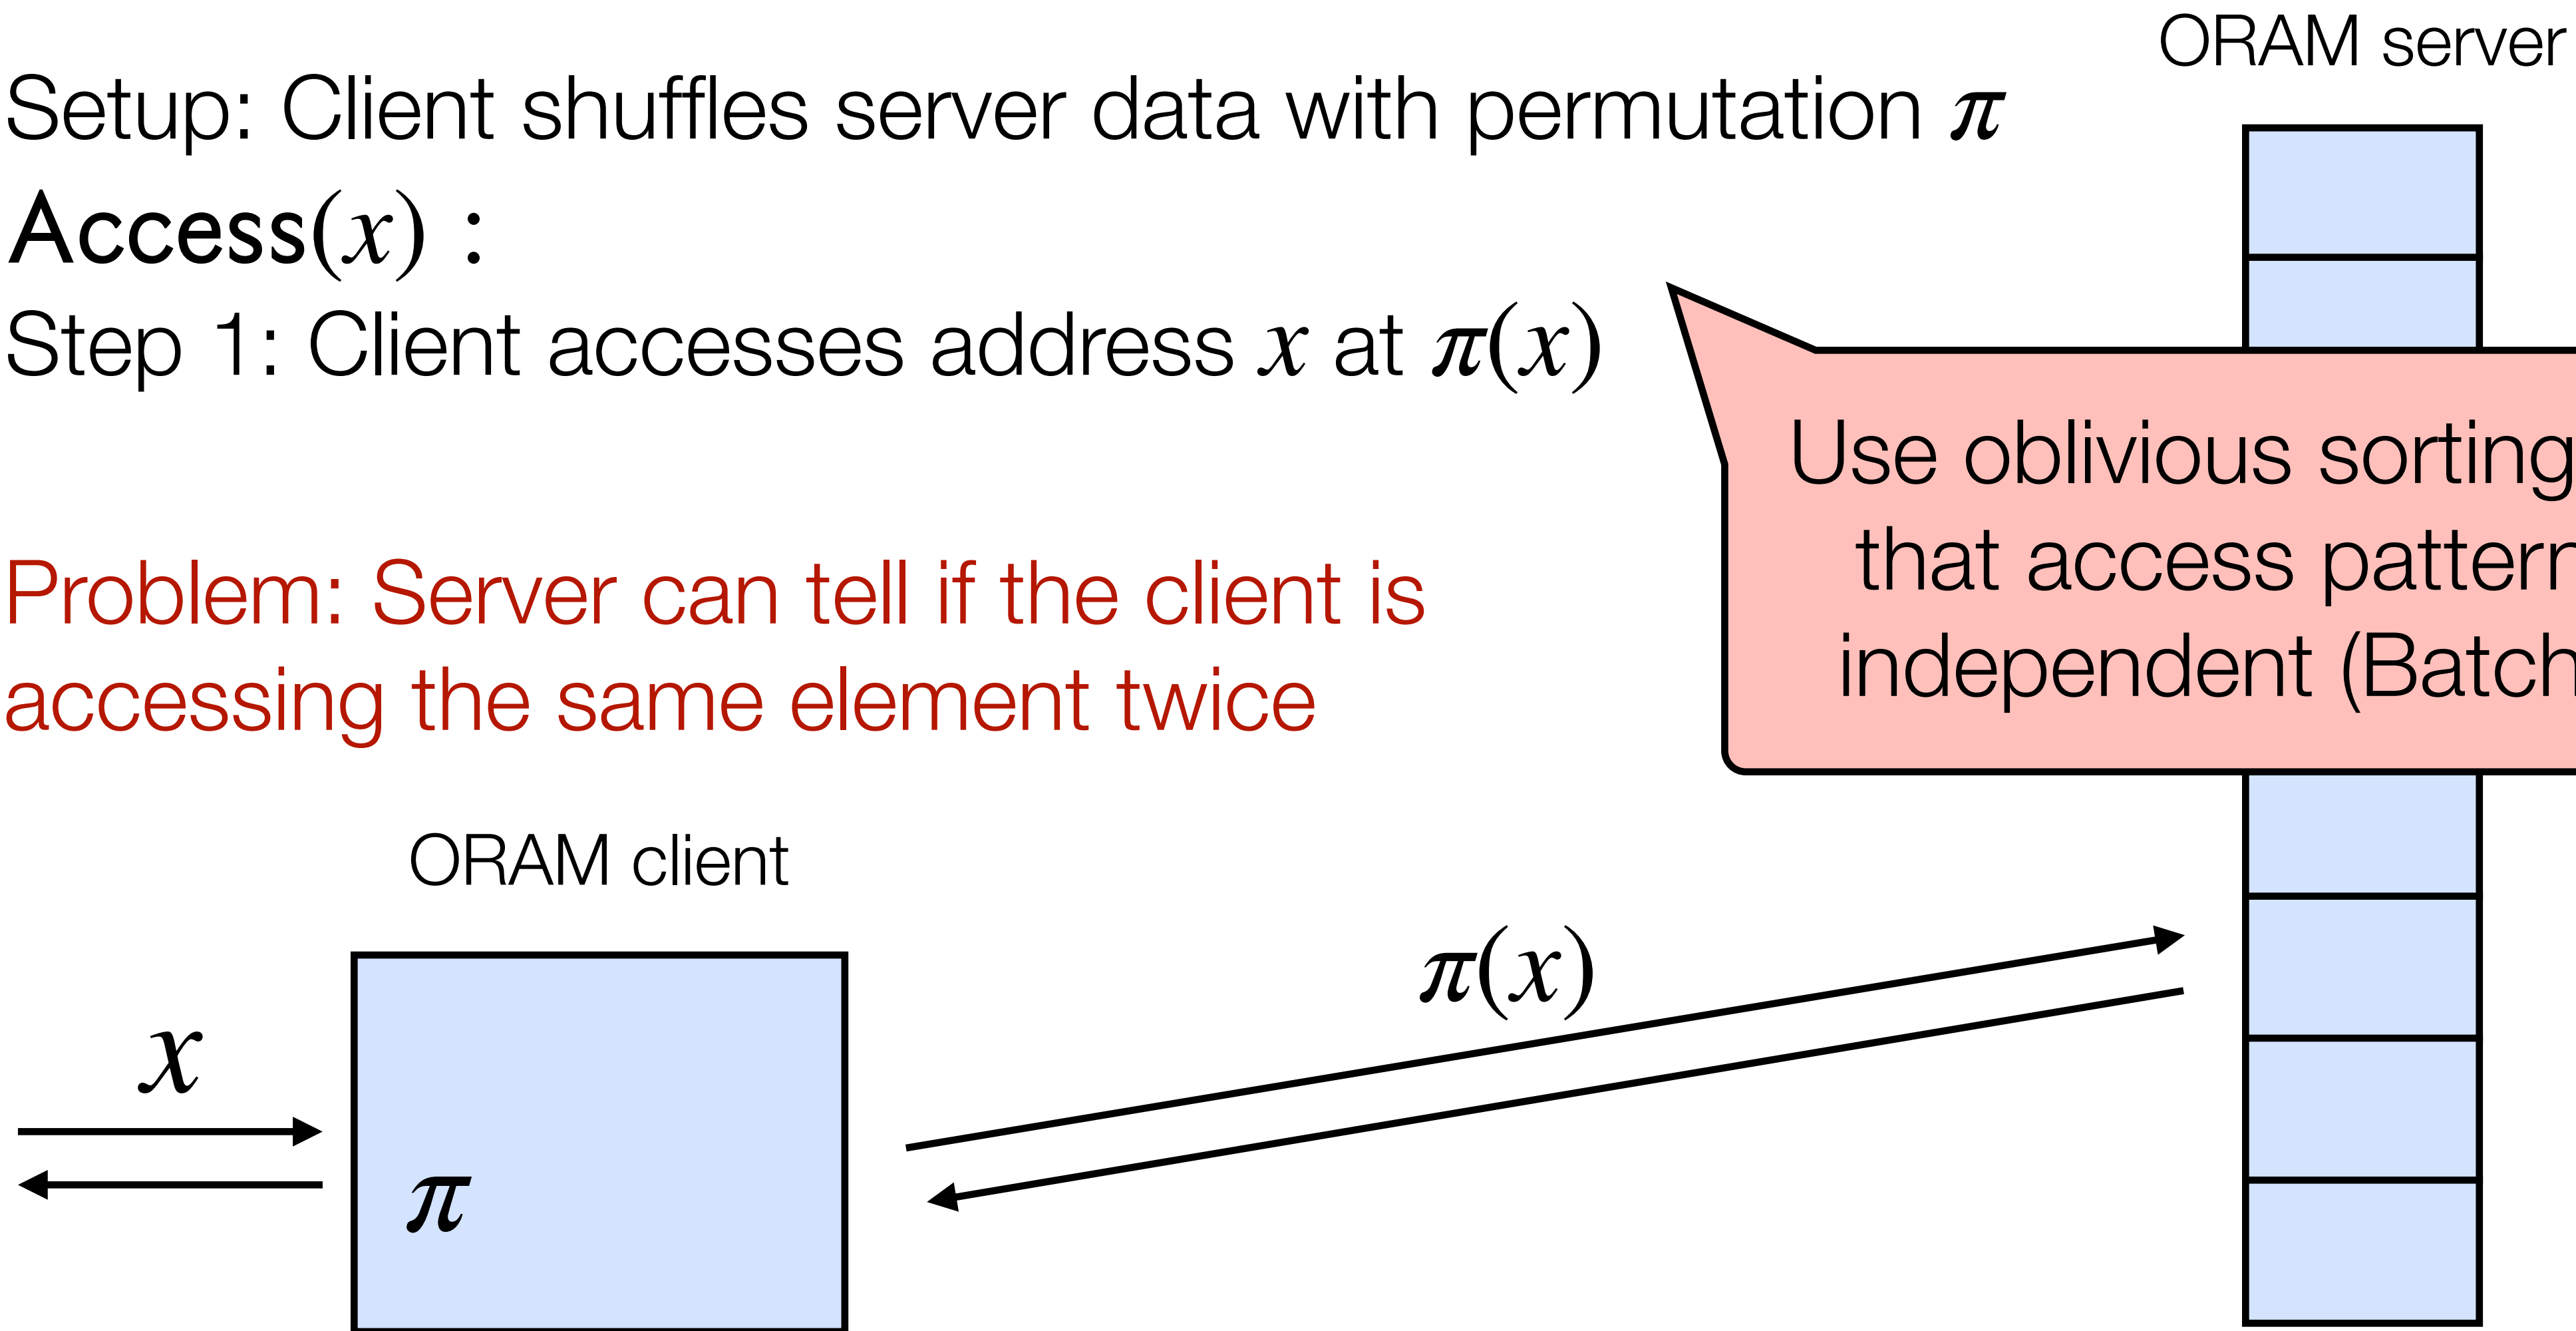
Setup: Client shuffles server data with permutation π

Access(x) :

Step 1: Client accesses address x at $\pi(x)$

Problem: Server can tell if the client is accessing the same element twice

Use oblivious sorting algorithm so that access patterns are data-independent (Batcher's sorting)



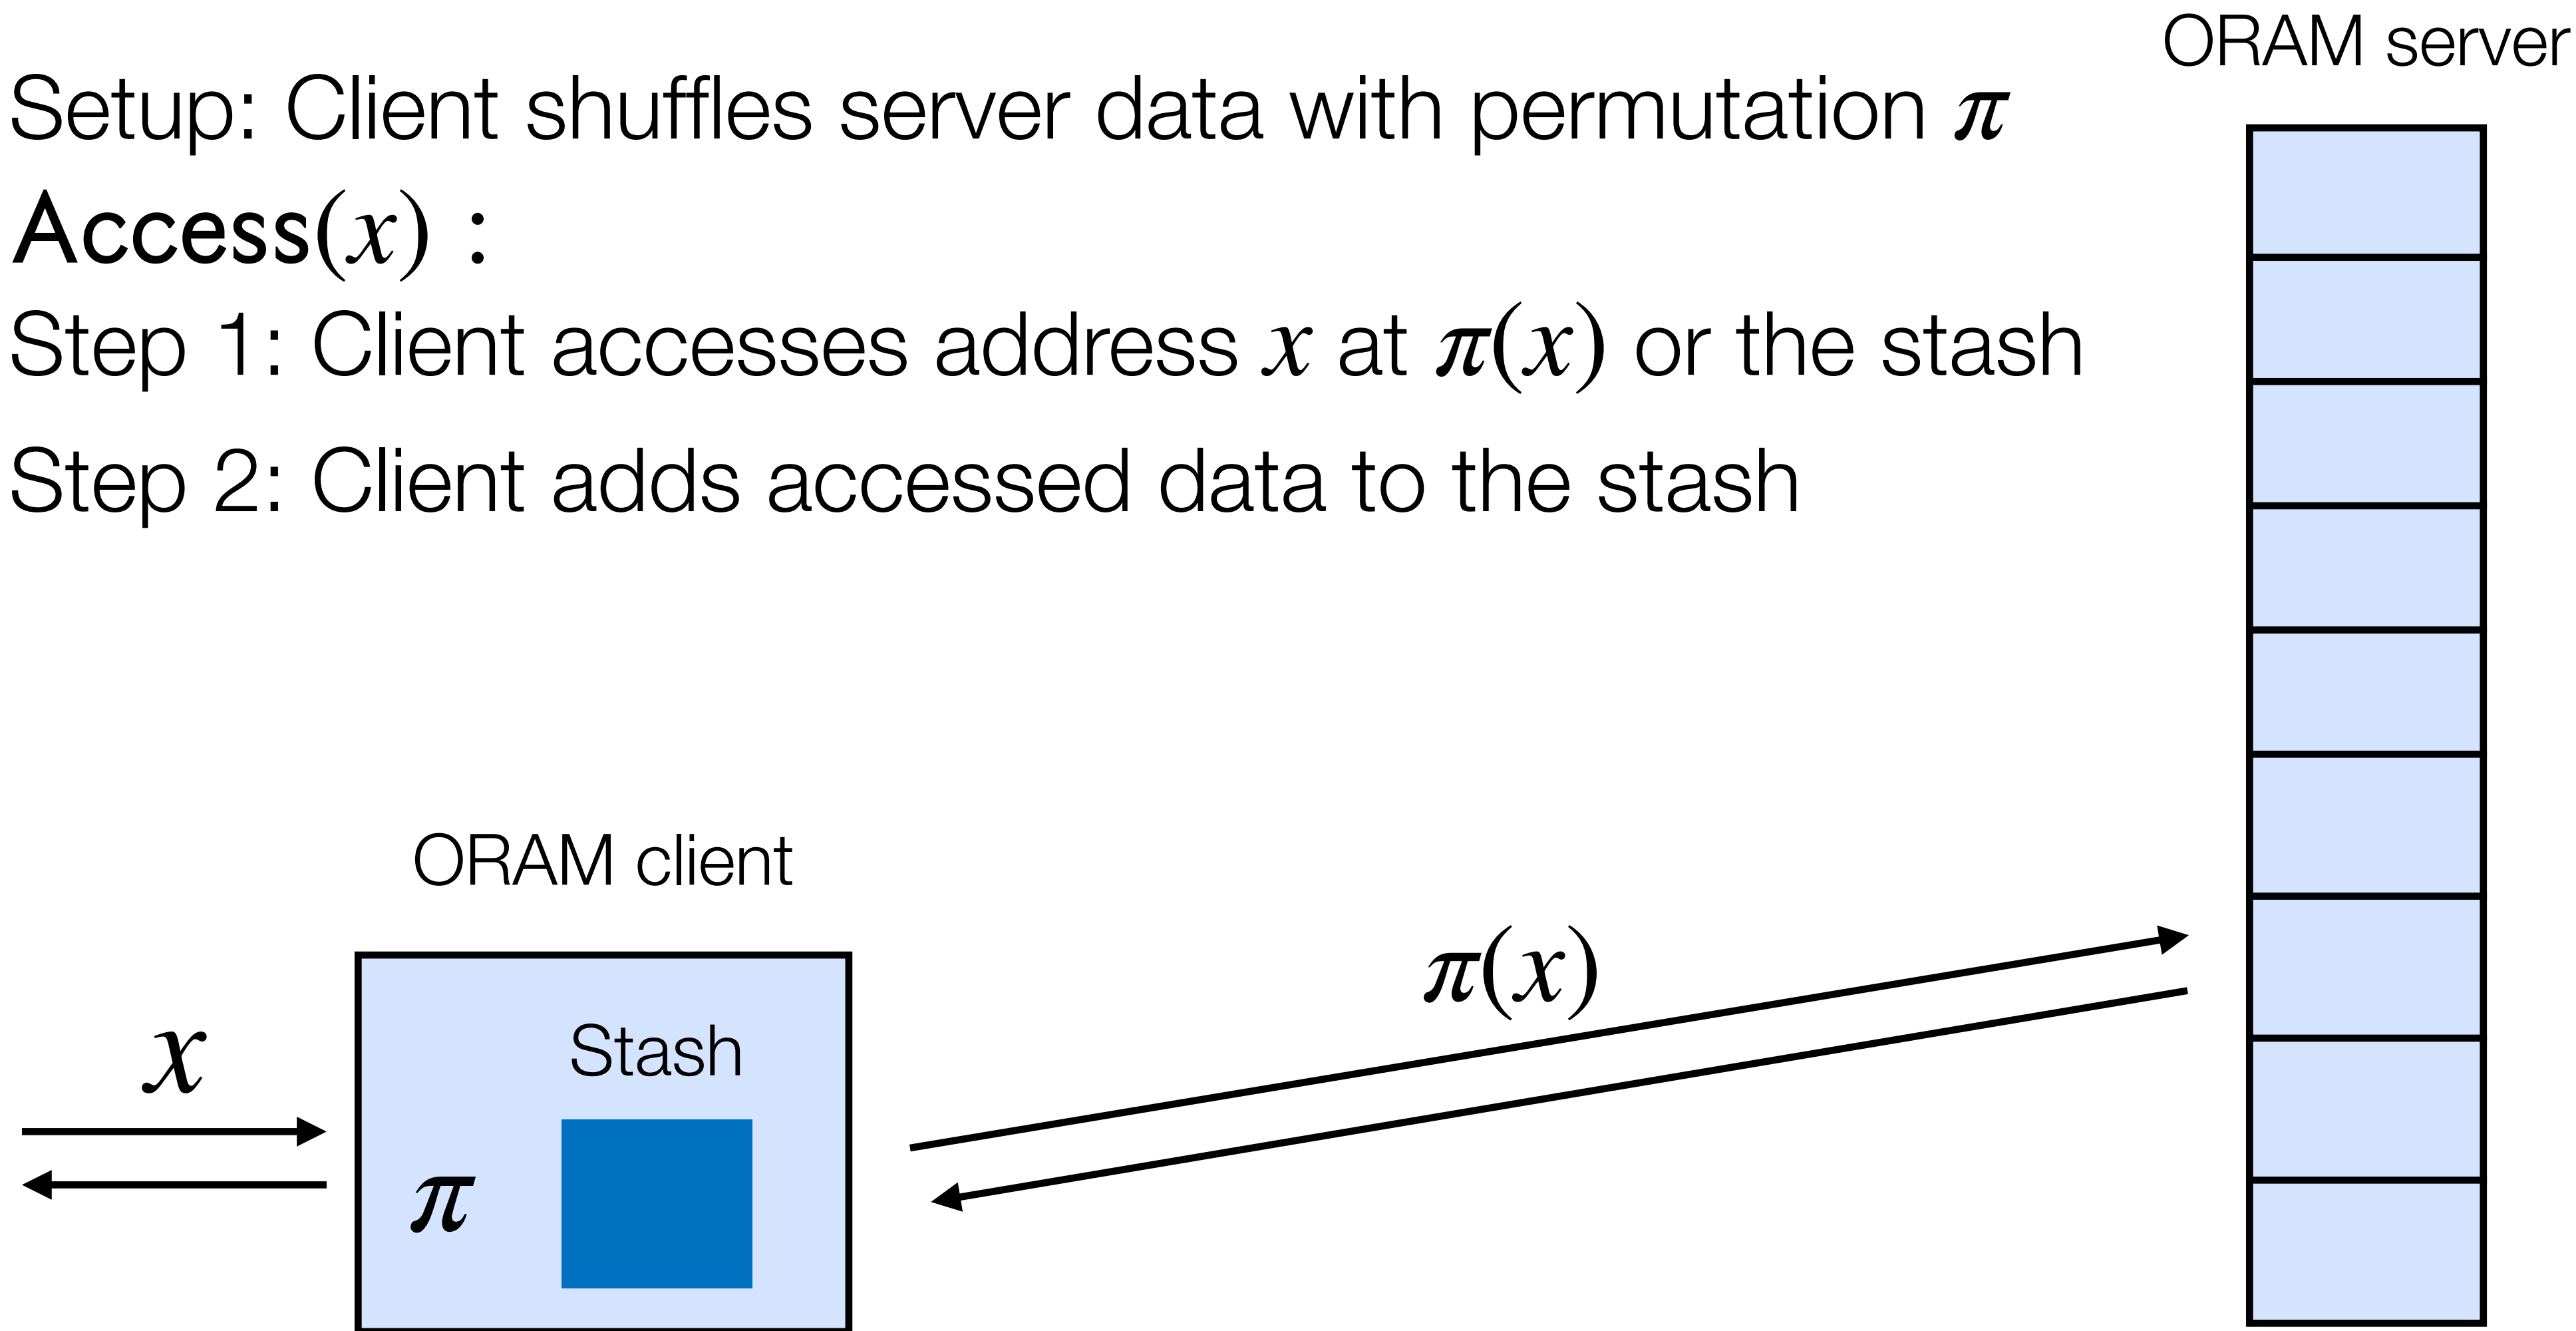
On the way to square-root ORAM construction

Setup: Client shuffles server data with permutation π

Access(x) :

Step 1: Client accesses address x at $\pi(x)$ or the stash

Step 2: Client adds accessed data to the stash



Stash stores all elements fetched with π


On the way to square-root ORAM construction

Setup: Client shuffles server data with permutation π

Access(x) :

Step 1: Client accesses address x at $\pi(x)$ or the stash

Step 2: Client adds accessed data to the stash



Problem: How to keep stash from growing indefinitely?

On the way to square-root ORAM construction

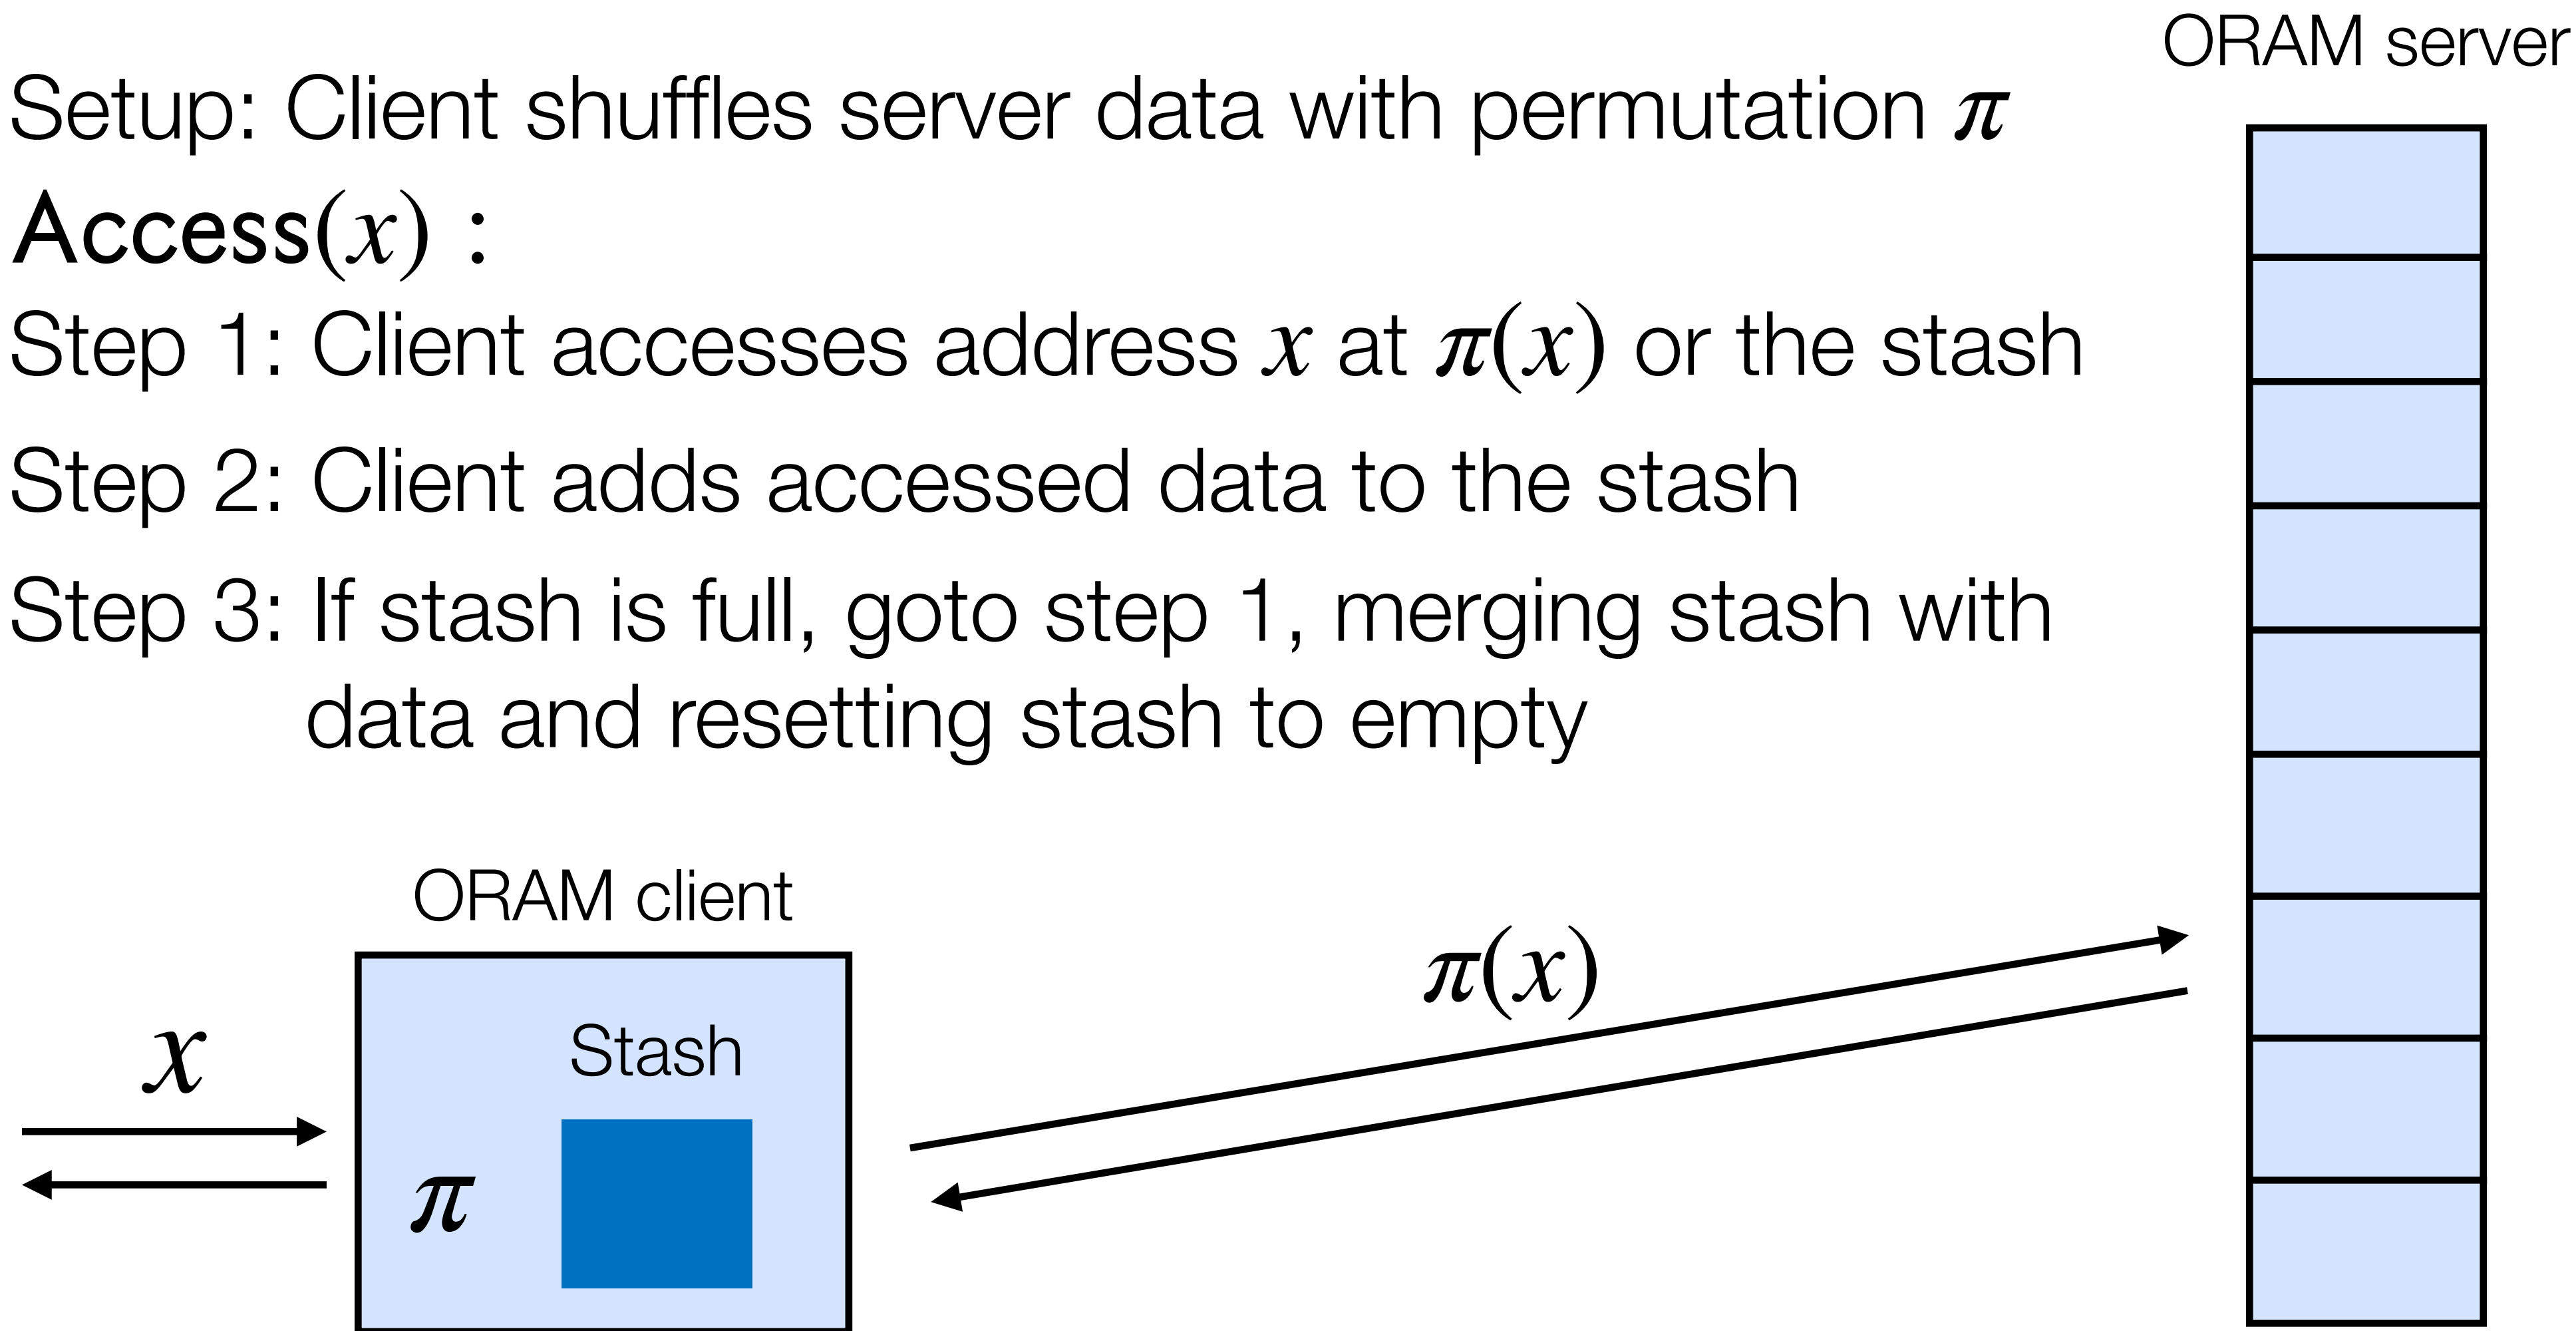
Setup: Client shuffles server data with permutation π

Access(x) :

Step 1: Client accesses address x at $\pi(x)$ or the stash

Step 2: Client adds accessed data to the stash

Step 3: If stash is full, goto step 1, merging stash with data and resetting stash to empty



Problem: How to hide if a request is in the stash?

On the way to square-root ORAM construction

Setup: Shuffle server data + dummies with permutation π


Access(x) :

Step 1: If x is in the stash, client accesses **dummy**.

Otherwise, client accesses $\pi(x)$

Step 2: Client adds accessed data to the stash

Step 3: If stash is full, goto step 1, merging stash with data and resetting stash to empty



On the way to square-root ORAM construction

Setup: Shuffle server data + dummies with permutation π


Access(x) :

Step 1: If x is in the stash, client accesses **dummy**.

Otherwise, client accesses $\pi(x)$

Step 2: Client adds accessed data to the stash

Step 3: If stash is full, goto step 1, merging stash with data and resetting stash to empty



Correctness

Setup: Shuffle server data + dummies with permutation π

Access(x) :

Step 1: If x is in the stash, client accesses **dummy**.

Otherwise, client accesses $\pi(x)$

Step 2: Client adds accessed data to the stash

Step 3: If stash is full, goto step 1, merging stash with data and resetting stash to empty

Correctness from:

- Correctness of shuffle
- Correct maintenance of the stash

Security

Setup: Shuffle server data + dummies with permutation π

Access(x) :

Step 1: If x is in the stash, client accesses **dummy**.

Otherwise, client accesses $\pi(x)$

Step 2: Client adds accessed data to the stash

Step 3: If stash is full, goto step 1, merging stash with data and resetting stash to empty

Security from:

- Permutation appears random to the server
- Each request accesses a unique element
- Server cannot distinguish a real request from a dummy request

Parameterizing square-root ORAM

Length n array, stash size s

Add s dummies

Setup: Shuffle server data + dummies with permutation π

$O(n \cdot \log^2 n)$
comparisons

Access(x) :

Step 1: If x is in the stash, client accesses **dummy**.

Otherwise, client accesses $\pi(x)$

Check s elements

Step 2: Client adds accessed data to the stash

Step 3: If stash is full, goto step 1, merging stash with data and resetting stash to empty

By setting stash size $s = \sqrt{n}$, amortized overhead is $O(\sqrt{n} \cdot \log^2 n)$

Client storage is $O(\sqrt{n})$

Outline

1. Overview
2. Square-root construction
- 3. Hierarchical construction**
4. Limitations
5. Student presentation: PathORAM


Hierarchical ORAM

[Goldreich, Ostrovsky]

High-level idea: Hierarchy of different-sized buffers to achieve logarithmic overhead


Square-root ORAM

$$n = 2^k$$




Hierarchical ORAM


$$n = 2^k$$




$$2^{k-1}$$



$$2^{k-2}$$



$$2^{k-3}$$




...

Shuffle smaller buffers
more frequently


Hierarchical ORAM

Access requires scanning over the stash and making a lookup in each buffer



Hierarchical ORAM

Shuffle smaller buffers more frequently




When stash is full:

- Let p be the index of the smallest unfilled buffer
- Shuffle together the stash and buffers $i < p$
- Put the results in buffer p

Hierarchical ORAM

Shuffle smaller buffers more frequently




When stash is full:

- Let p be the index of the smallest unfilled buffer
- Shuffle together the stash and buffers $i < p$
- Put the results in buffer p

Hierarchical ORAM

Shuffle smaller buffers more frequently




When stash is full:

- Let p be the index of the smallest unfilled buffer
- Shuffle together the stash and buffers $i < p$
- Put the results in buffer p

Hierarchical ORAM

Shuffle smaller buffers more frequently



When stash is full:

- Let p be the index of the smallest unfilled buffer
- Shuffle together the stash and buffers $i < p$
- Put the results in buffer p

Outline

1. Overview
2. Square-root construction
3. Hierarchical construction
- 4. Limitations**
5. Student presentation: PathORAM

Limitations of ORAM

- All elements must be the same length
- Increased cost compared to plaintext RAM accesses: lower bound of $O(\log n)$ accesses per operation [Larsen, Nielsen]
- ORAM is designed for a single client; does not directly support multiple clients
- Accesses where shuffling is required take longer than accesses without shuffling (not the case for tree-based ORAMs)
- Elements must be fetched in sequence, not in parallel (addressed in subsequent work, e.g., TaoStore)
- Only supports key-value lookups, but applications need other types of queries

Outline

1. Overview
2. Square-root construction
3. Hierarchical construction
4. Limitations
- 5. Student presentation: PathORAM**

References

Goldreich, Oded, and Rafail Ostrovsky. "Software protection and simulation on oblivious RAMs." *Journal of the ACM (JACM)* 43.3 (1996): 431-473.

Larsen, Kasper Green, and Jesper Buus Nielsen. "Yes, there is an oblivious RAM lower bound!." *Annual International Cryptology Conference*. Cham: Springer International Publishing, 2018.

Sahin, Cetin, et al. "Taostore: Overcoming asynchronicity in oblivious data storage." *2016 IEEE Symposium on Security and Privacy (SP)*. IEEE, 2016.

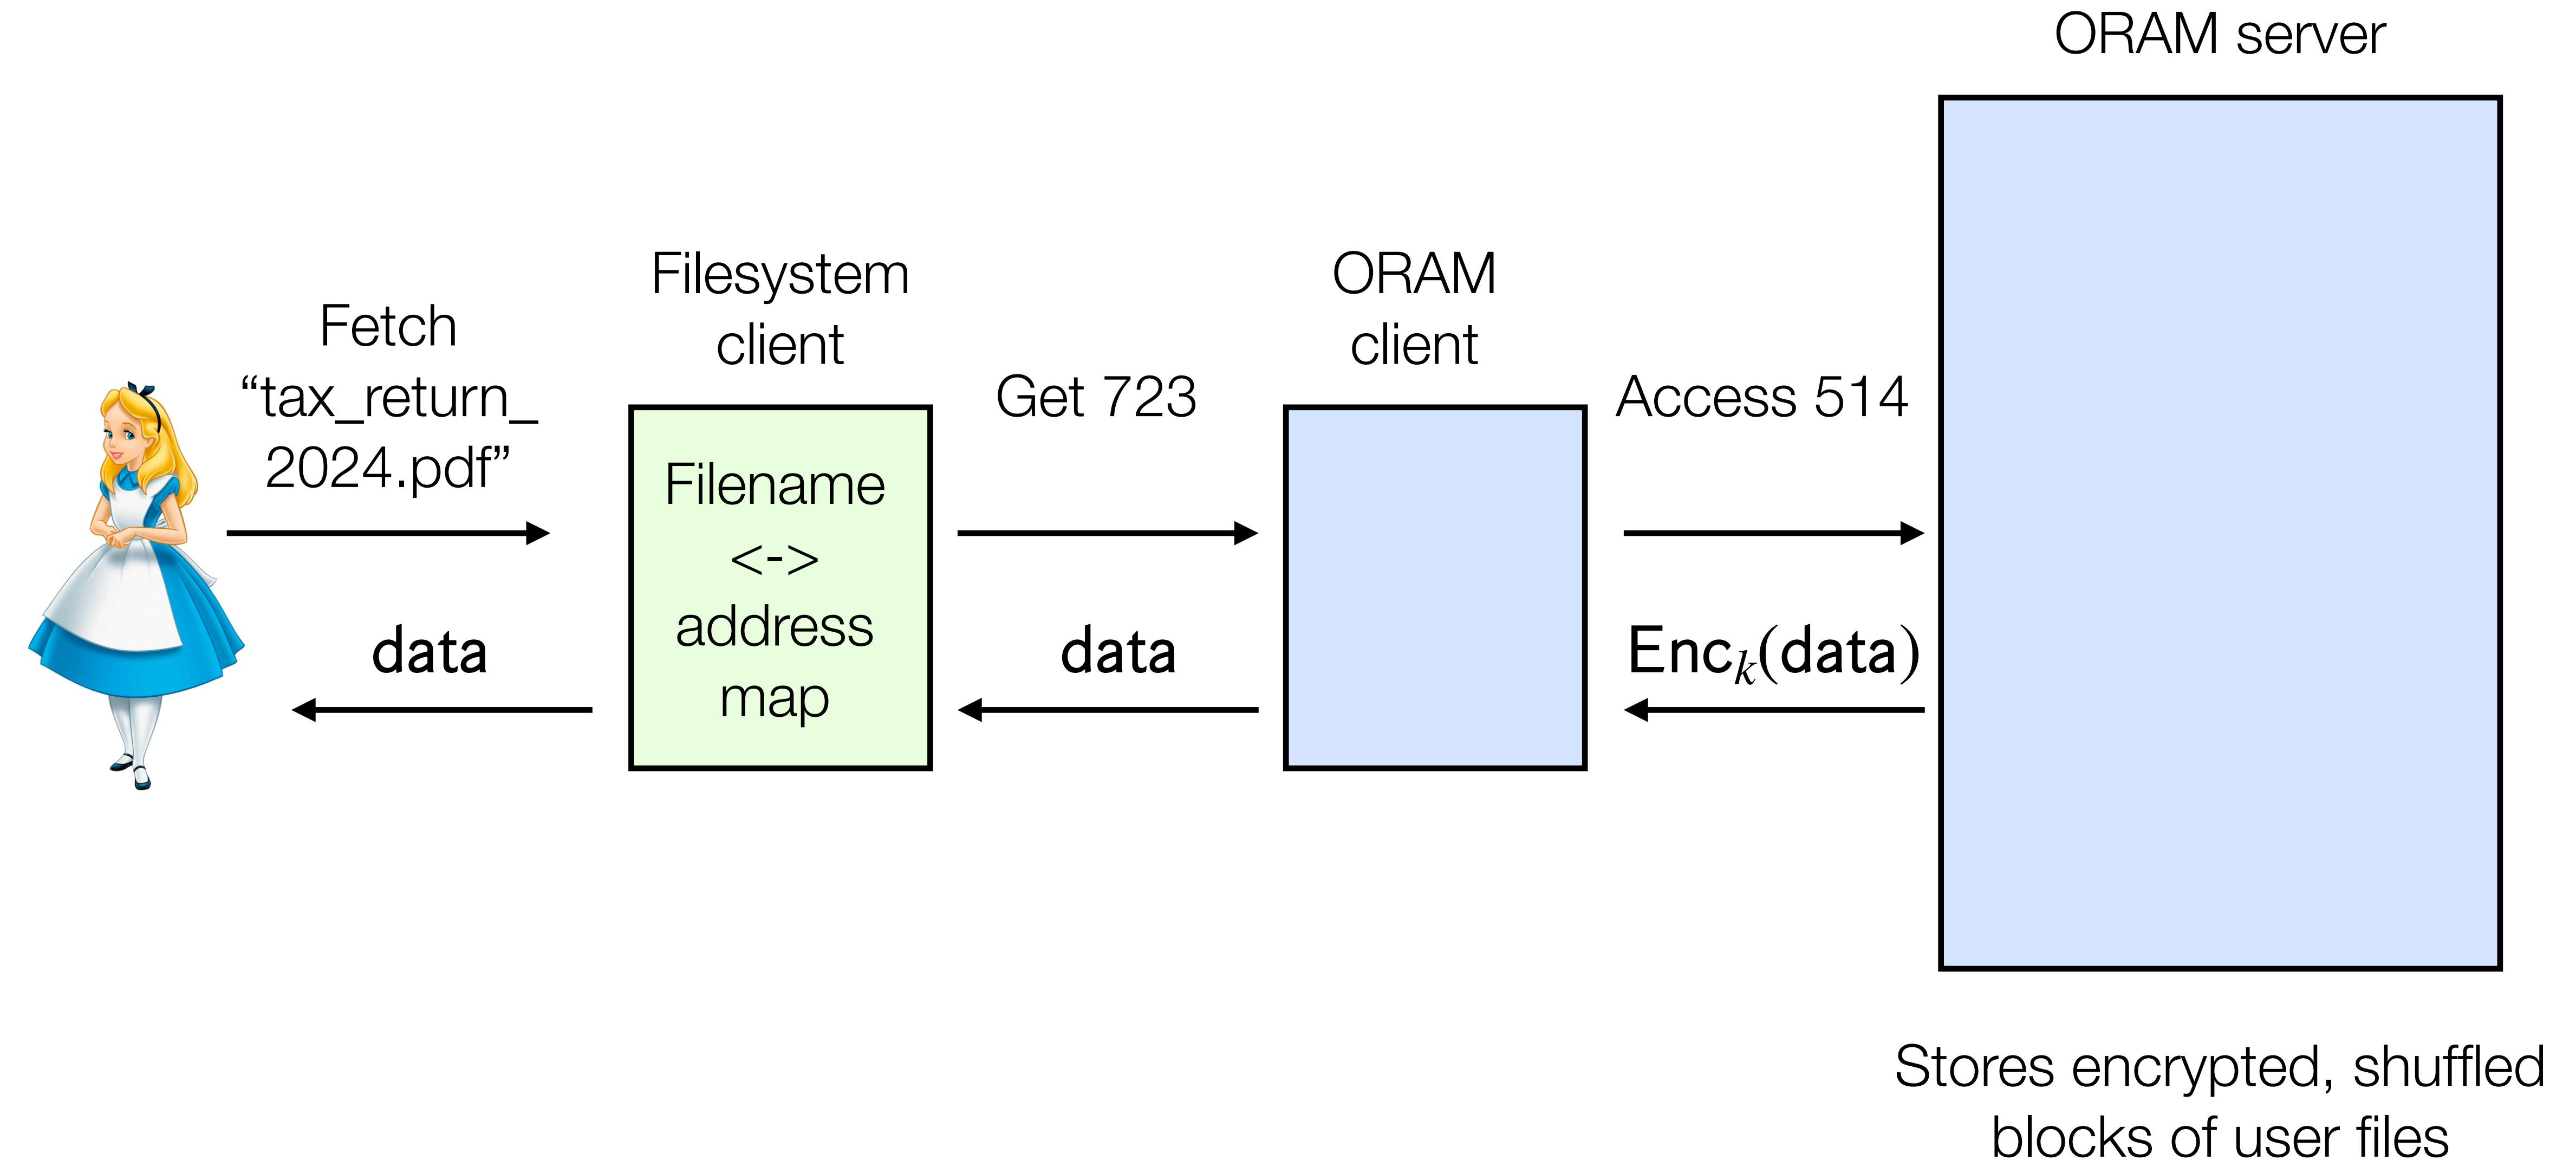
Stefanov, Emil, et al. "Path ORAM: an extremely simple oblivious RAM protocol." *Journal of the ACM (JACM)* 65.4 (2018): 1-26

<https://6893.csail.mit.edu/lec7.pdf>

Applications

- Oblivious filesystem: hide both the client's documents and the access patterns to these documents
- Oblivious search over files: hide both the client's documents and the client's queries

Application: Oblivious filesystem



Application: Oblivious search over encrypted files

ORAM server

