

# CS 350S: Privacy-Preserving Systems

Oblivious RAM


# Outline

1. Overview
2. Square-root construction
3. Hierarchical construction
4. Limitations
5. Student presentation: PathORAM

# Access patterns can reveal sensitive information



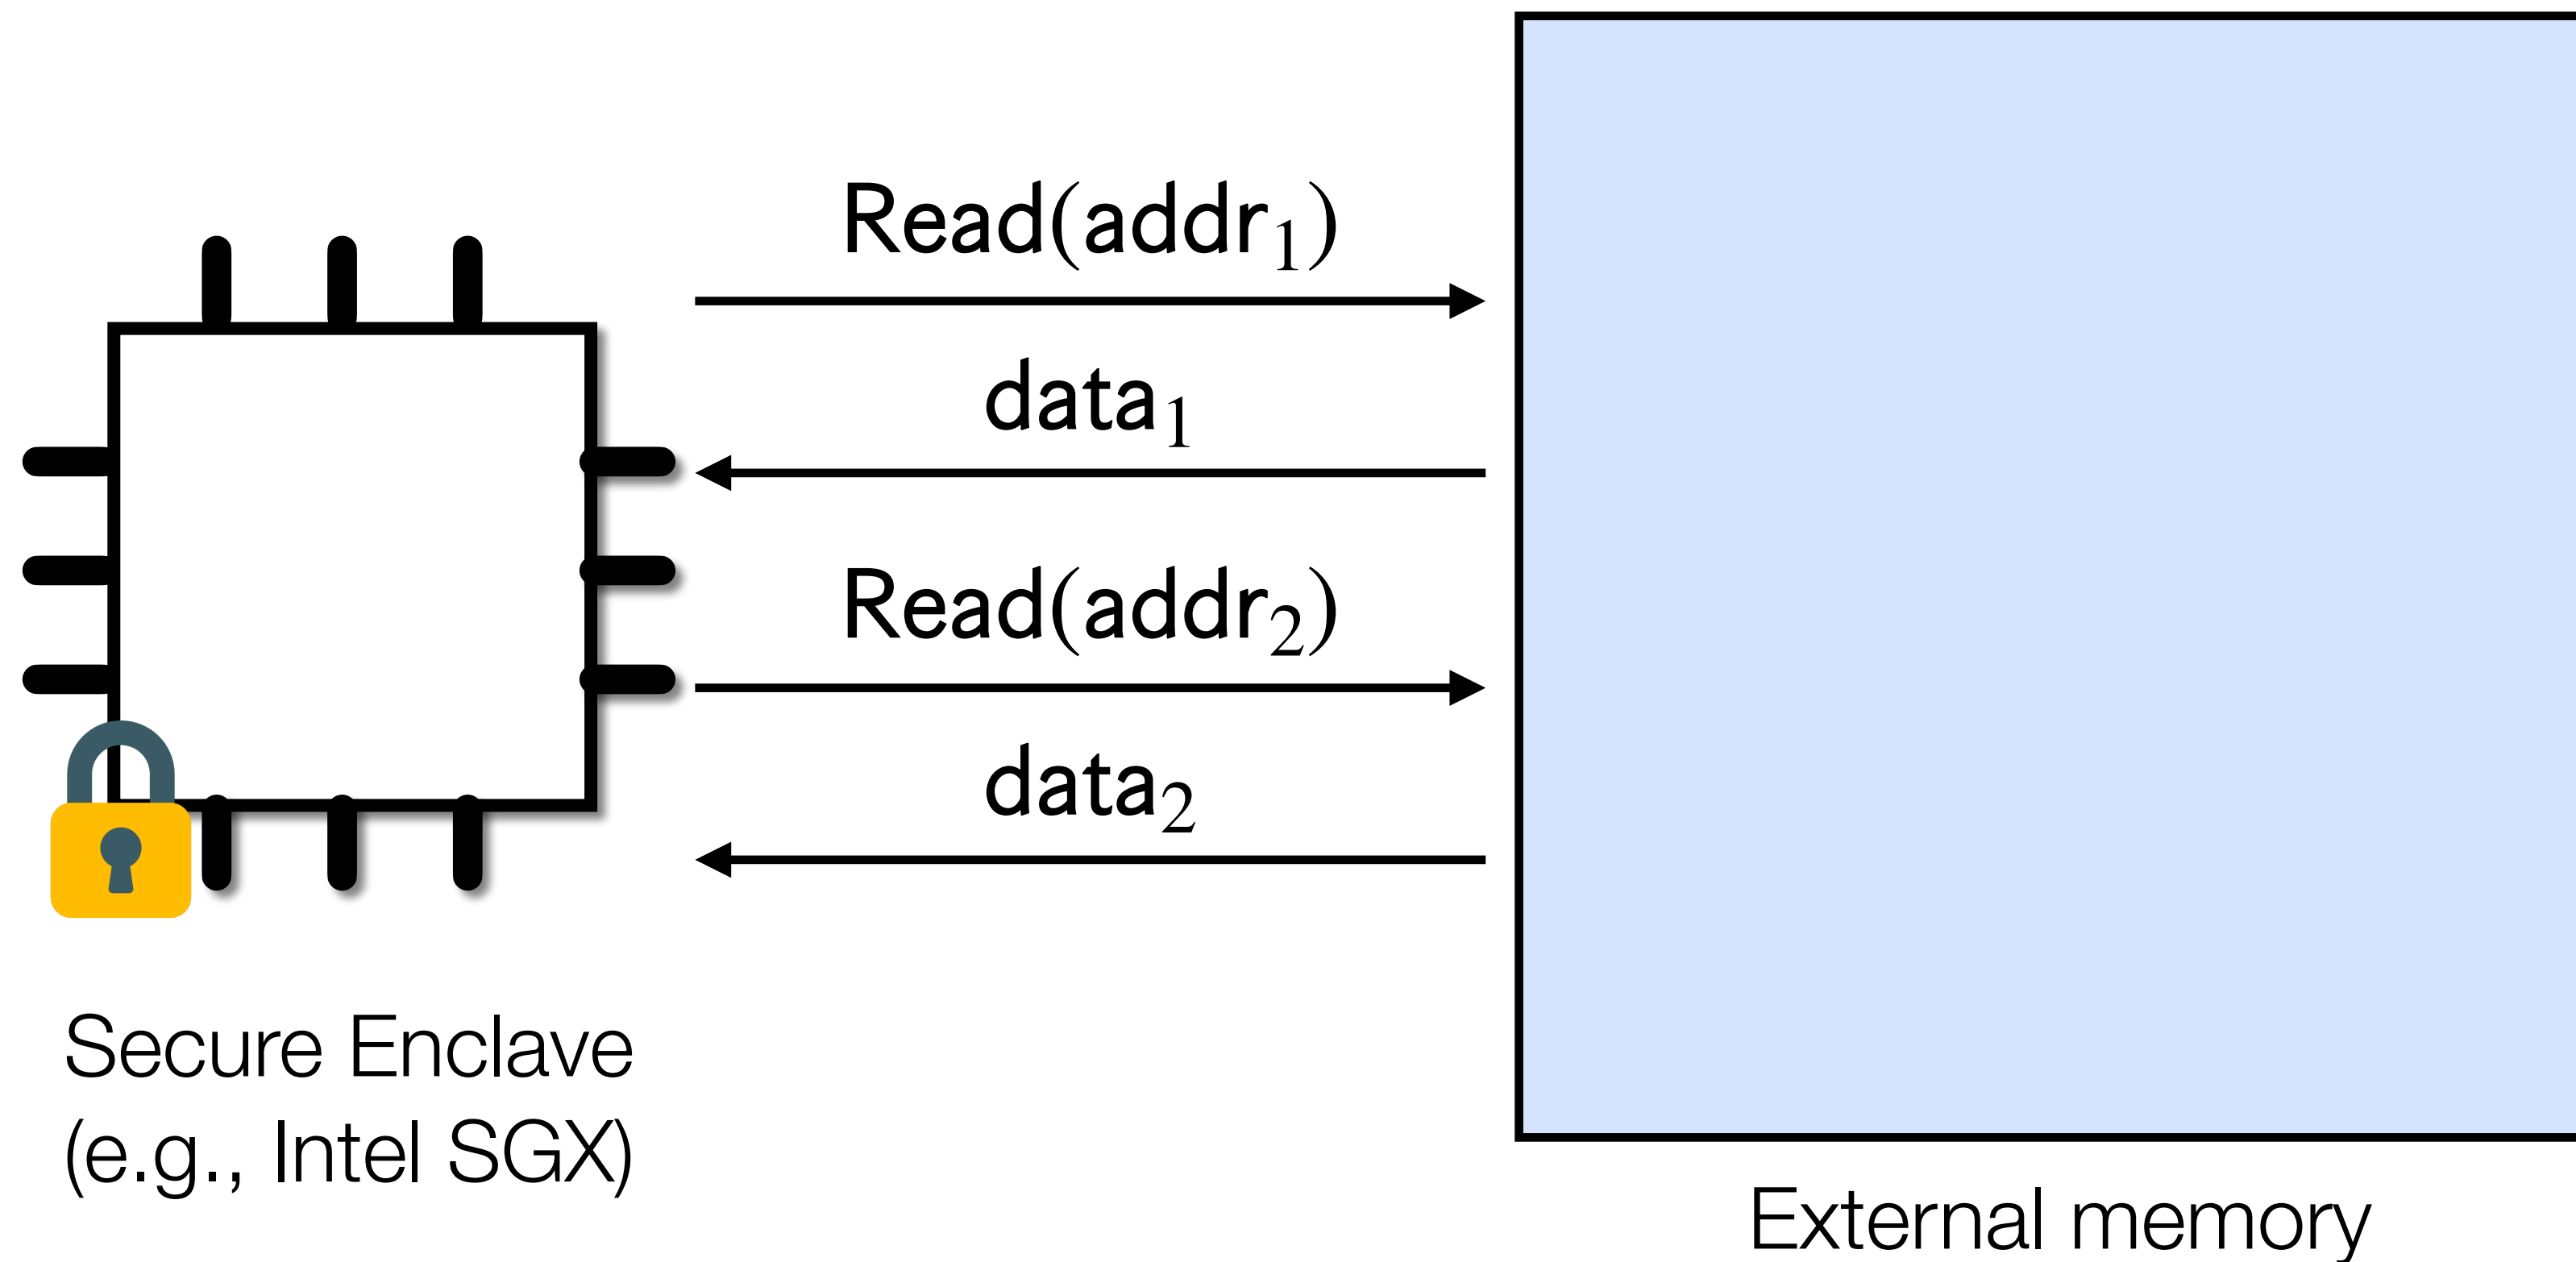
$\text{Enc}_k(\text{heart disease})$	$\text{Enc}_k(\text{info}_1)$
$\text{Enc}_k(\text{cancer})$	$\text{Enc}_k(\text{info}_2)$
$\text{Enc}_k(\text{diabetes})$	$\text{Enc}_k(\text{info}_3)$
...	...



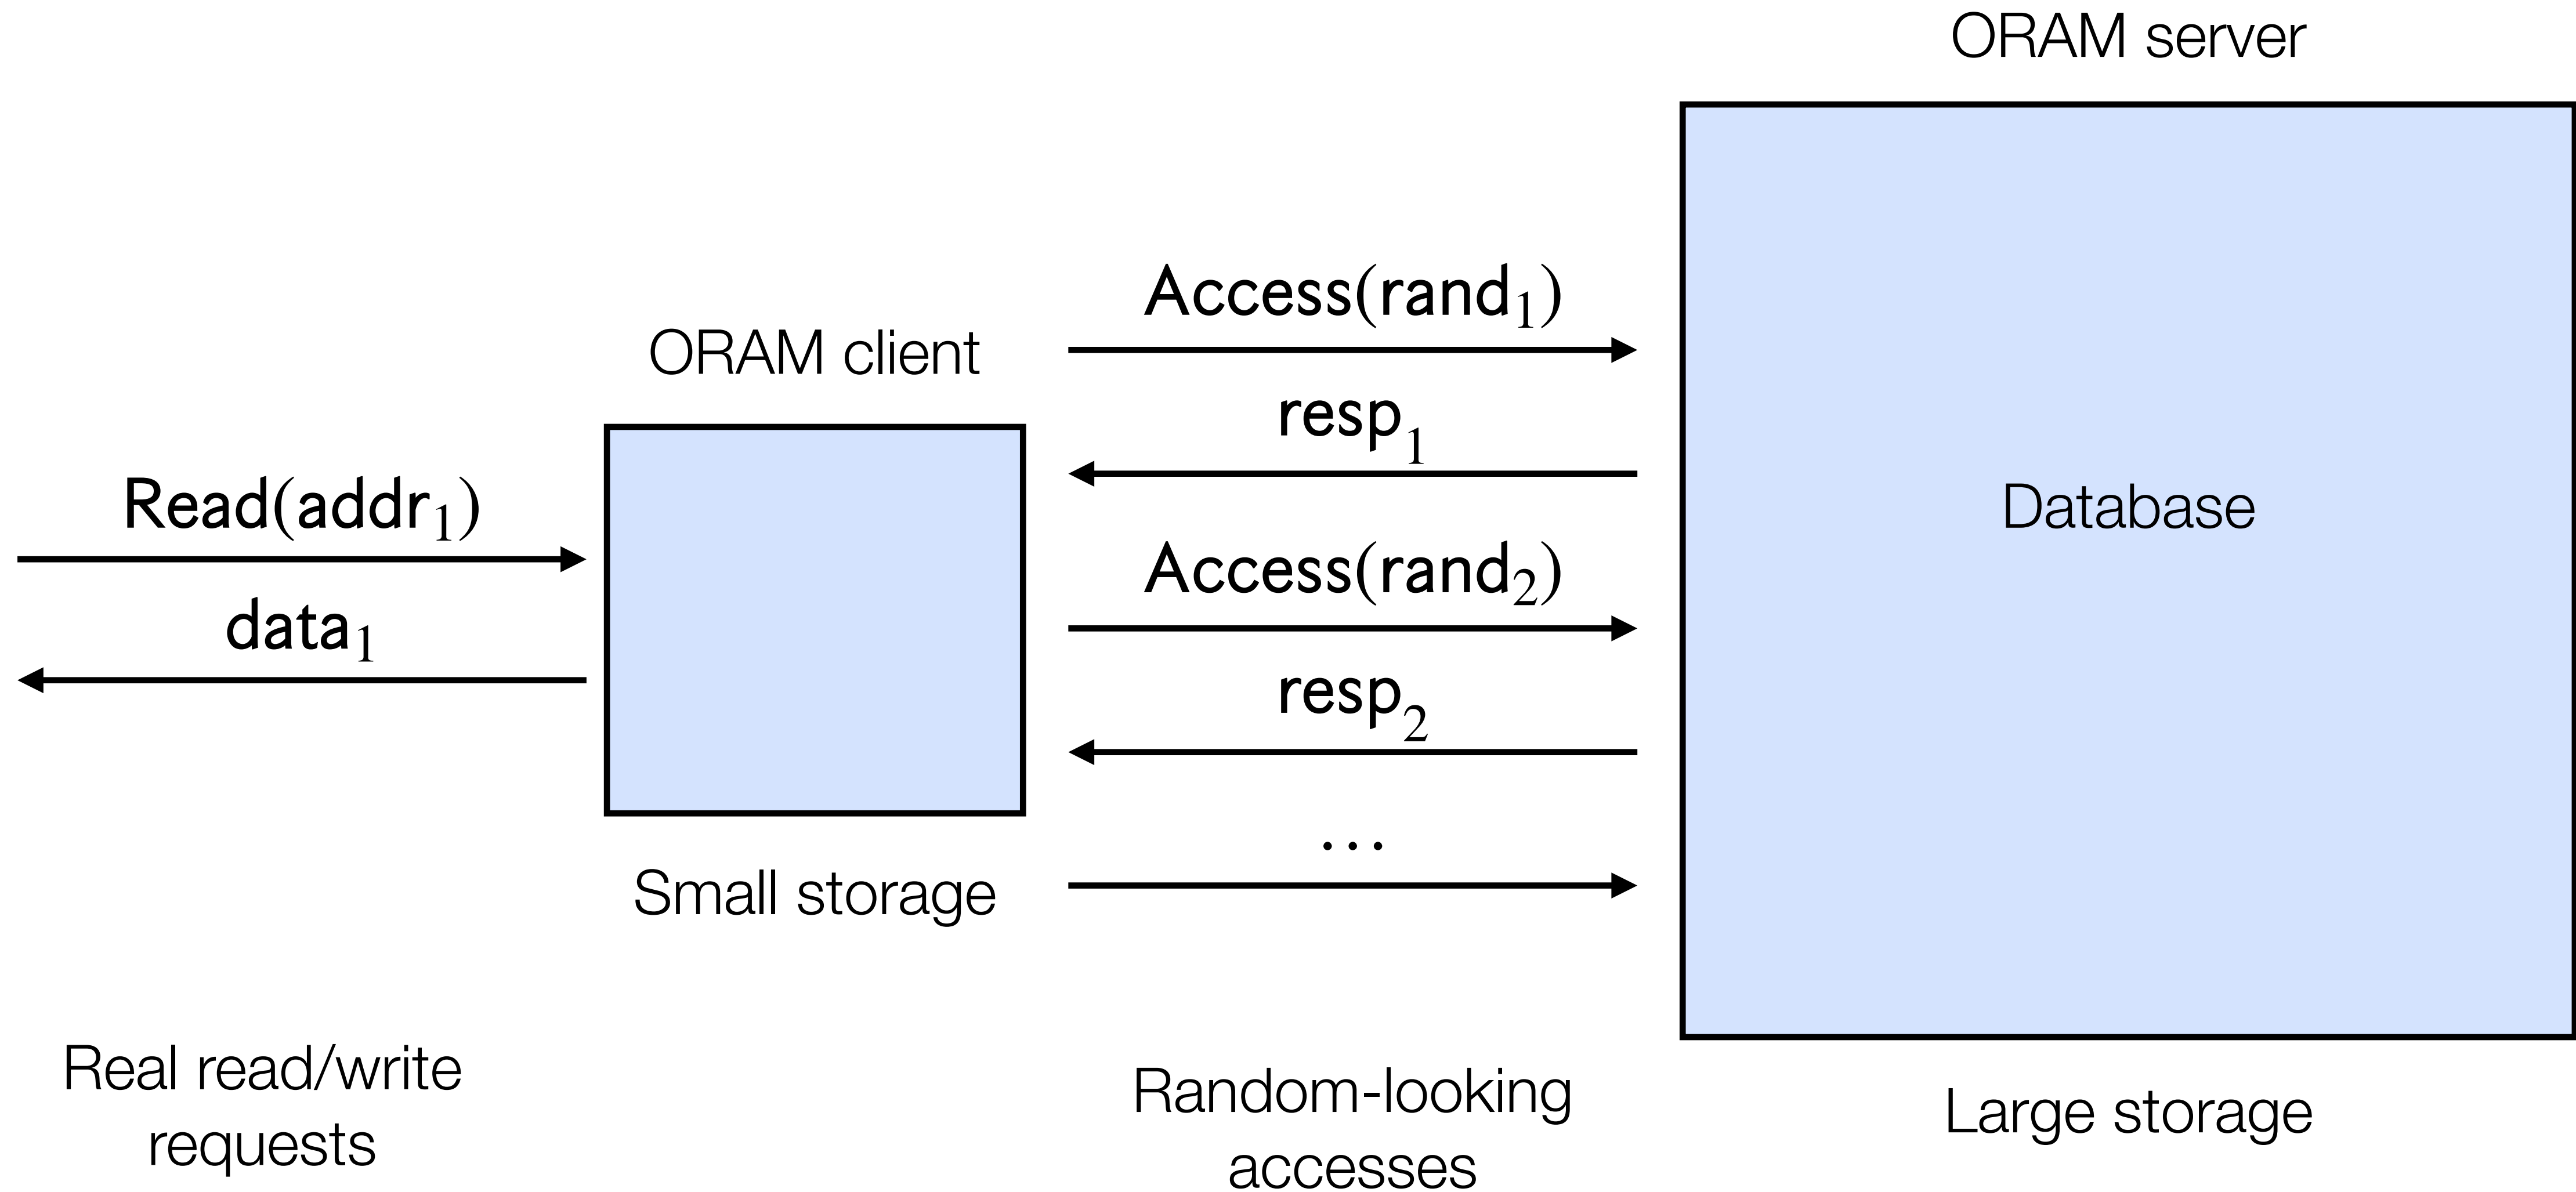
Attacker can infer that Alice and Bob  
have the same medical condition



# Access patterns can reveal sensitive information



# Oblivious RAM

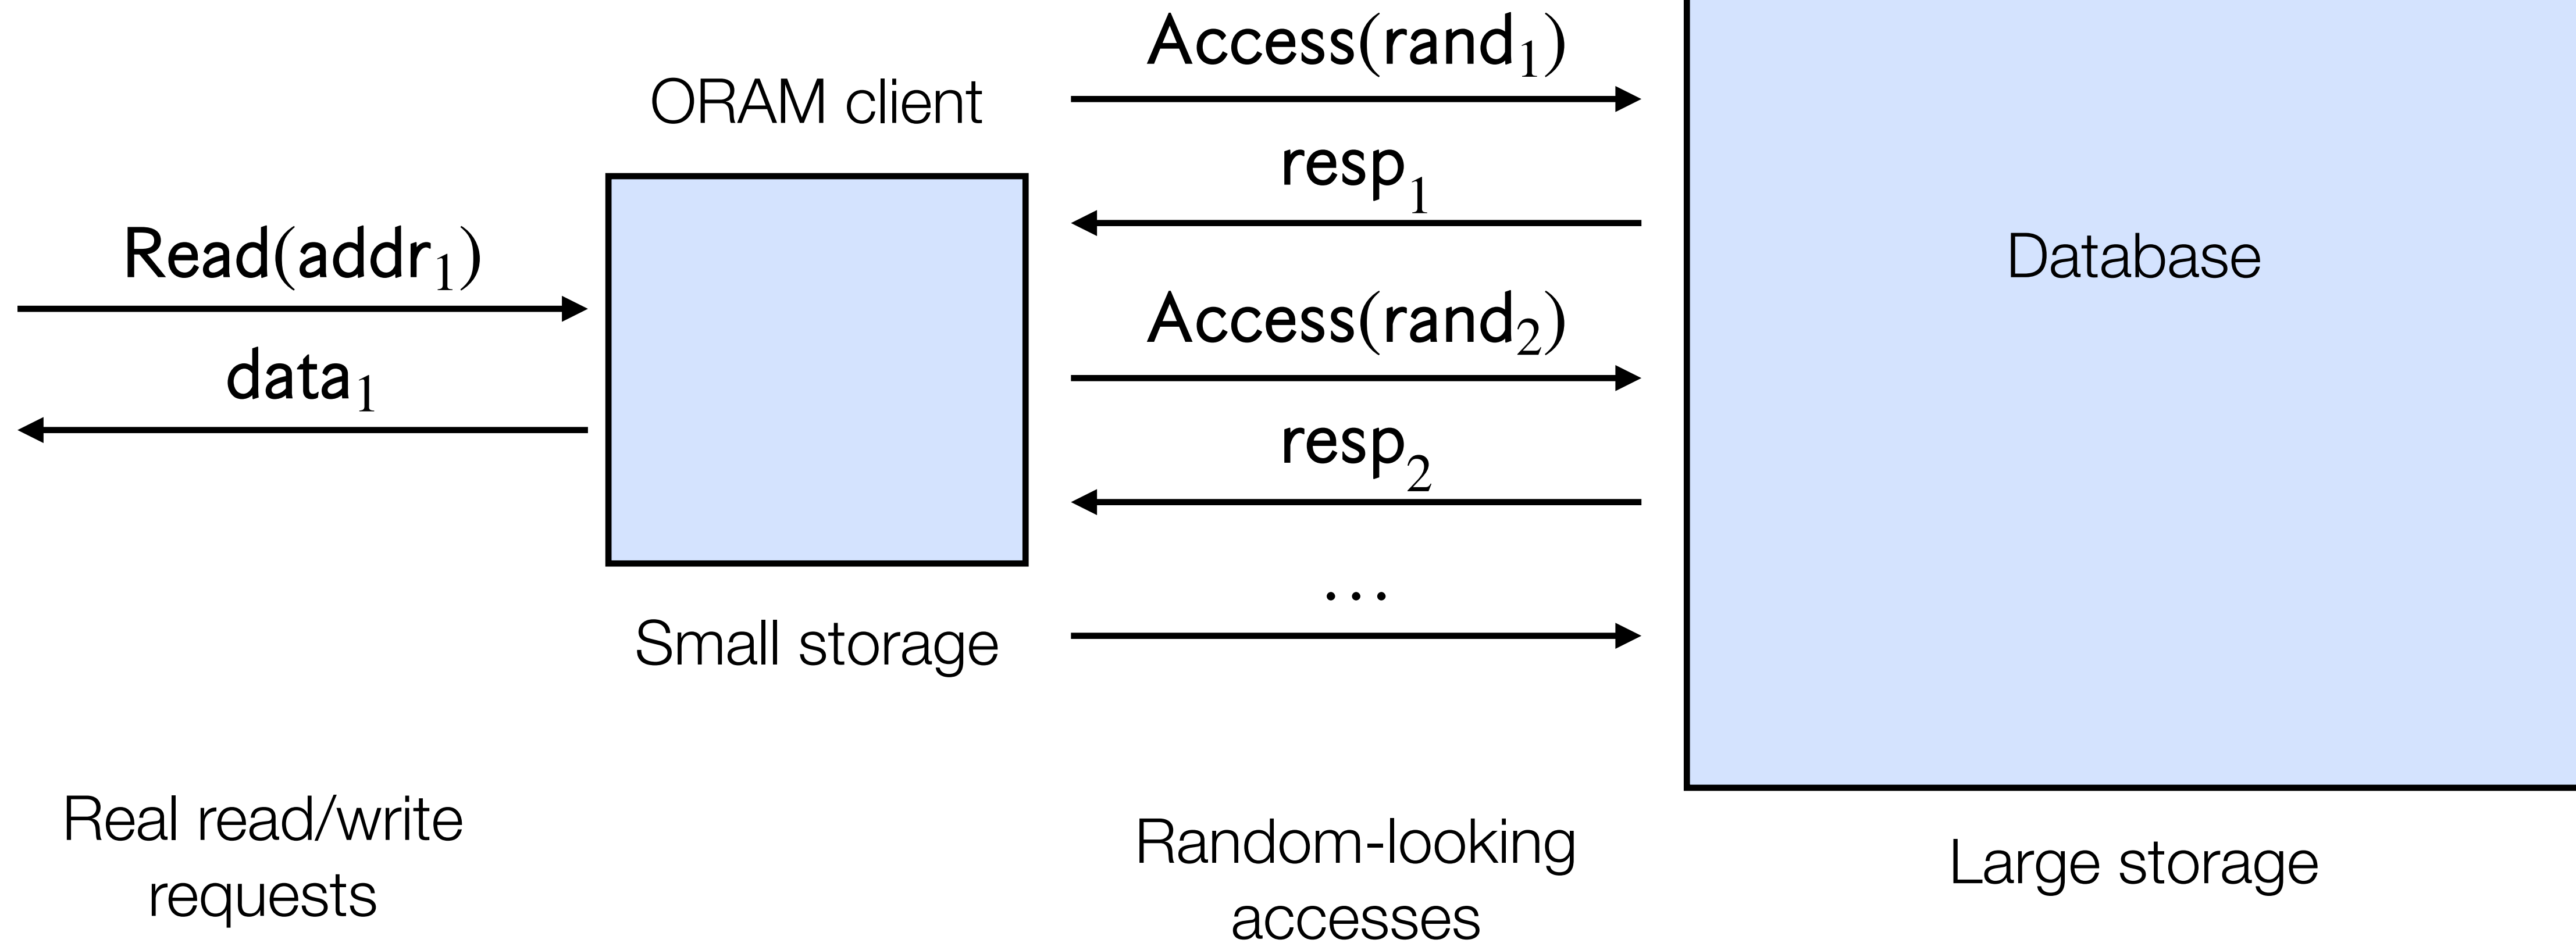


# Oblivious RAM

Attacker that observes memory accesses “learns nothing” about the real read/write requests



ORAM server



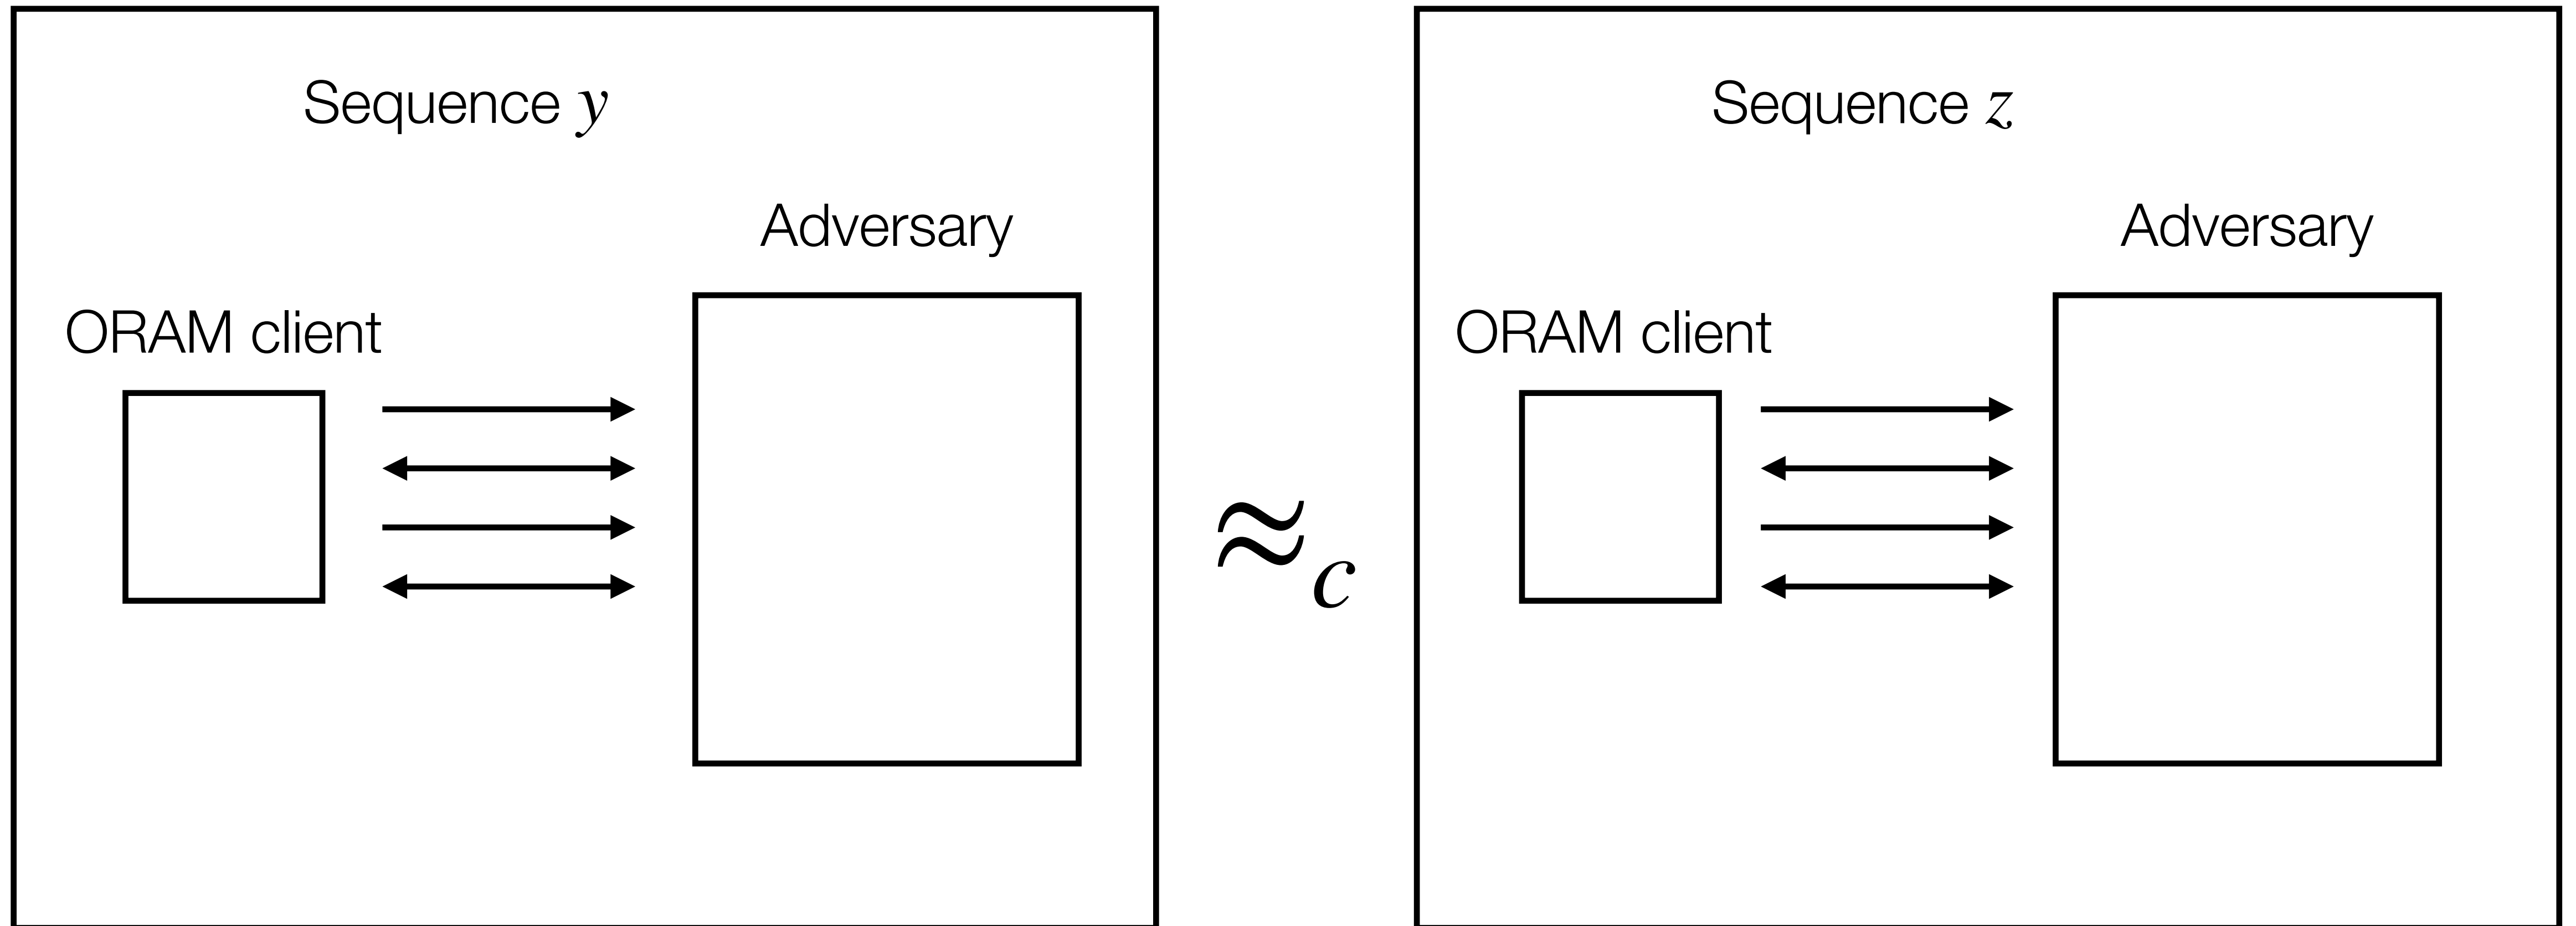
# Definitions

Sequence of operations  $y = ((op_1, addr_1, data_1), (op_2, addr_2, data_2), \dots)$   
where  $op \in \{\text{read}, \text{write}\}$

**Correctness:** An ORAM construction is correct if the responses from the ORAM for a sequence of operations  $y$  matches the responses from a standard RAM (with overwhelming probability).

**Security:** For any two request sequences  $y, z$  of the same length, their access patterns (i.e., interactions between ORAM client and server) are indistinguishable.

# Security definition






# Outline

1. Overview
- 2. Square-root construction**
3. Hierarchical construction
4. Limitations
5. Student presentation: PathORAM

# A very simple, very expensive construction



Why private?

Server touched every element, and so the client could be accessing any element

# Square-root ORAM construction

Goldreich and Ostrovsky

- Can we reduce the costs?
- Take advantage of:
  - Randomness
  - Large, mutable server state
  - Small, mutable client state

# On the way to square-root ORAM construction

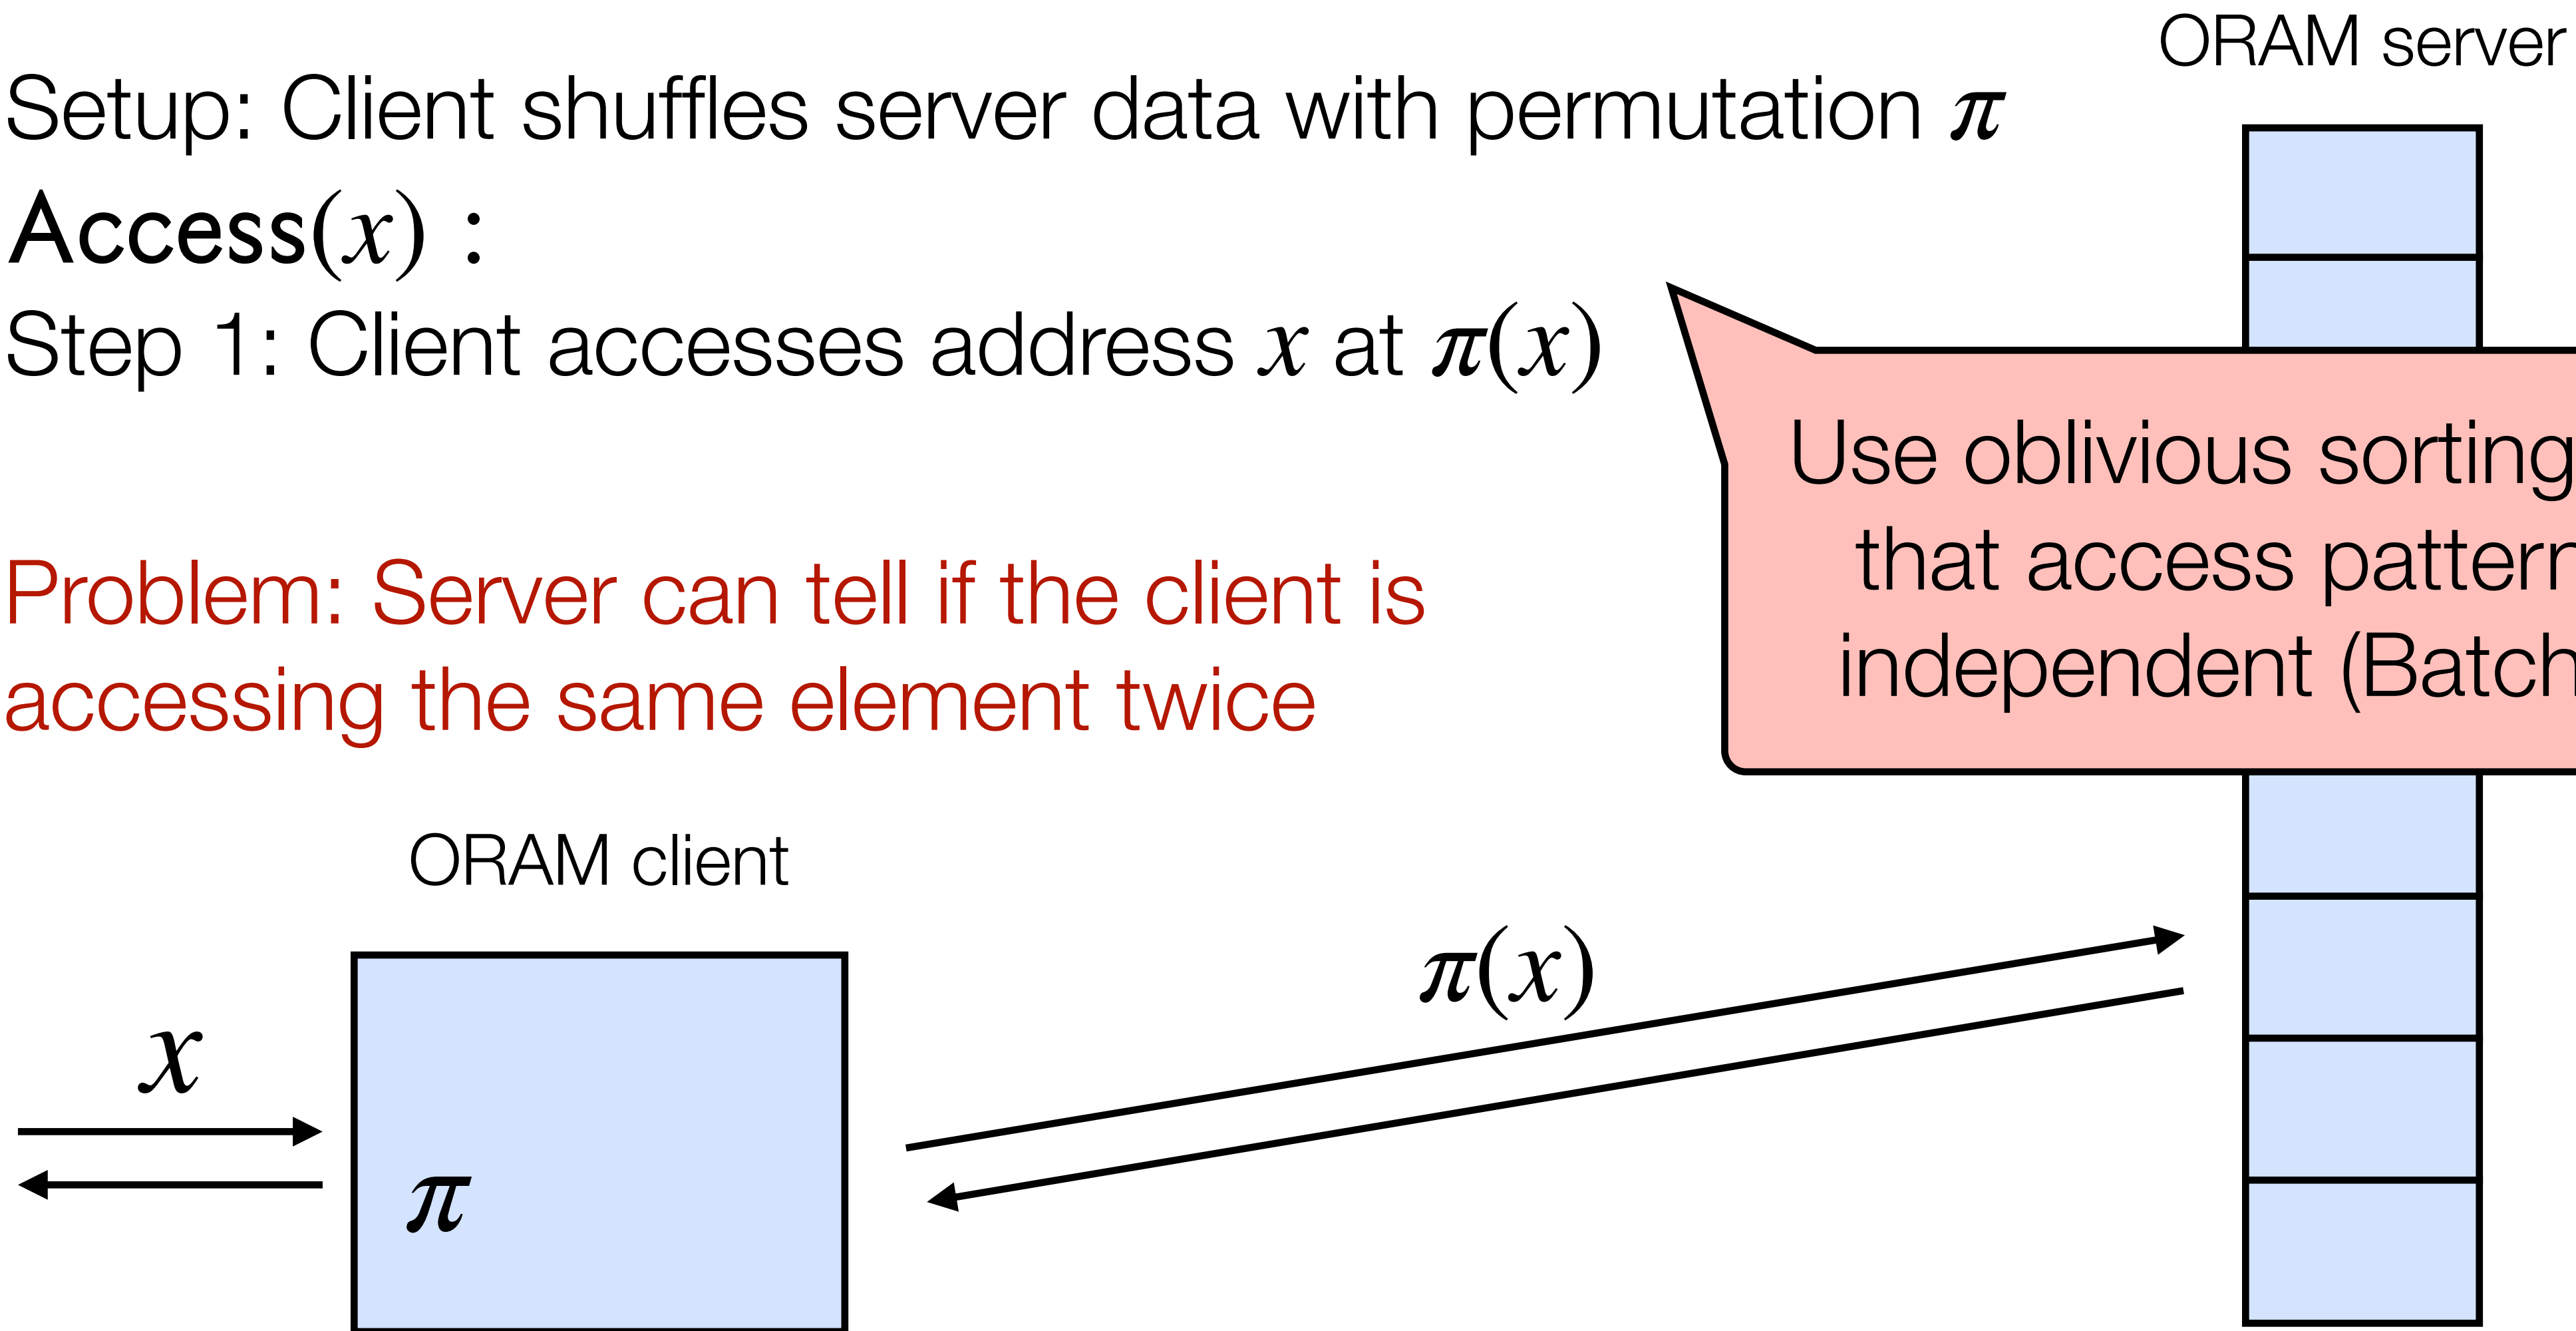
Setup: Client shuffles server data with permutation  $\pi$

**Access( $x$ ) :**

Step 1: Client accesses address  $x$  at  $\pi(x)$

Problem: Server can tell if the client is accessing the same element twice

Use oblivious sorting algorithm so that access patterns are data-independent (Batcher's sorting)



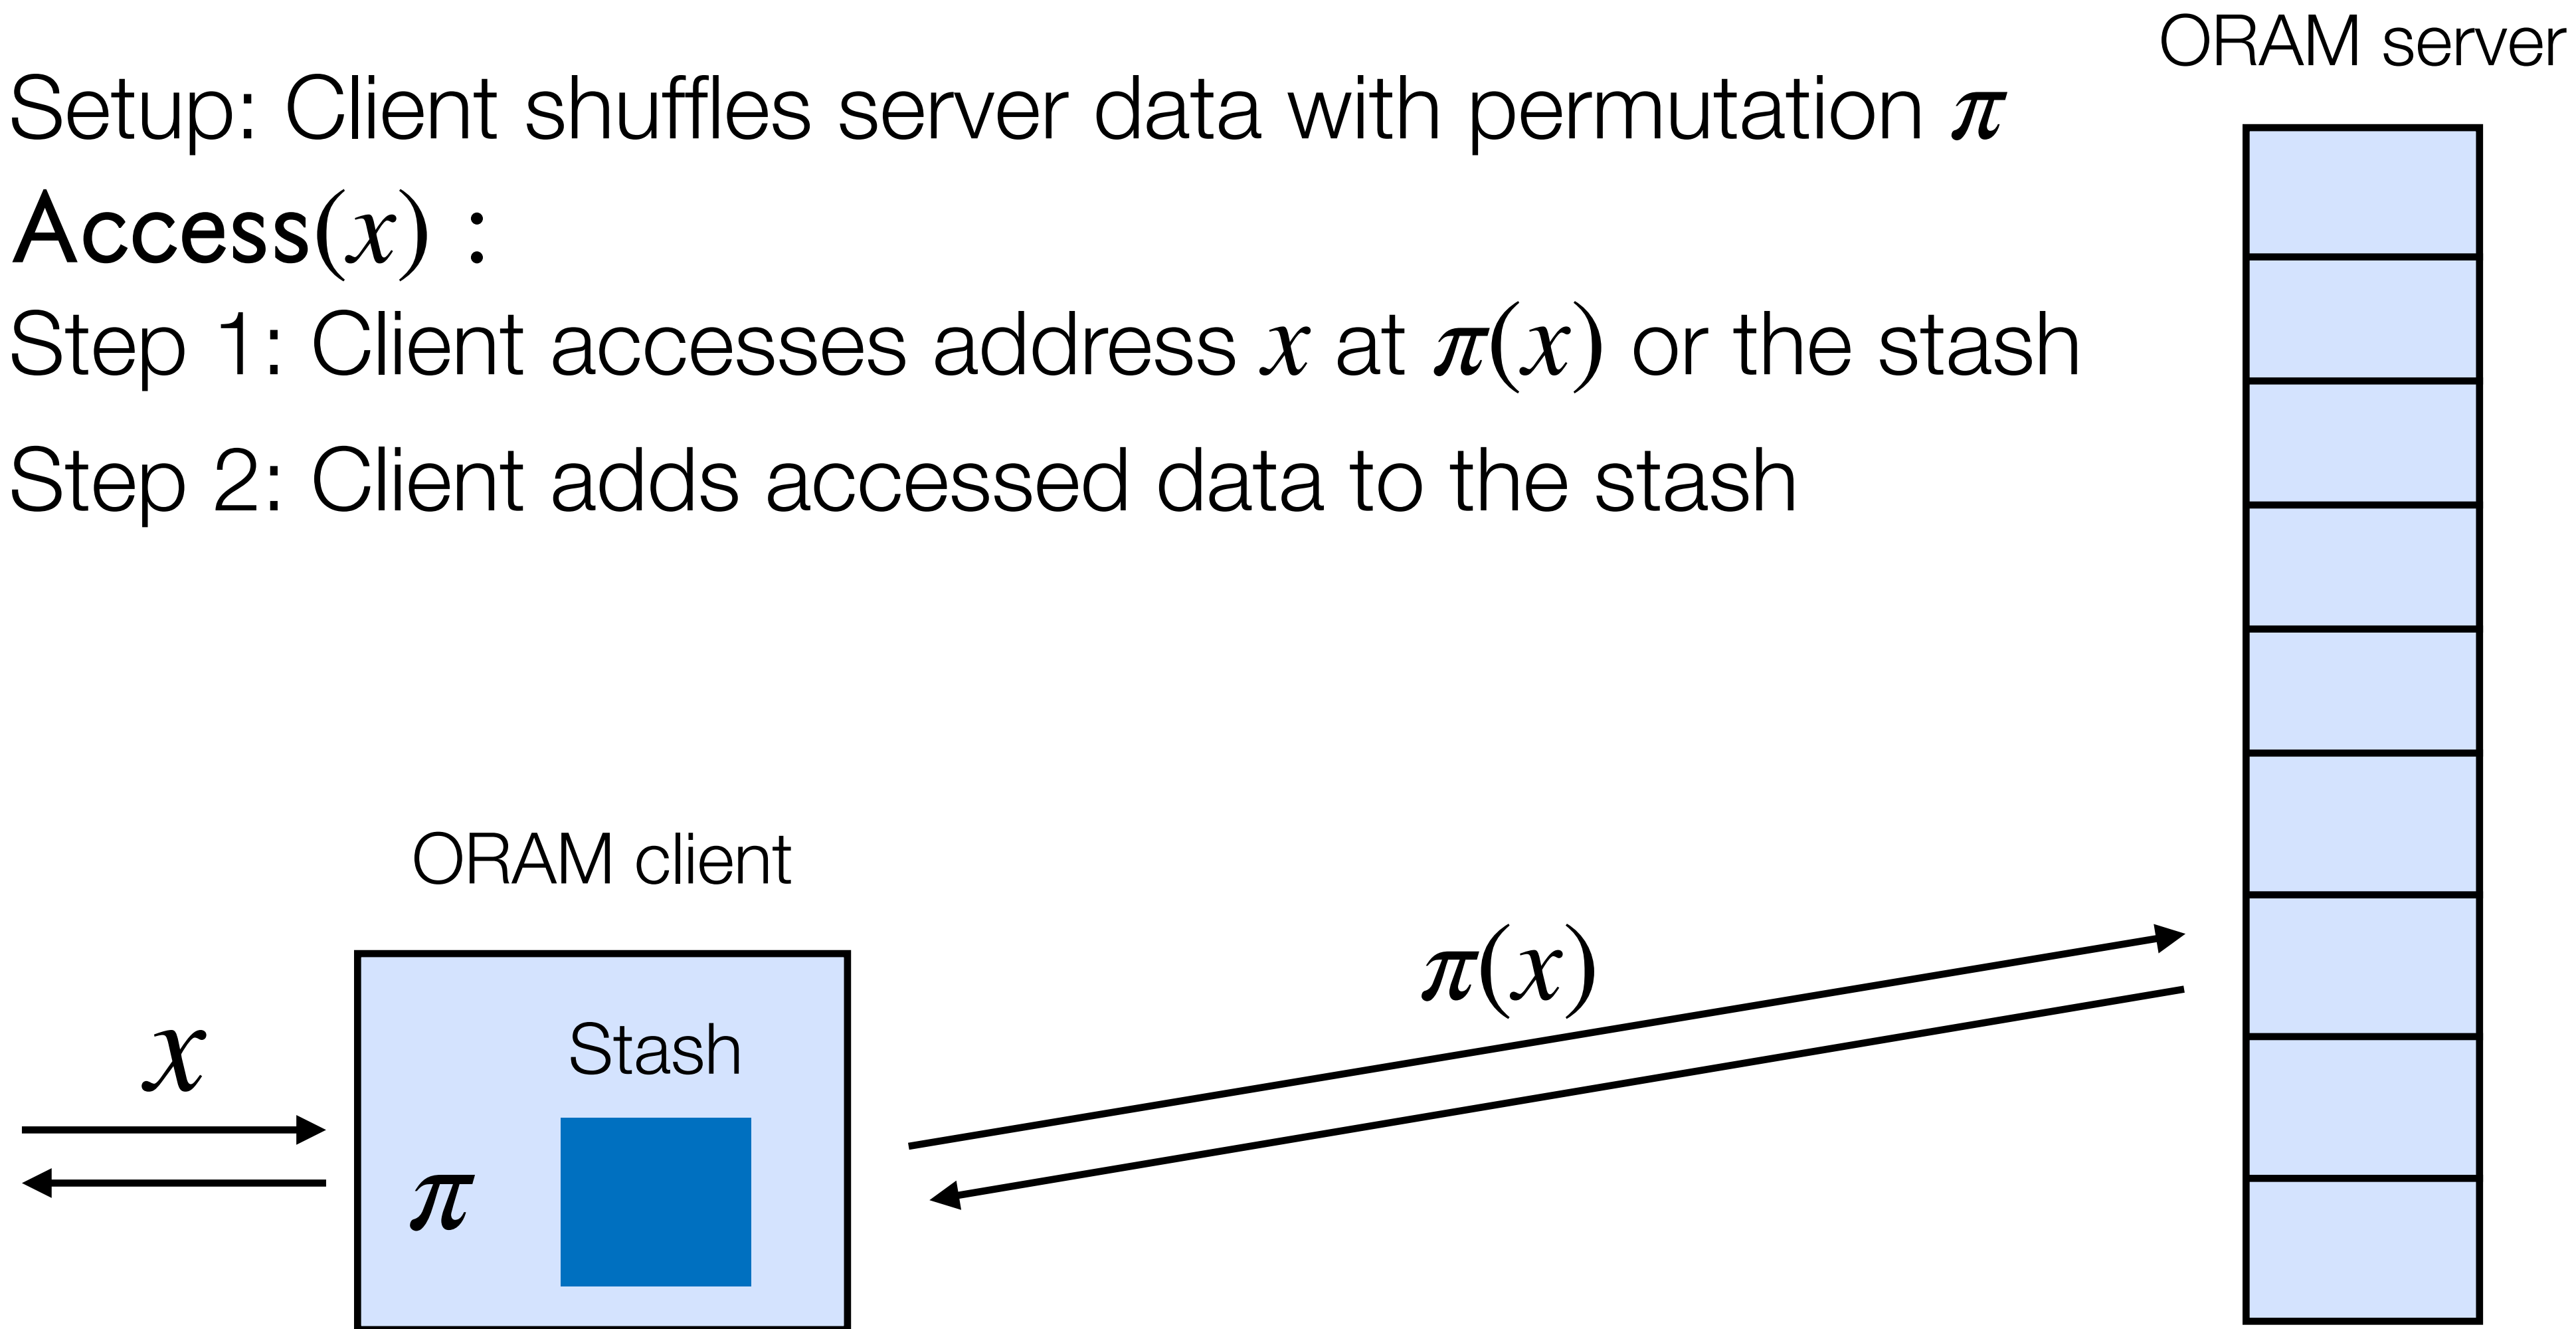
# On the way to square-root ORAM construction

Setup: Client shuffles server data with permutation  $\pi$

**Access( $x$ ) :**

Step 1: Client accesses address  $x$  at  $\pi(x)$  or the stash

Step 2: Client adds accessed data to the stash



Stash stores all elements fetched with  $\pi$

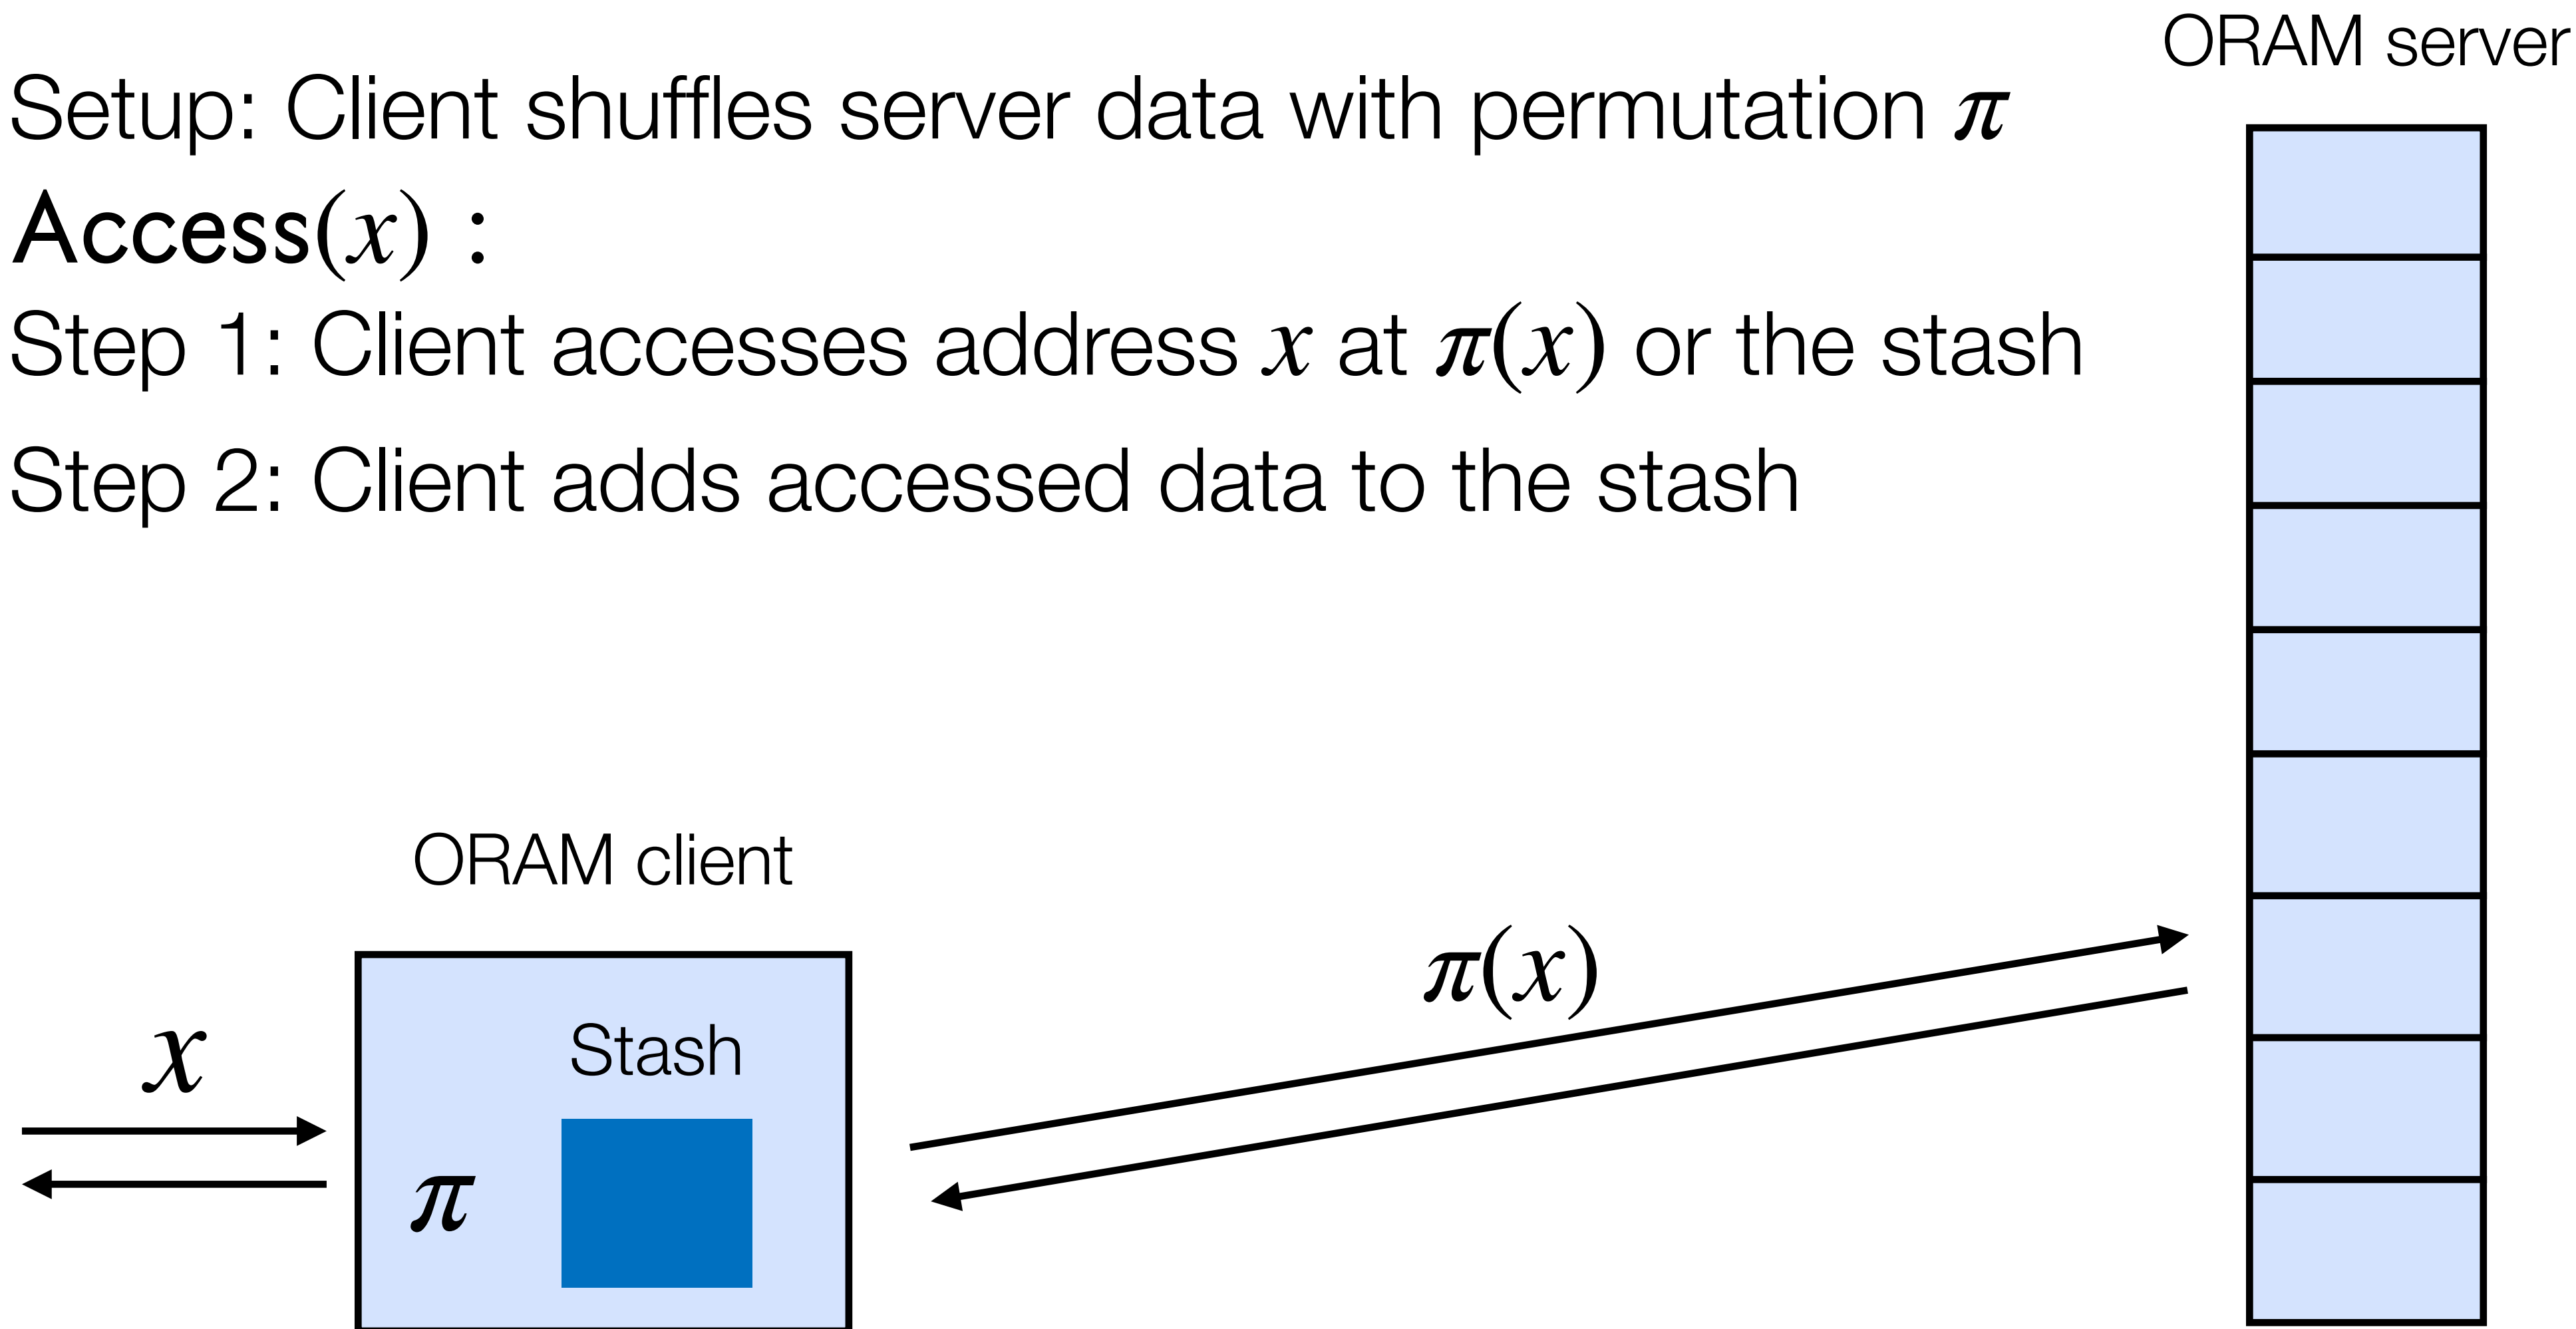
# On the way to square-root ORAM construction

Setup: Client shuffles server data with permutation  $\pi$

**Access( $x$ ) :**

Step 1: Client accesses address  $x$  at  $\pi(x)$  or the stash

Step 2: Client adds accessed data to the stash



Problem: How to keep stash from growing indefinitely?

# On the way to square-root ORAM construction

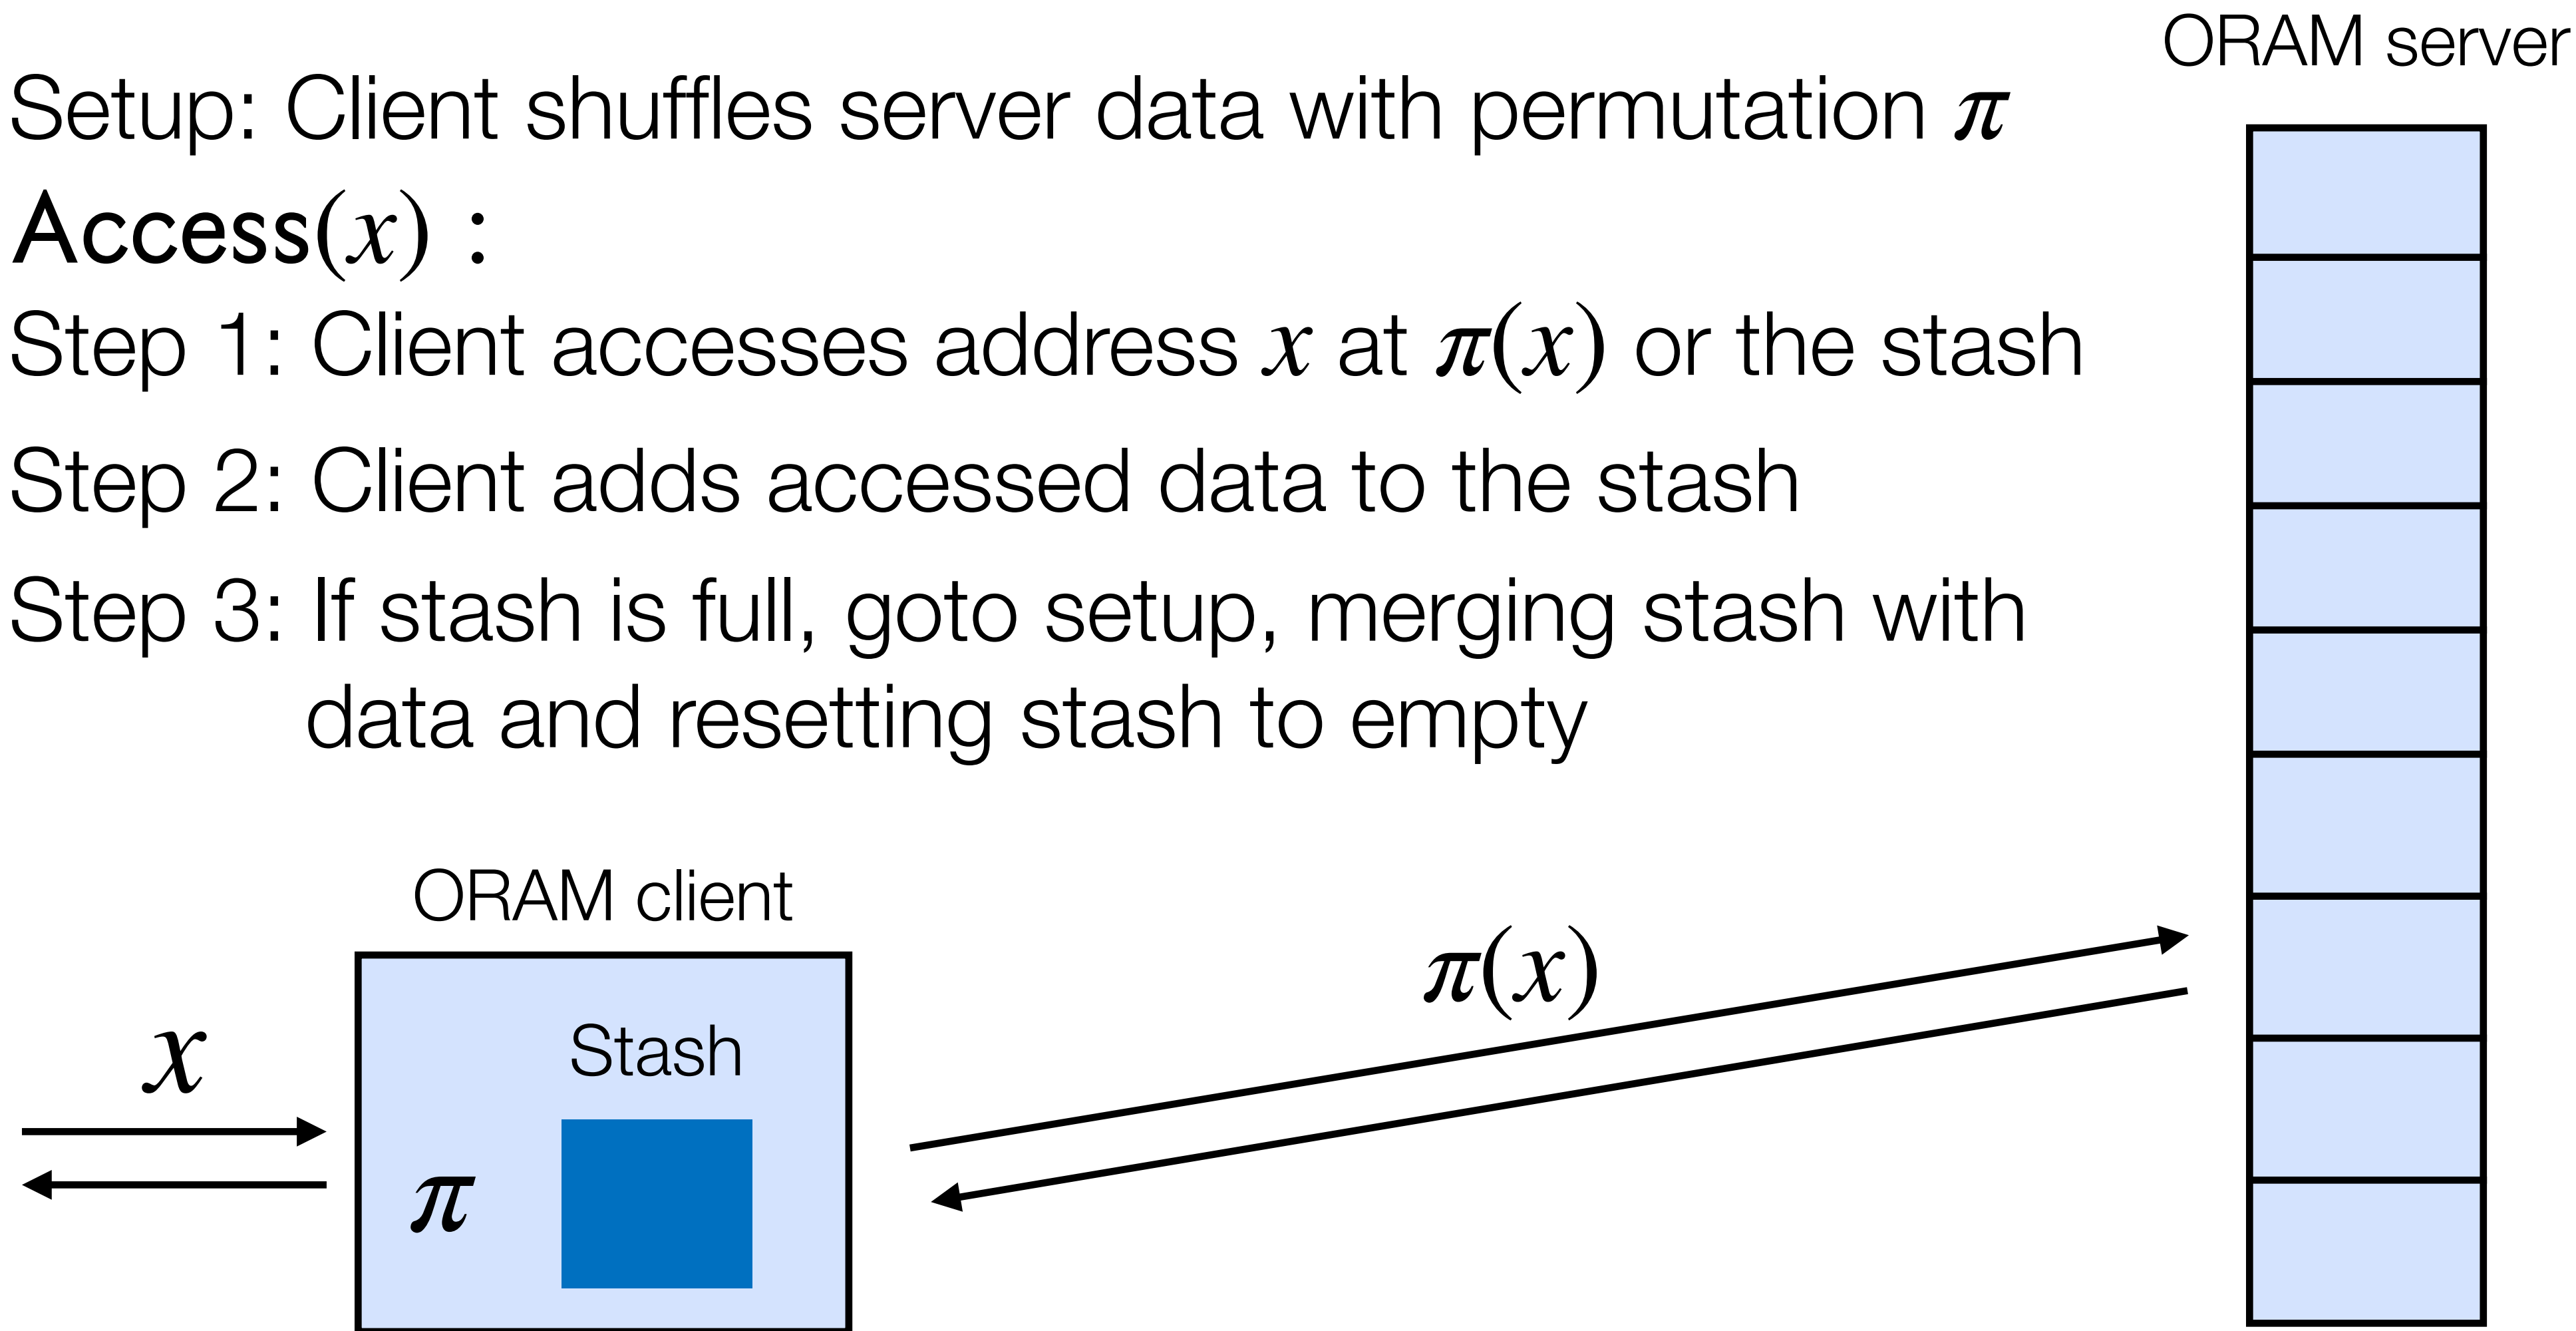
Setup: Client shuffles server data with permutation  $\pi$

**Access( $x$ ) :**

Step 1: Client accesses address  $x$  at  $\pi(x)$  or the stash

Step 2: Client adds accessed data to the stash

Step 3: If stash is full, goto setup, merging stash with data and resetting stash to empty



Problem: How to hide if a request is in the stash?

# On the way to square-root ORAM construction

Setup: Shuffle server data + dummies with permutation  $\pi$


**Access( $x$ ) :**

Step 1: If  $x$  is in the stash, client accesses **dummy**.

Otherwise, client accesses  $\pi(x)$

Step 2: Client adds accessed data to the stash

Step 3: If stash is full, goto setup, merging stash with data and resetting stash to empty





# On the way to square-root ORAM construction

Setup: Shuffle server data + dummies with permutation  $\pi$


**Access( $x$ ) :**

Step 1: If  $x$  is in the stash, client accesses **dummy**.

Otherwise, client accesses  $\pi(x)$

Step 2: Client adds accessed data to the stash

Step 3: If stash is full, goto setup, merging stash with data and resetting stash to empty



# Correctness

Setup: Shuffle server data + dummies with permutation  $\pi$

**Access( $x$ ) :**

Step 1: If  $x$  is in the stash, client accesses **dummy**.

Otherwise, client accesses  $\pi(x)$

Step 2: Client adds accessed data to the stash

Step 3: If stash is full, goto setup, merging stash with data and resetting stash to empty

Correctness from:

- Correctness of shuffle
- Correct maintenance of the stash

# Security

Setup: Shuffle server data + dummies with permutation  $\pi$

**Access( $x$ ) :**

Step 1: If  $x$  is in the stash, client accesses **dummy**.

Otherwise, client accesses  $\pi(x)$

Step 2: Client adds accessed data to the stash

Step 3: If stash is full, goto setup, merging stash with data and resetting stash to empty

Security from:

- Permutation appears random to the server
- Each request accesses a unique element
- Server cannot distinguish a real request from a dummy request

# Parameterizing square-root ORAM

Length  $n$  array, stash size  $s$

Add  $s$  dummies

Setup: Shuffle server data + dummies with permutation  $\pi$

$O(n \cdot \log^2 n)$   
comparisons

**Access( $x$ ) :**

Step 1: If  $x$  is in the stash, client accesses **dummy**.

Otherwise, client accesses  $\pi(x)$

Check  $s$  elements

Step 2: Client adds accessed data to the stash

Step 3: If stash is full, goto setup, merging stash with data and resetting stash to empty

By setting stash size  $s = \sqrt{n}$ , amortized overhead is  $O(\sqrt{n} \cdot \log^2 n)$

Client storage is  $O(\sqrt{n})$

# Outline

1. Overview
2. Square-root construction
- 3. Hierarchical construction**
4. Limitations
5. Student presentation: PathORAM


# Hierarchical ORAM

[Goldreich, Ostrovsky]

High-level idea: Hierarchy of different-sized buffers to achieve logarithmic overhead

Square-root ORAM

$$n = 2^k$$




Hierarchical ORAM

$$n = 2^k$$

$$2^{k-1}$$

$$2^{k-2}$$


$$2^{k-3}$$



Shuffle smaller buffers more frequently


# Hierarchical ORAM

Access requires scanning over the stash and making a lookup in each buffer



# Hierarchical ORAM

Shuffle smaller buffers more frequently




When stash is full:

- Let  $p$  be the index of the smallest unfilled buffer
- Shuffle together the stash and buffers  $i < p$
- Put the results in buffer  $p$



# Hierarchical ORAM

Shuffle smaller buffers more frequently




When stash is full:

- Let  $p$  be the index of the smallest unfilled buffer
- Shuffle together the stash and buffers  $i < p$
- Put the results in buffer  $p$

# Hierarchical ORAM

Shuffle smaller buffers more frequently




When stash is full:

- Let  $p$  be the index of the smallest unfilled buffer
- Shuffle together the stash and buffers  $i < p$
- Put the results in buffer  $p$

# Hierarchical ORAM

Shuffle smaller buffers more frequently



When stash is full:

- Let  $p$  be the index of the smallest unfilled buffer
- Shuffle together the stash and buffers  $i < p$
- Put the results in buffer  $p$

# Outline

1. Overview
2. Square-root construction
3. Hierarchical construction
- 4. Limitations**
5. Student presentation: PathORAM

# Limitations of ORAM

- All elements must be the same length
- Increased cost compared to plaintext RAM accesses: lower bound of  $O(\log n)$  accesses per operation [Larsen, Nielsen]
- ORAM is designed for a single client; does not directly support multiple clients
- Accesses where shuffling is required take longer than accesses without shuffling (not the case for tree-based ORAMs)
- Elements must be fetched in sequence, not in parallel (addressed in subsequent work, e.g., TaoStore)
- Only supports key-value lookups, but applications need other types of queries

# Outline

1. Overview
2. Square-root construction
3. Hierarchical construction
4. Limitations
- 5. Student presentation: PathORAM**

# References

Goldreich, Oded, and Rafail Ostrovsky. "Software protection and simulation on oblivious RAMs." *Journal of the ACM (JACM)* 43.3 (1996): 431-473.

Larsen, Kasper Green, and Jesper Buus Nielsen. "Yes, there is an oblivious RAM lower bound!." *Annual International Cryptology Conference*. Cham: Springer International Publishing, 2018.

Sahin, Cetin, et al. "Taostore: Overcoming asynchronicity in oblivious data storage." *2016 IEEE Symposium on Security and Privacy (SP)*. IEEE, 2016.

Stefanov, Emil, et al. "Path ORAM: an extremely simple oblivious RAM protocol." *Journal of the ACM (JACM)* 65.4 (2018): 1-26

<https://6893.csail.mit.edu/lec7.pdf>