


# CS 350S: Privacy-Preserving Systems

## Private Information Retrieval II

# Recap: Private information retrieval


Goal: Attacker that controls the server learns nothing about a user's query



|               |          |
|---------------|----------|
| Heart disease | Record 1 |
| Lung cancer   | Record 2 |
| Diabetes      | Record 3 |
| ...           | ...      |

Can also use to fetch articles, images, podcasts, movies, etc. from a remote server

# Trivial solution: download the whole database

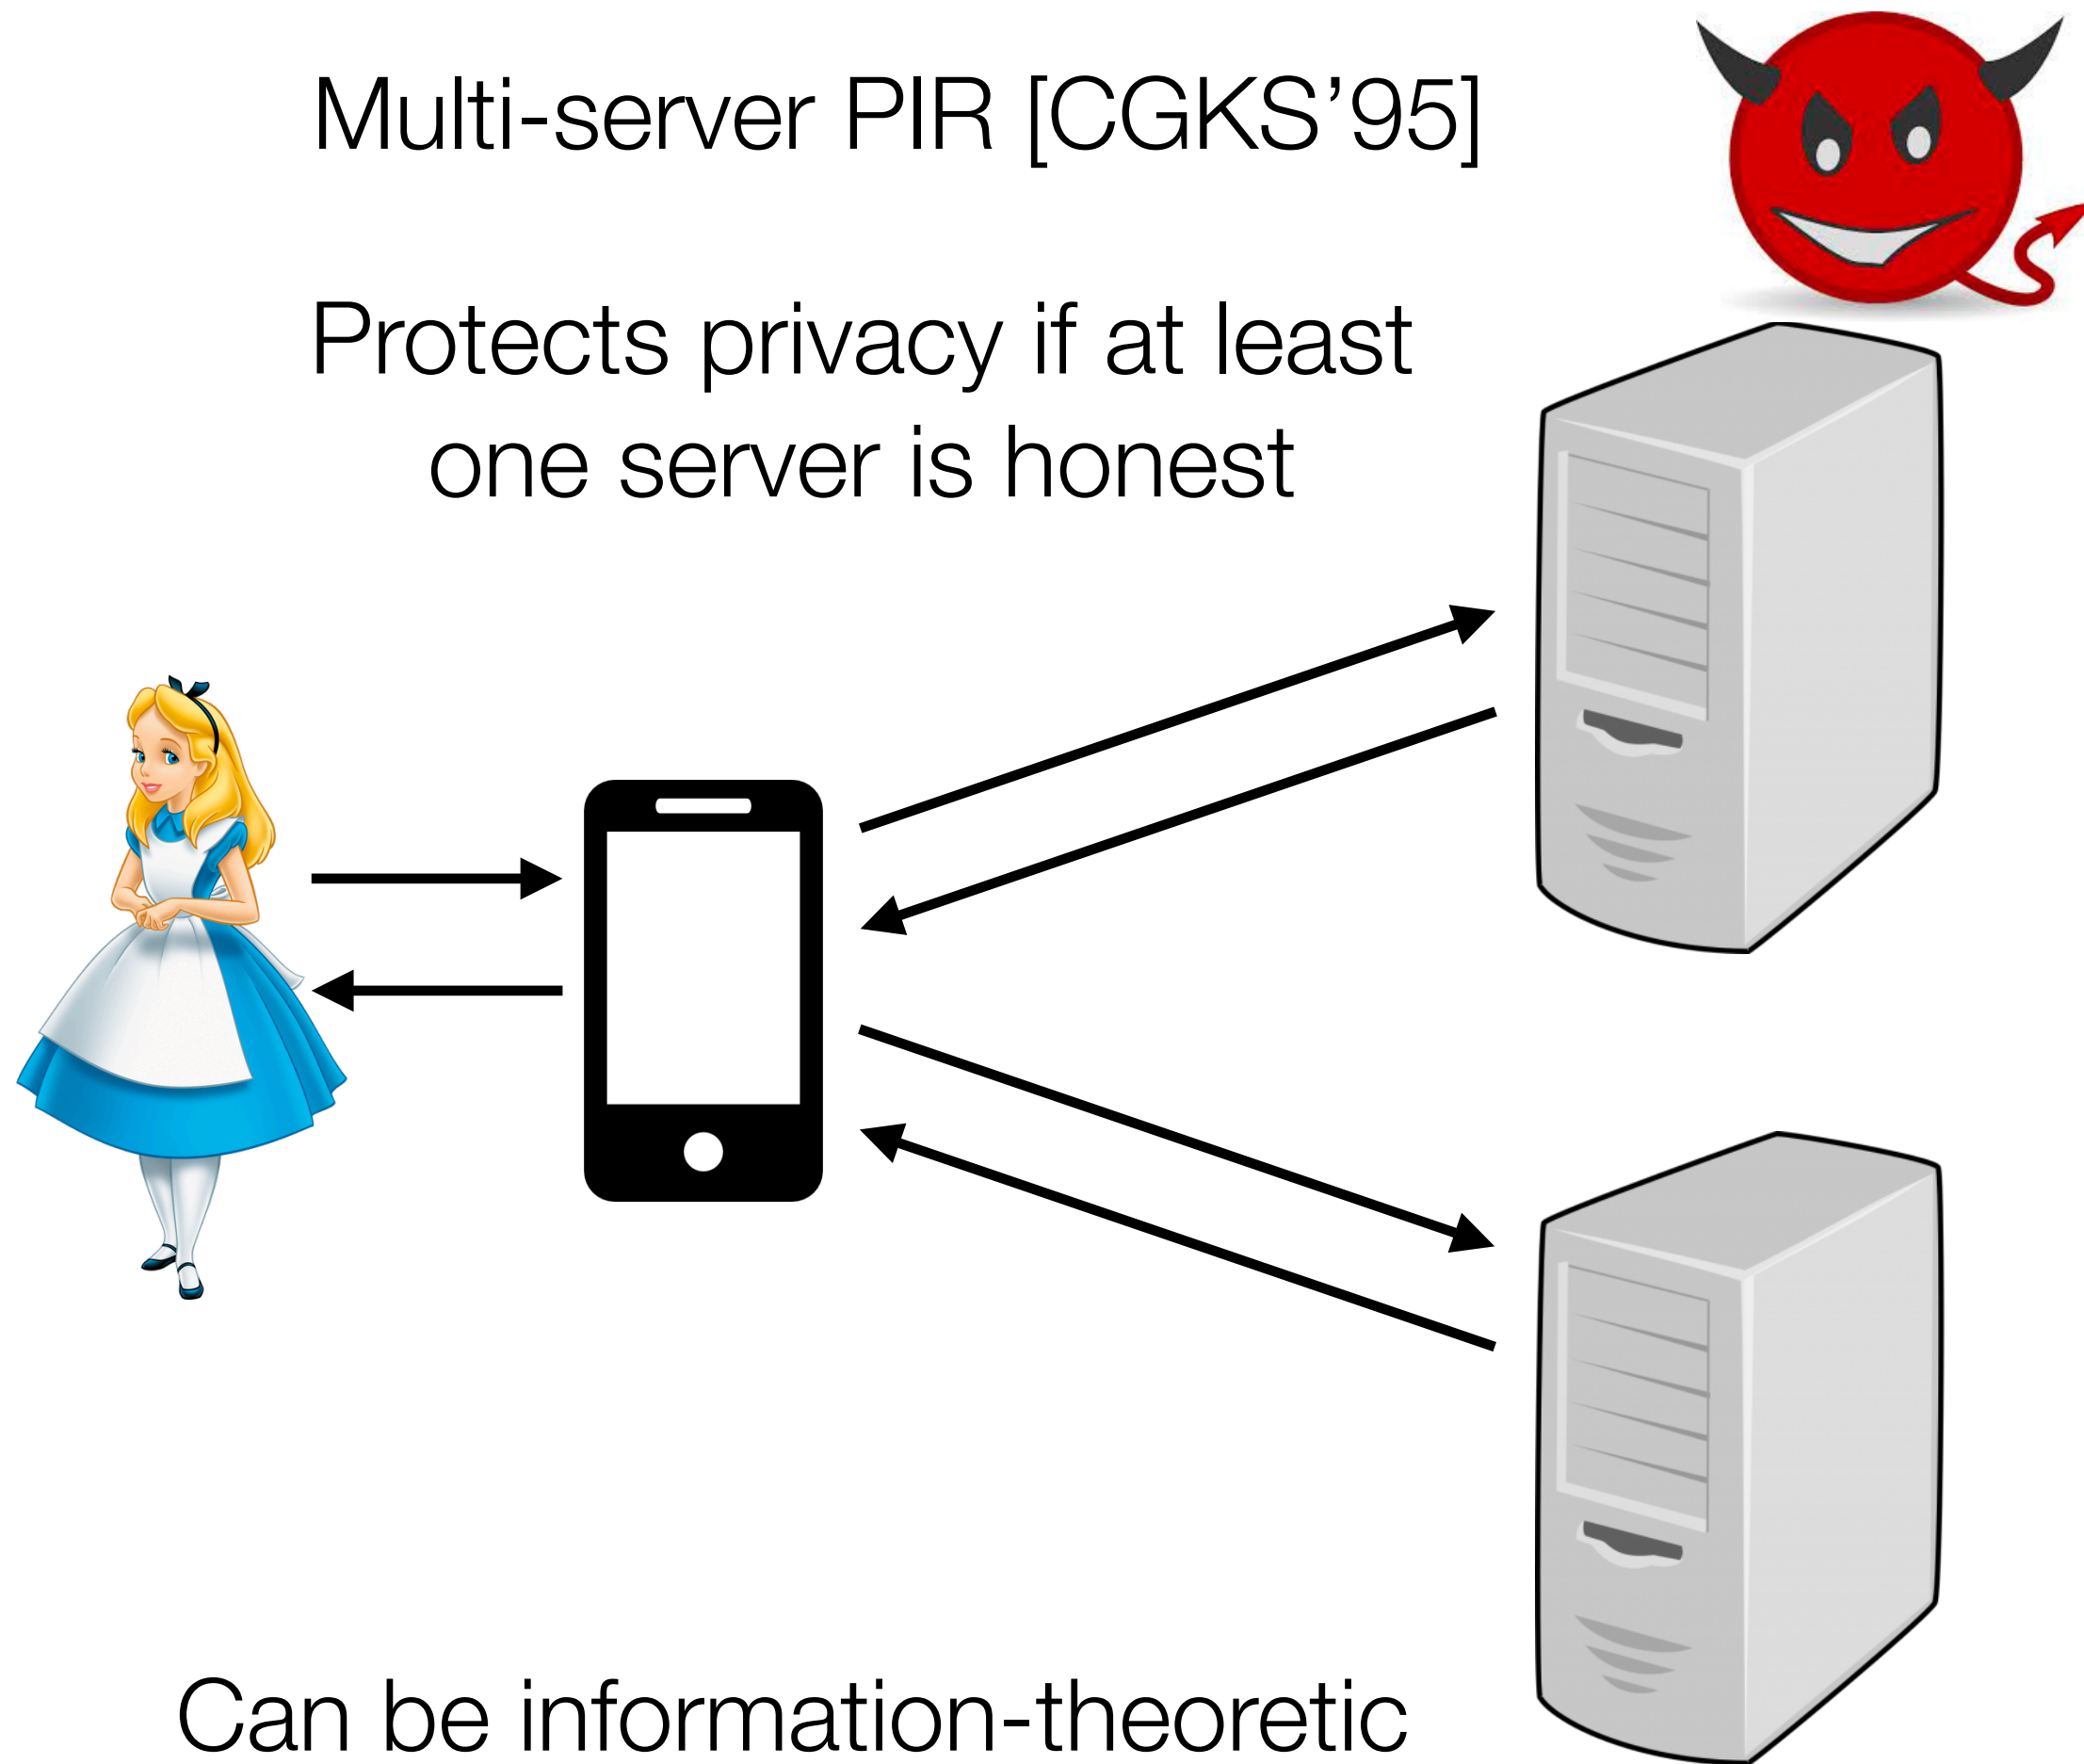


Want to privately query a database with communication *sublinear* in database size

# Two models

Multi-server PIR [CGKS'95]

Protects privacy if at least one server is honest

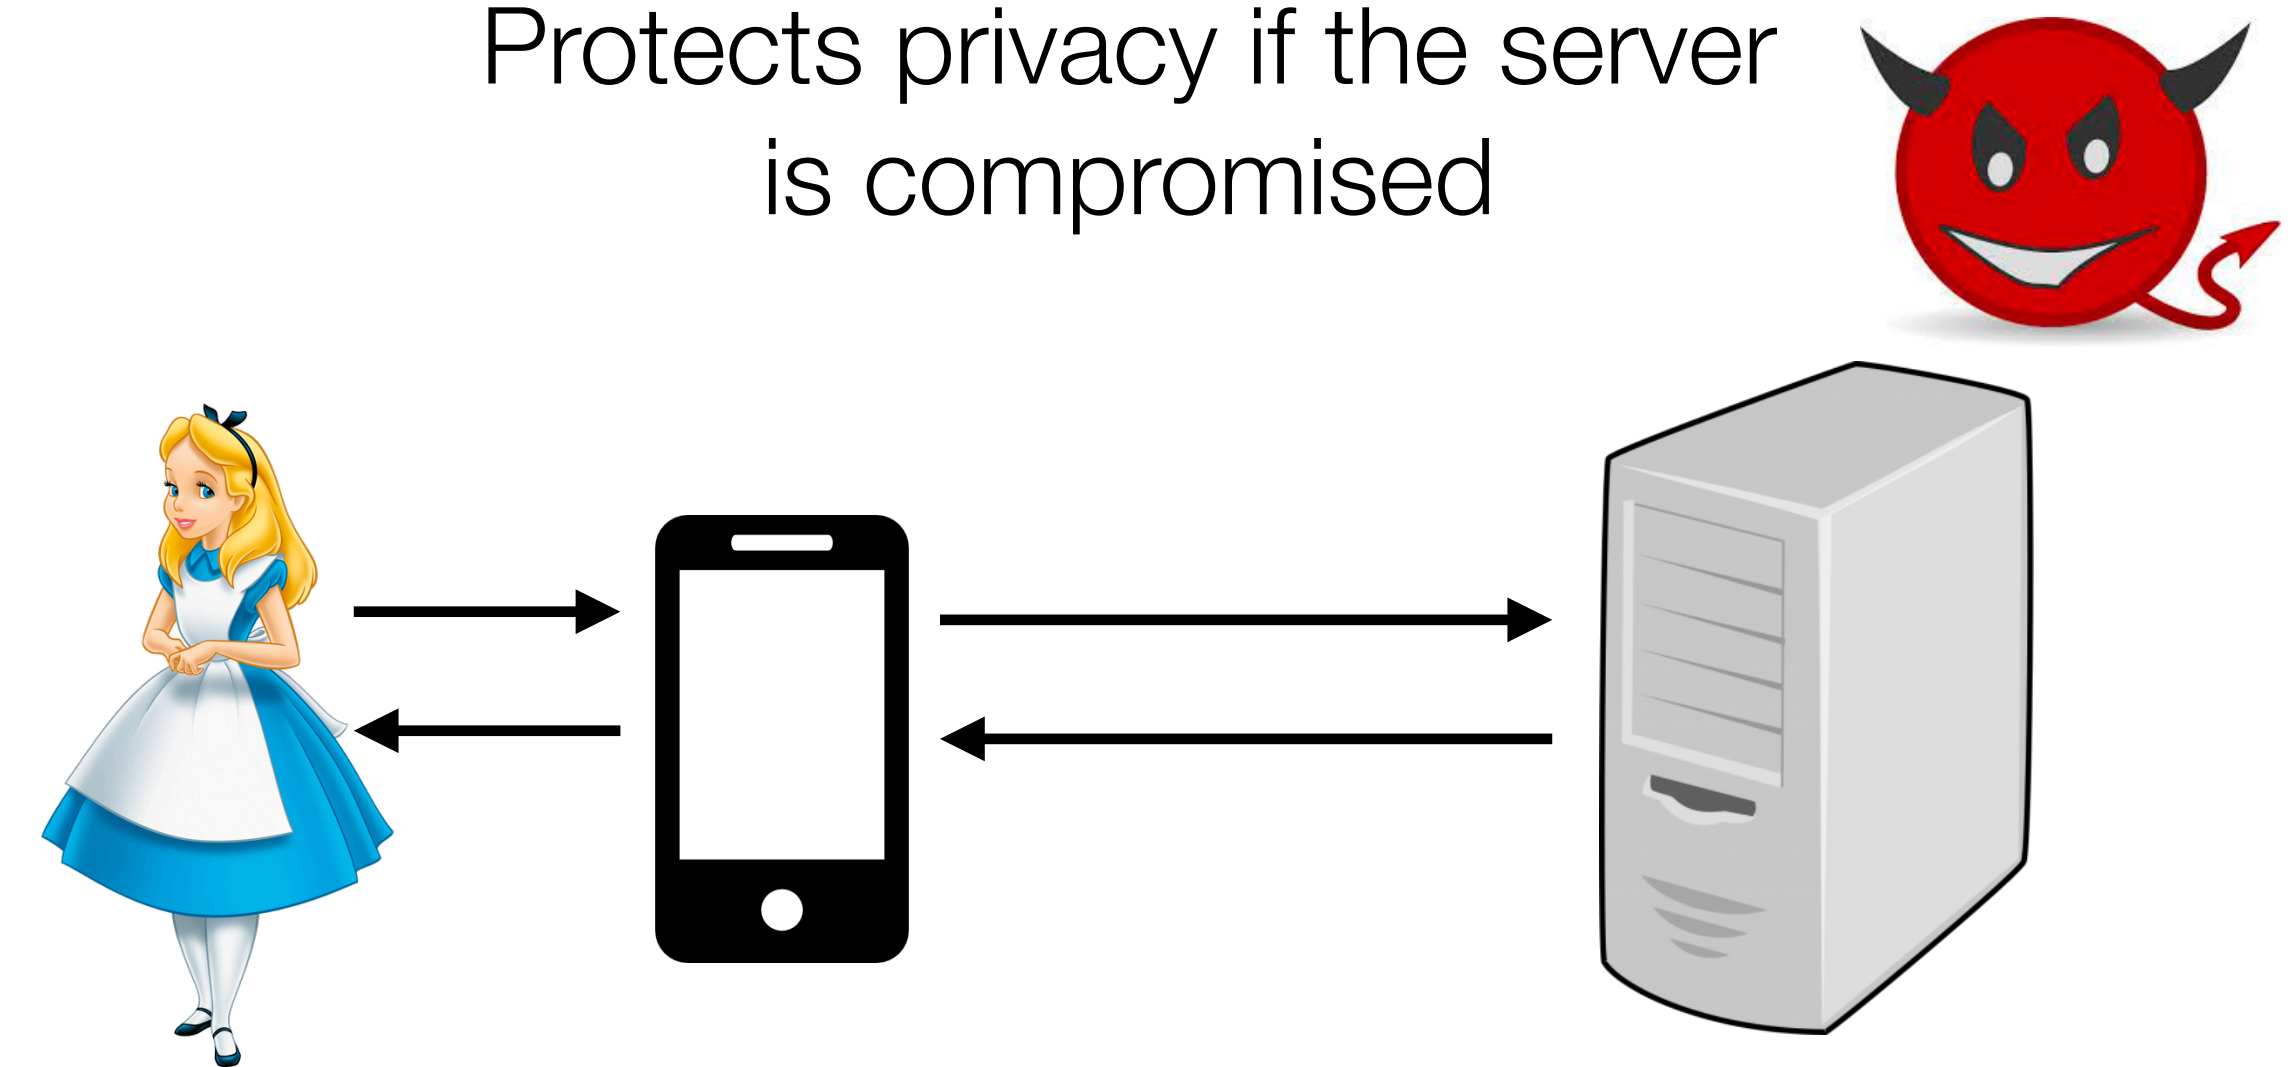


Can be information-theoretic

**TODAY**

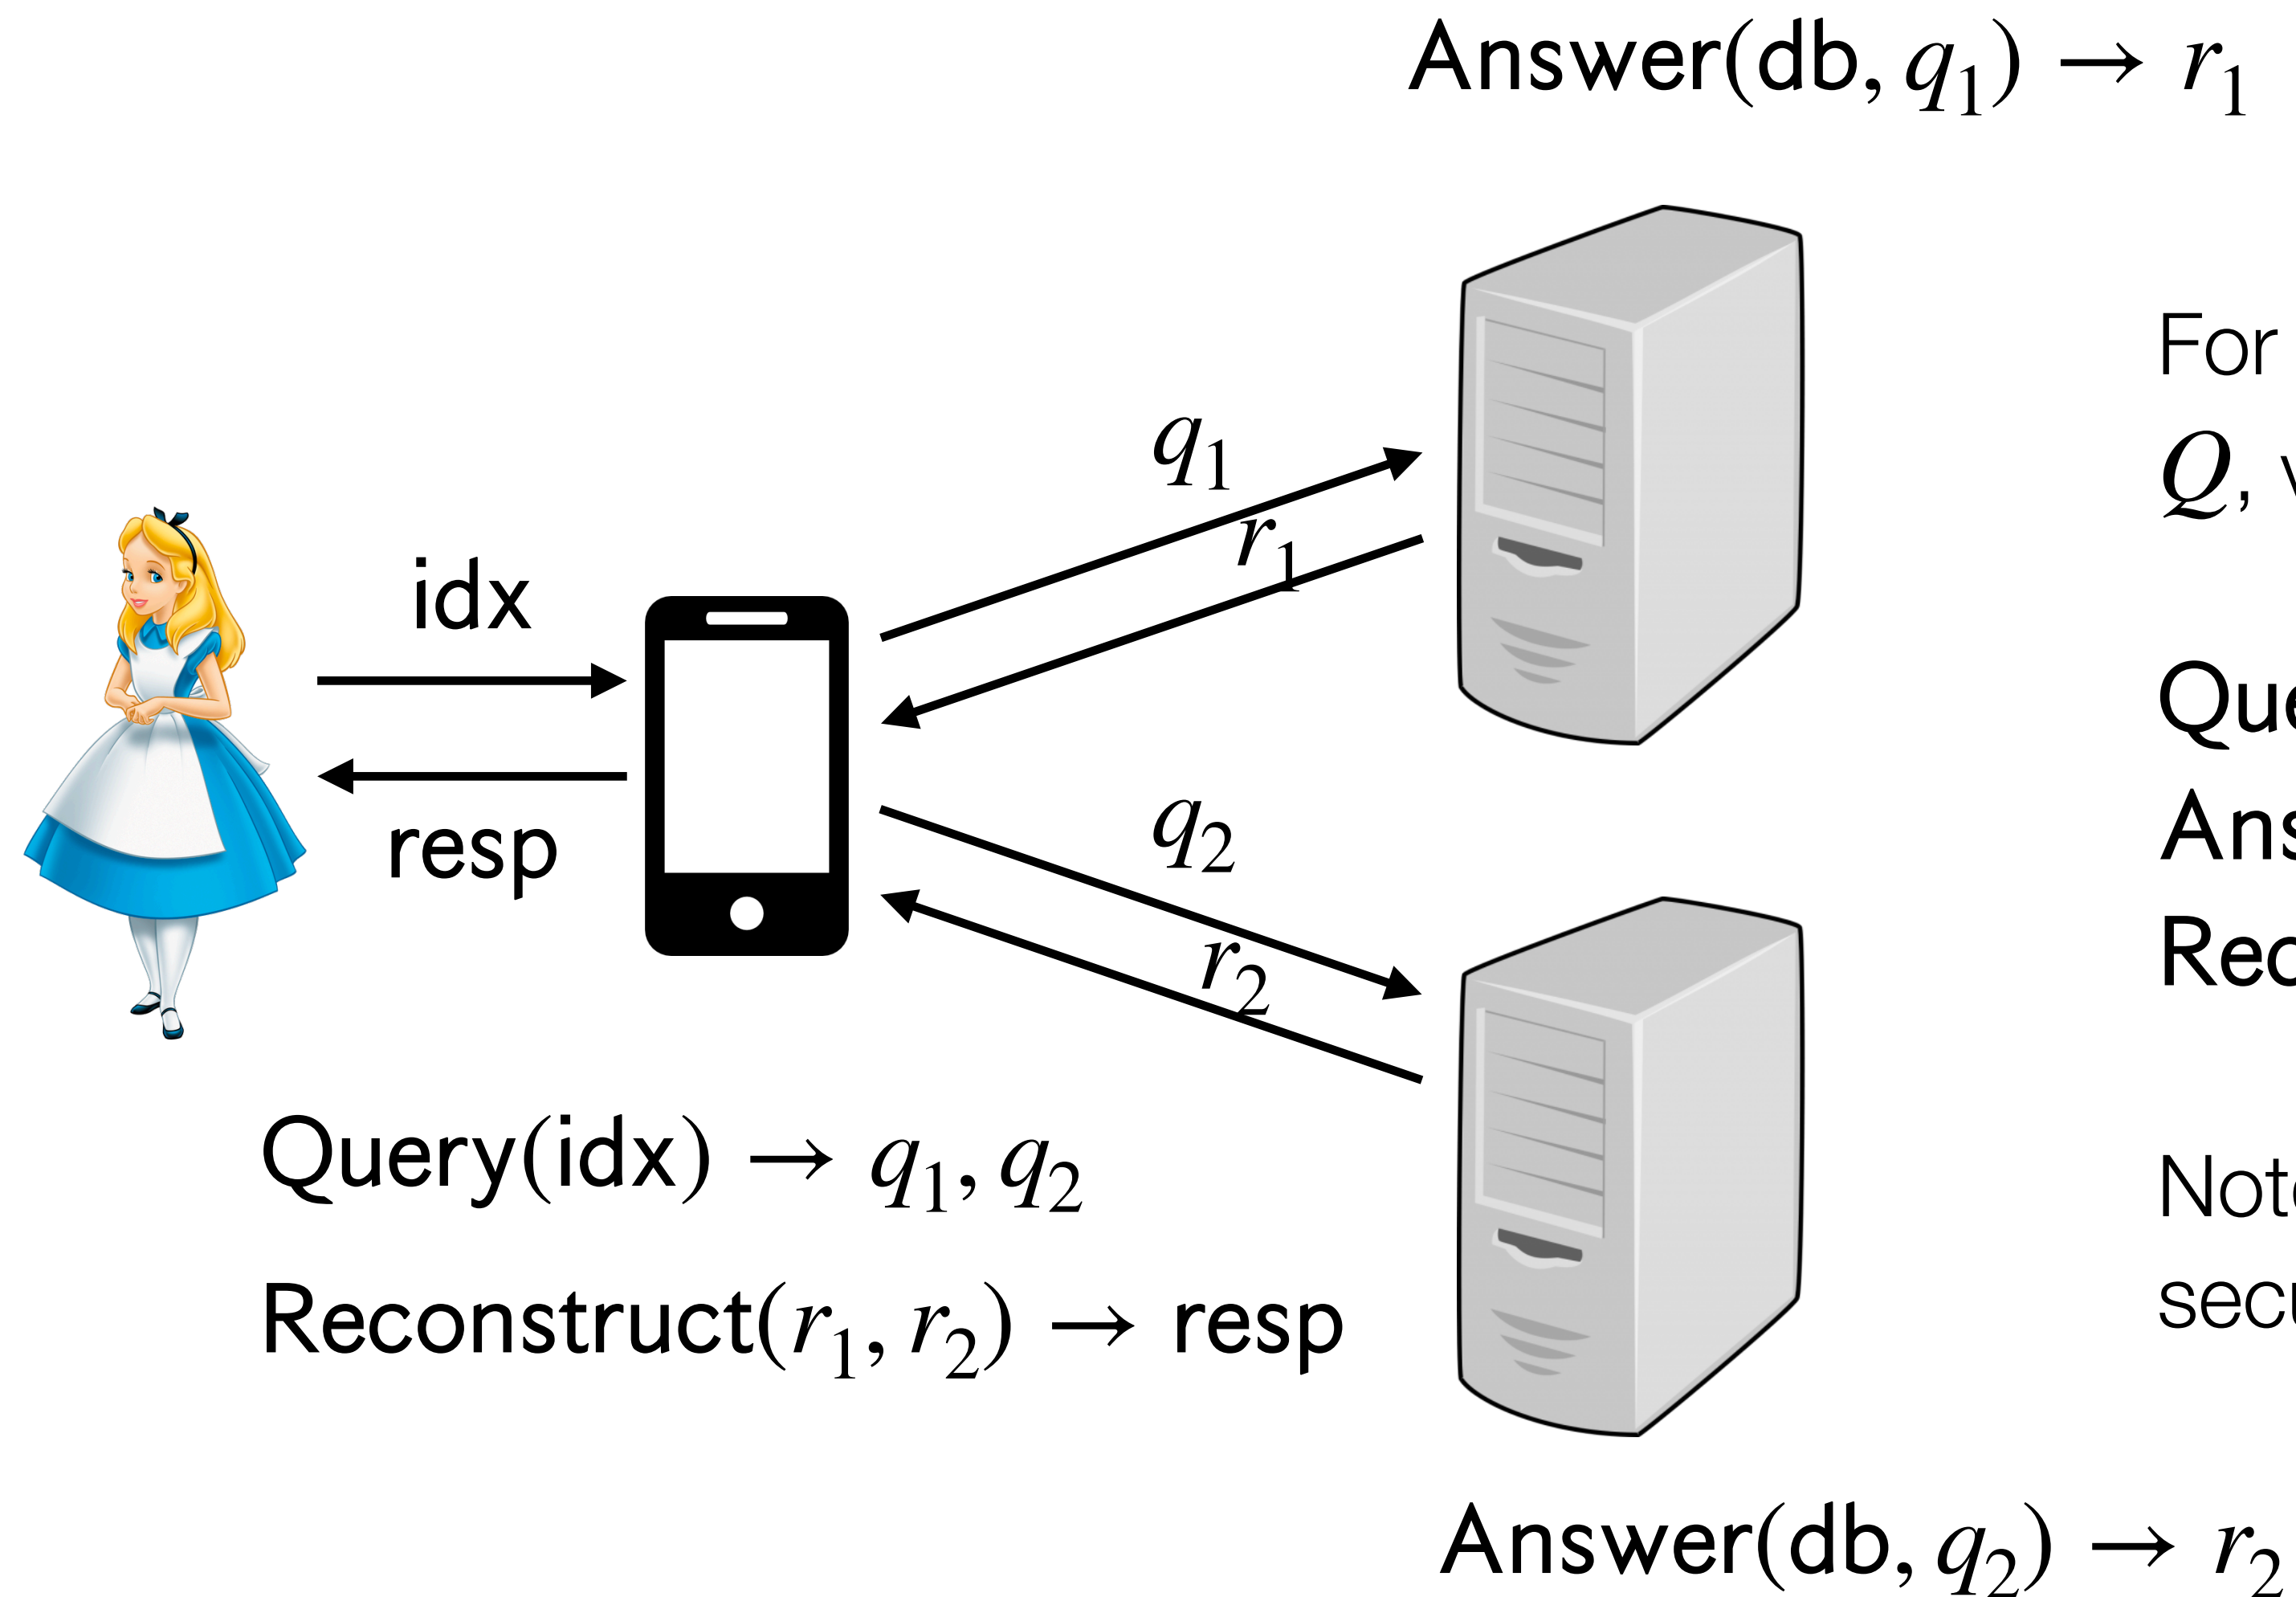
Single-server PIR [KO'97]

Protects privacy if the server is compromised



Requires a cryptographic assumption

# Two-server PIR



For database with  $n$  elements, query space  $Q$ , values in  $\{0,1\}$ , and response space  $R$

Query :  $[n] \rightarrow Q^2$

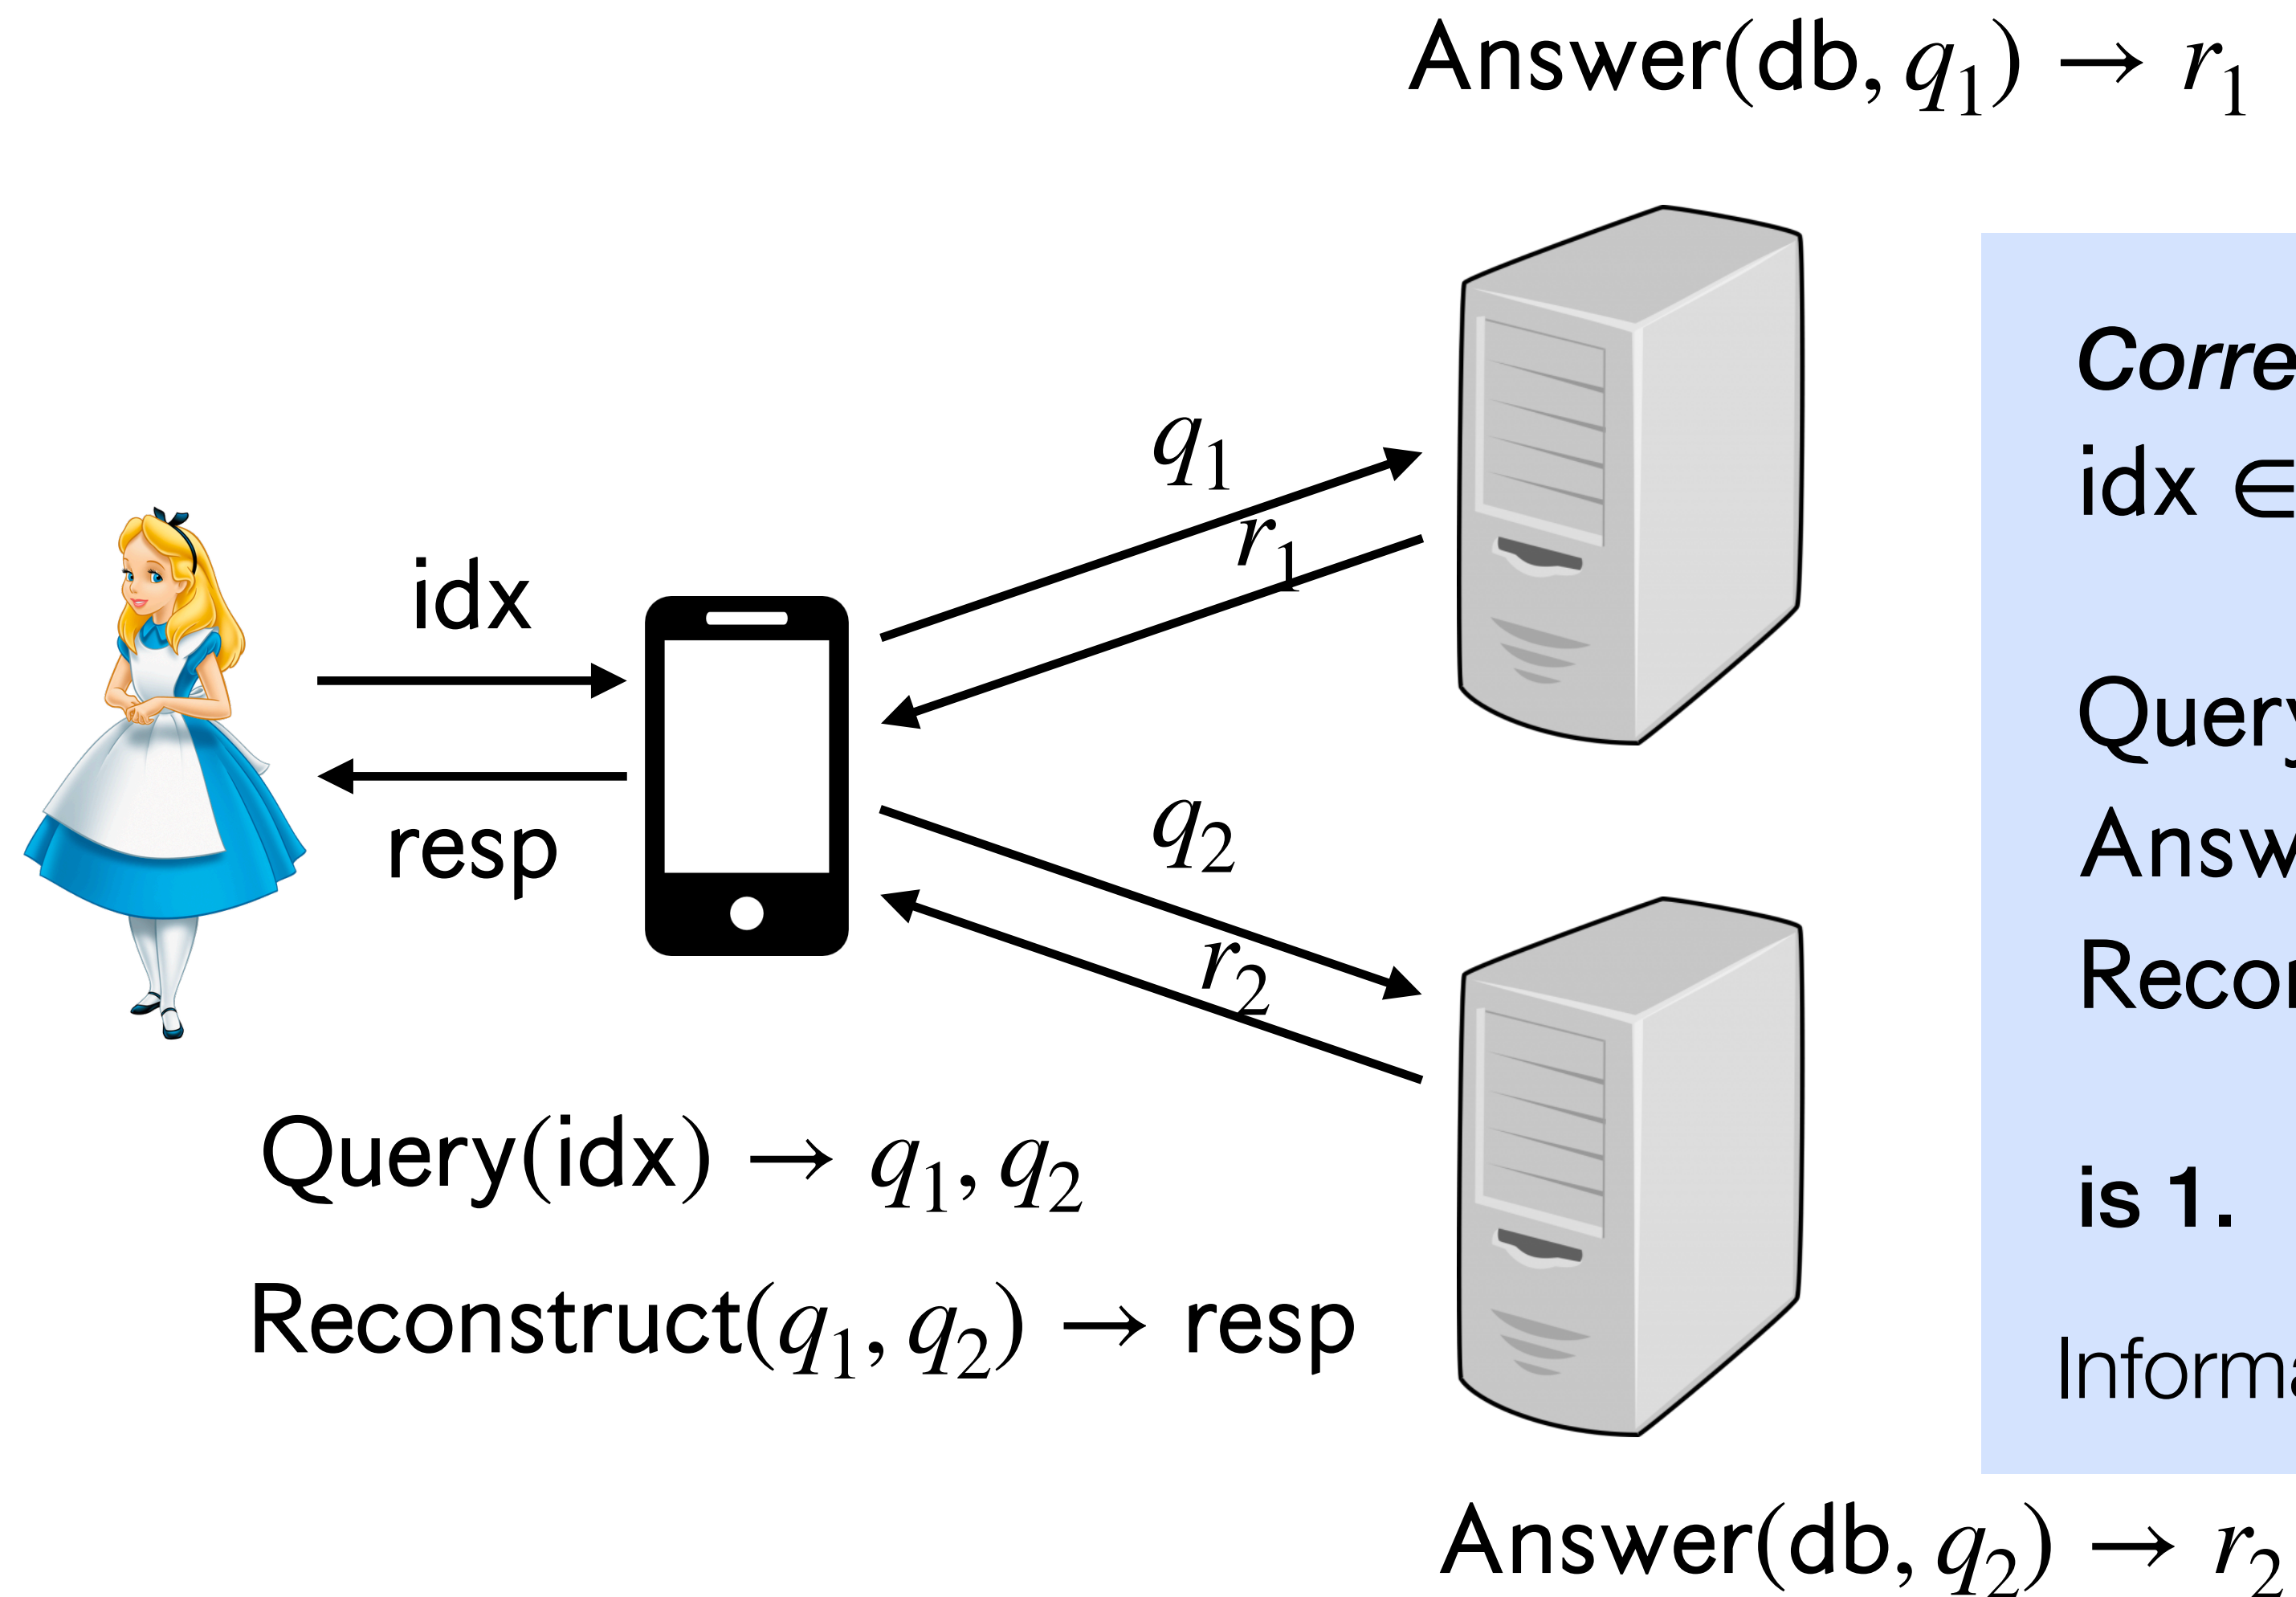
Answer :  $\{0,1\}^n \times Q \rightarrow R$

Reconstruct :  $R^2 \rightarrow \{0,1\}$

Note: **Query** is randomized and takes security parameter as implicit input



# Two-server PIR definitions



**Correctness:** For all  $n \in \mathbb{N}$ ,  $db \in \{0,1\}^n$ ,  $idx \in [n]$ , the probability that:

Query( $idx$ )  $\rightarrow q_1, q_2$

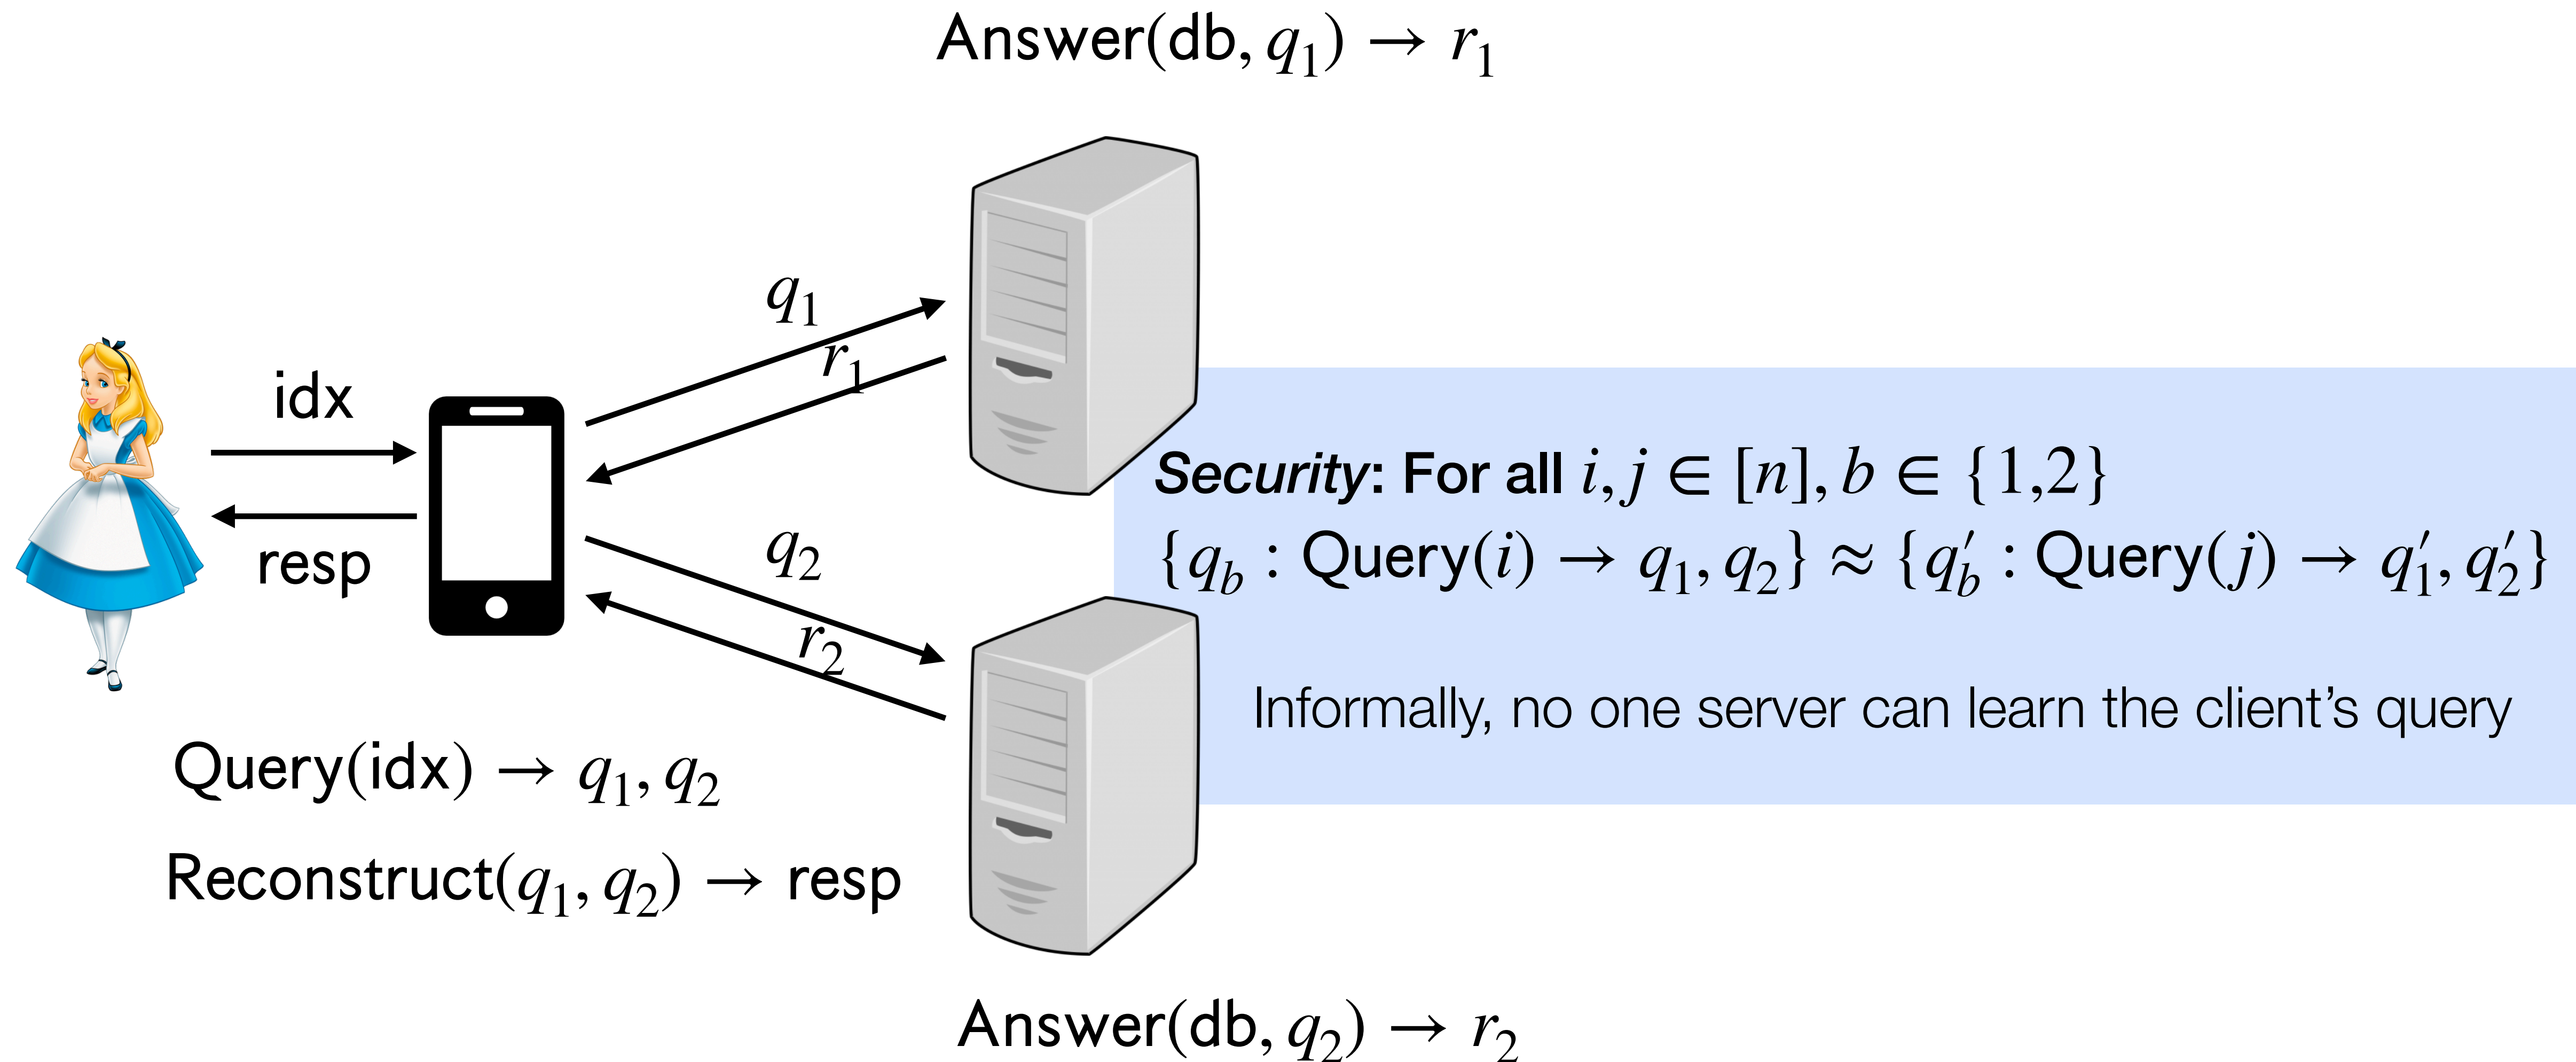
Answer( $db, q_i$ )  $\rightarrow r_i$  for  $i \in \{1,2\}$

Reconstruct( $q_1, q_2$ ) =  $db_{idx}$

**is 1.**

Informally, the client gets the requested record

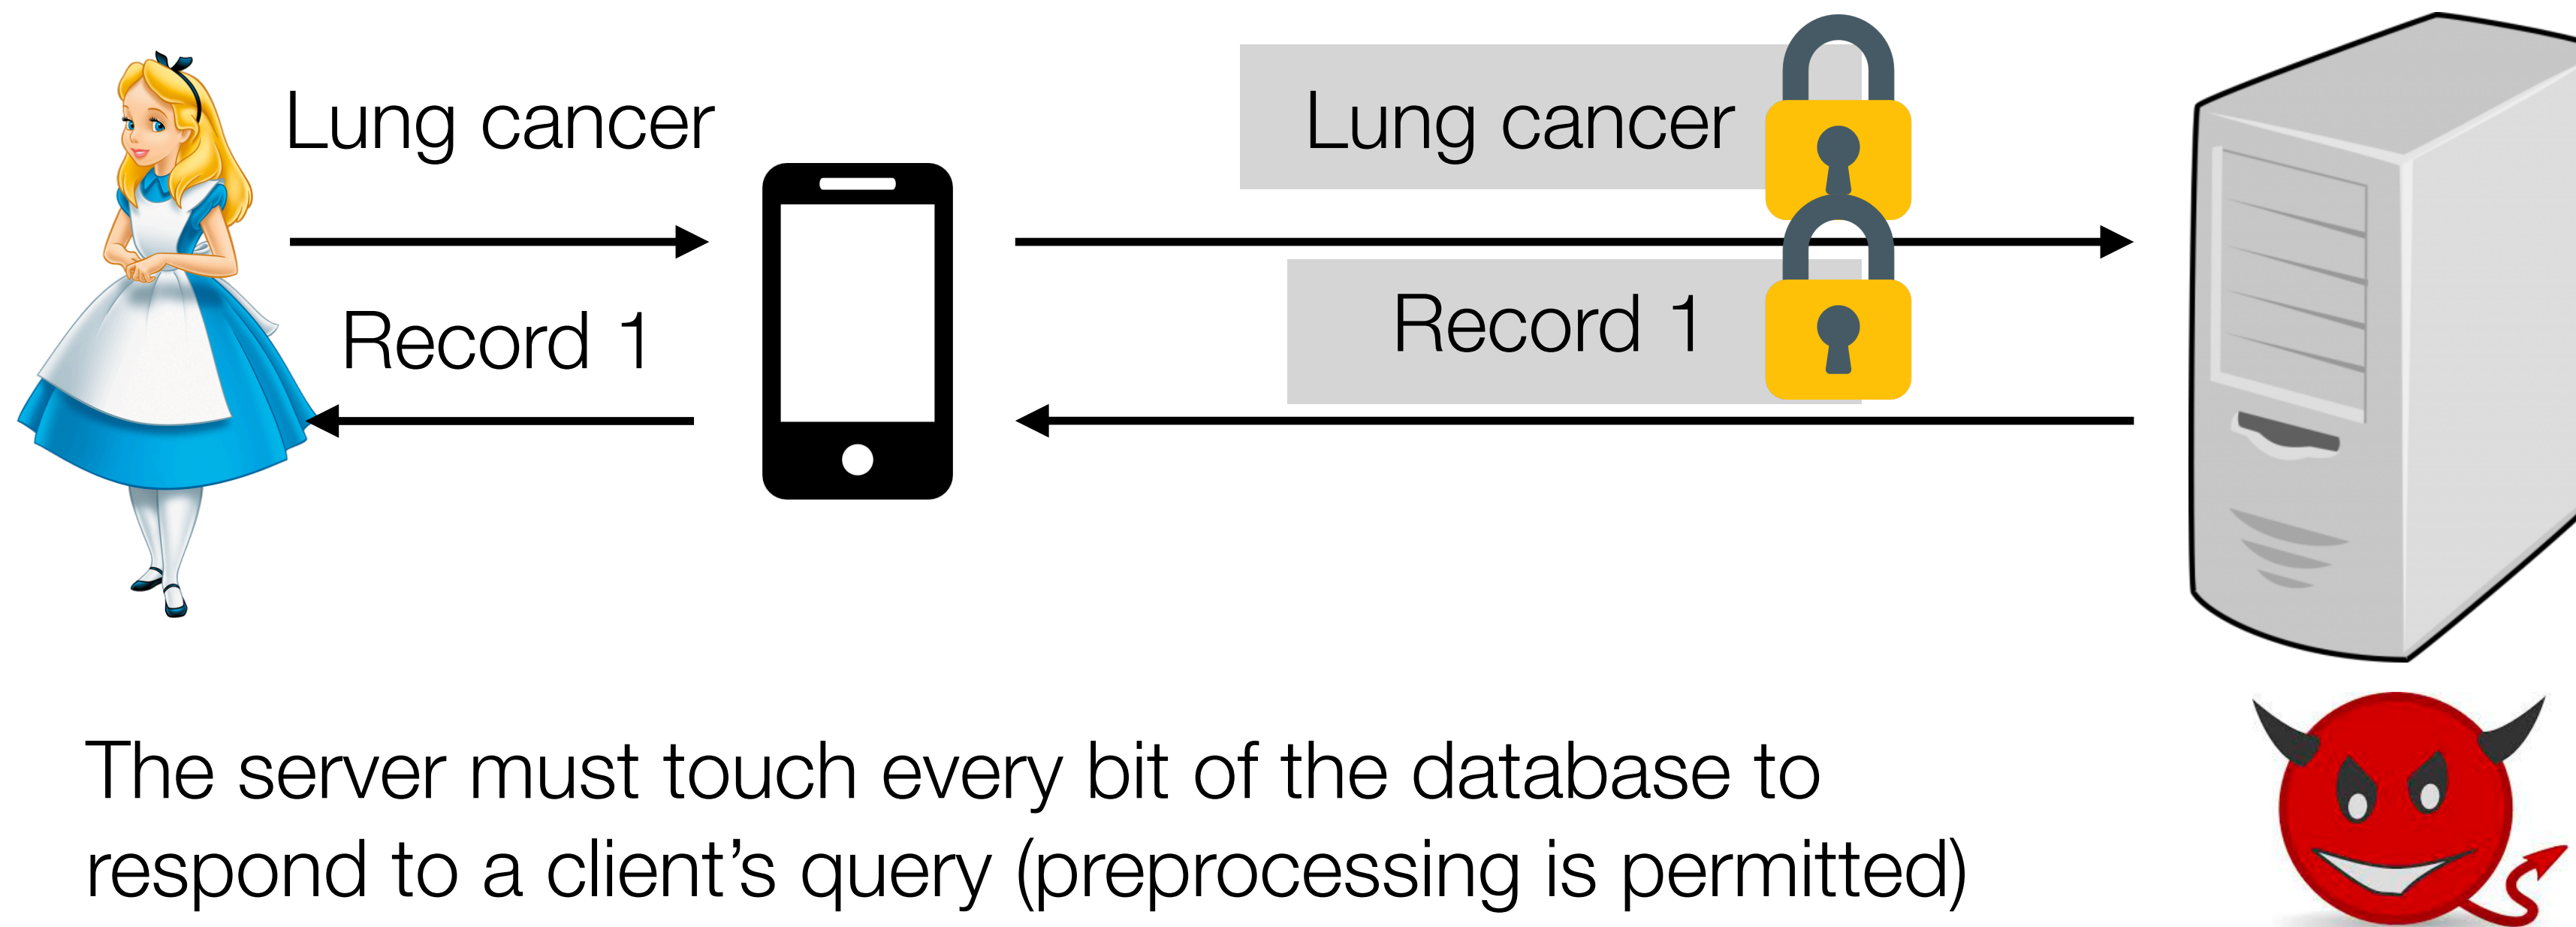
# Two-server PIR definitions



# What we can hope for with PIR

Is it possible to build a PIR scheme that only touches part of the database?

No! If query execution only touches part of the database, then an attacker can learn which part of the database is *not* accessed by the query



The server must touch every bit of the database to respond to a client's query (preprocessing is permitted)  
[BIM'00]

|               |          |
|---------------|----------|
| Heart disease | Record 1 |
| Lung cancer   | Record 2 |
| Diabetes      | Record 3 |
| ...           | ...      |



# Outline

- 1. Function secret sharing**
2. Splinter
3. Searching on encrypted data

# Recap: secret sharing

For some group  $\mathbb{G}$ , split a value  $x \in \mathbb{G}$  into secret shares  $x_1, \dots, x_k \in \mathbb{G}$  such that  $\sum_{i=1}^k x_i = x$

Information-theoretic privacy: Given just  $[x]_b$  for  $b \in [k]$ , adversary learns no information about  $x$

Operations:

- Given  $x$ , generate secret shares by randomly sampling  $x_1, \dots, x_{k-1}$  and setting  $x_k = x - \sum_{i=1}^{k-1} x_i$
- Given  $x_1, \dots, x_k$ , reconstruct  $x$  by computing  $x = \sum_{i=1}^k x_i$

Computing on secret shares:

- Can add secret shares:  $[x] + [y] = [x + y]$
- Can multiply by a constant:  $c \cdot [x] = [c \cdot x]$  (by extension)

# Function secret sharing (FSS)

[Boyle, Gilboa, Ishai]

Idea: Secret share a value rather than a function

Informal properties for  $n$ -party function secret sharing:

- Split function  $f$  into functions  $f_1, \dots, f_n$  where  $f(x) = \sum_{i=1}^n f_i(x)$
- Describe  $f_i$  using a short key  $K_i$
- Key  $K_i$  reveals nothing about  $f$

# Function secret sharing

[Boyle, Gilboa, Ishai]

Security parameter  $\lambda$ , function  $f$

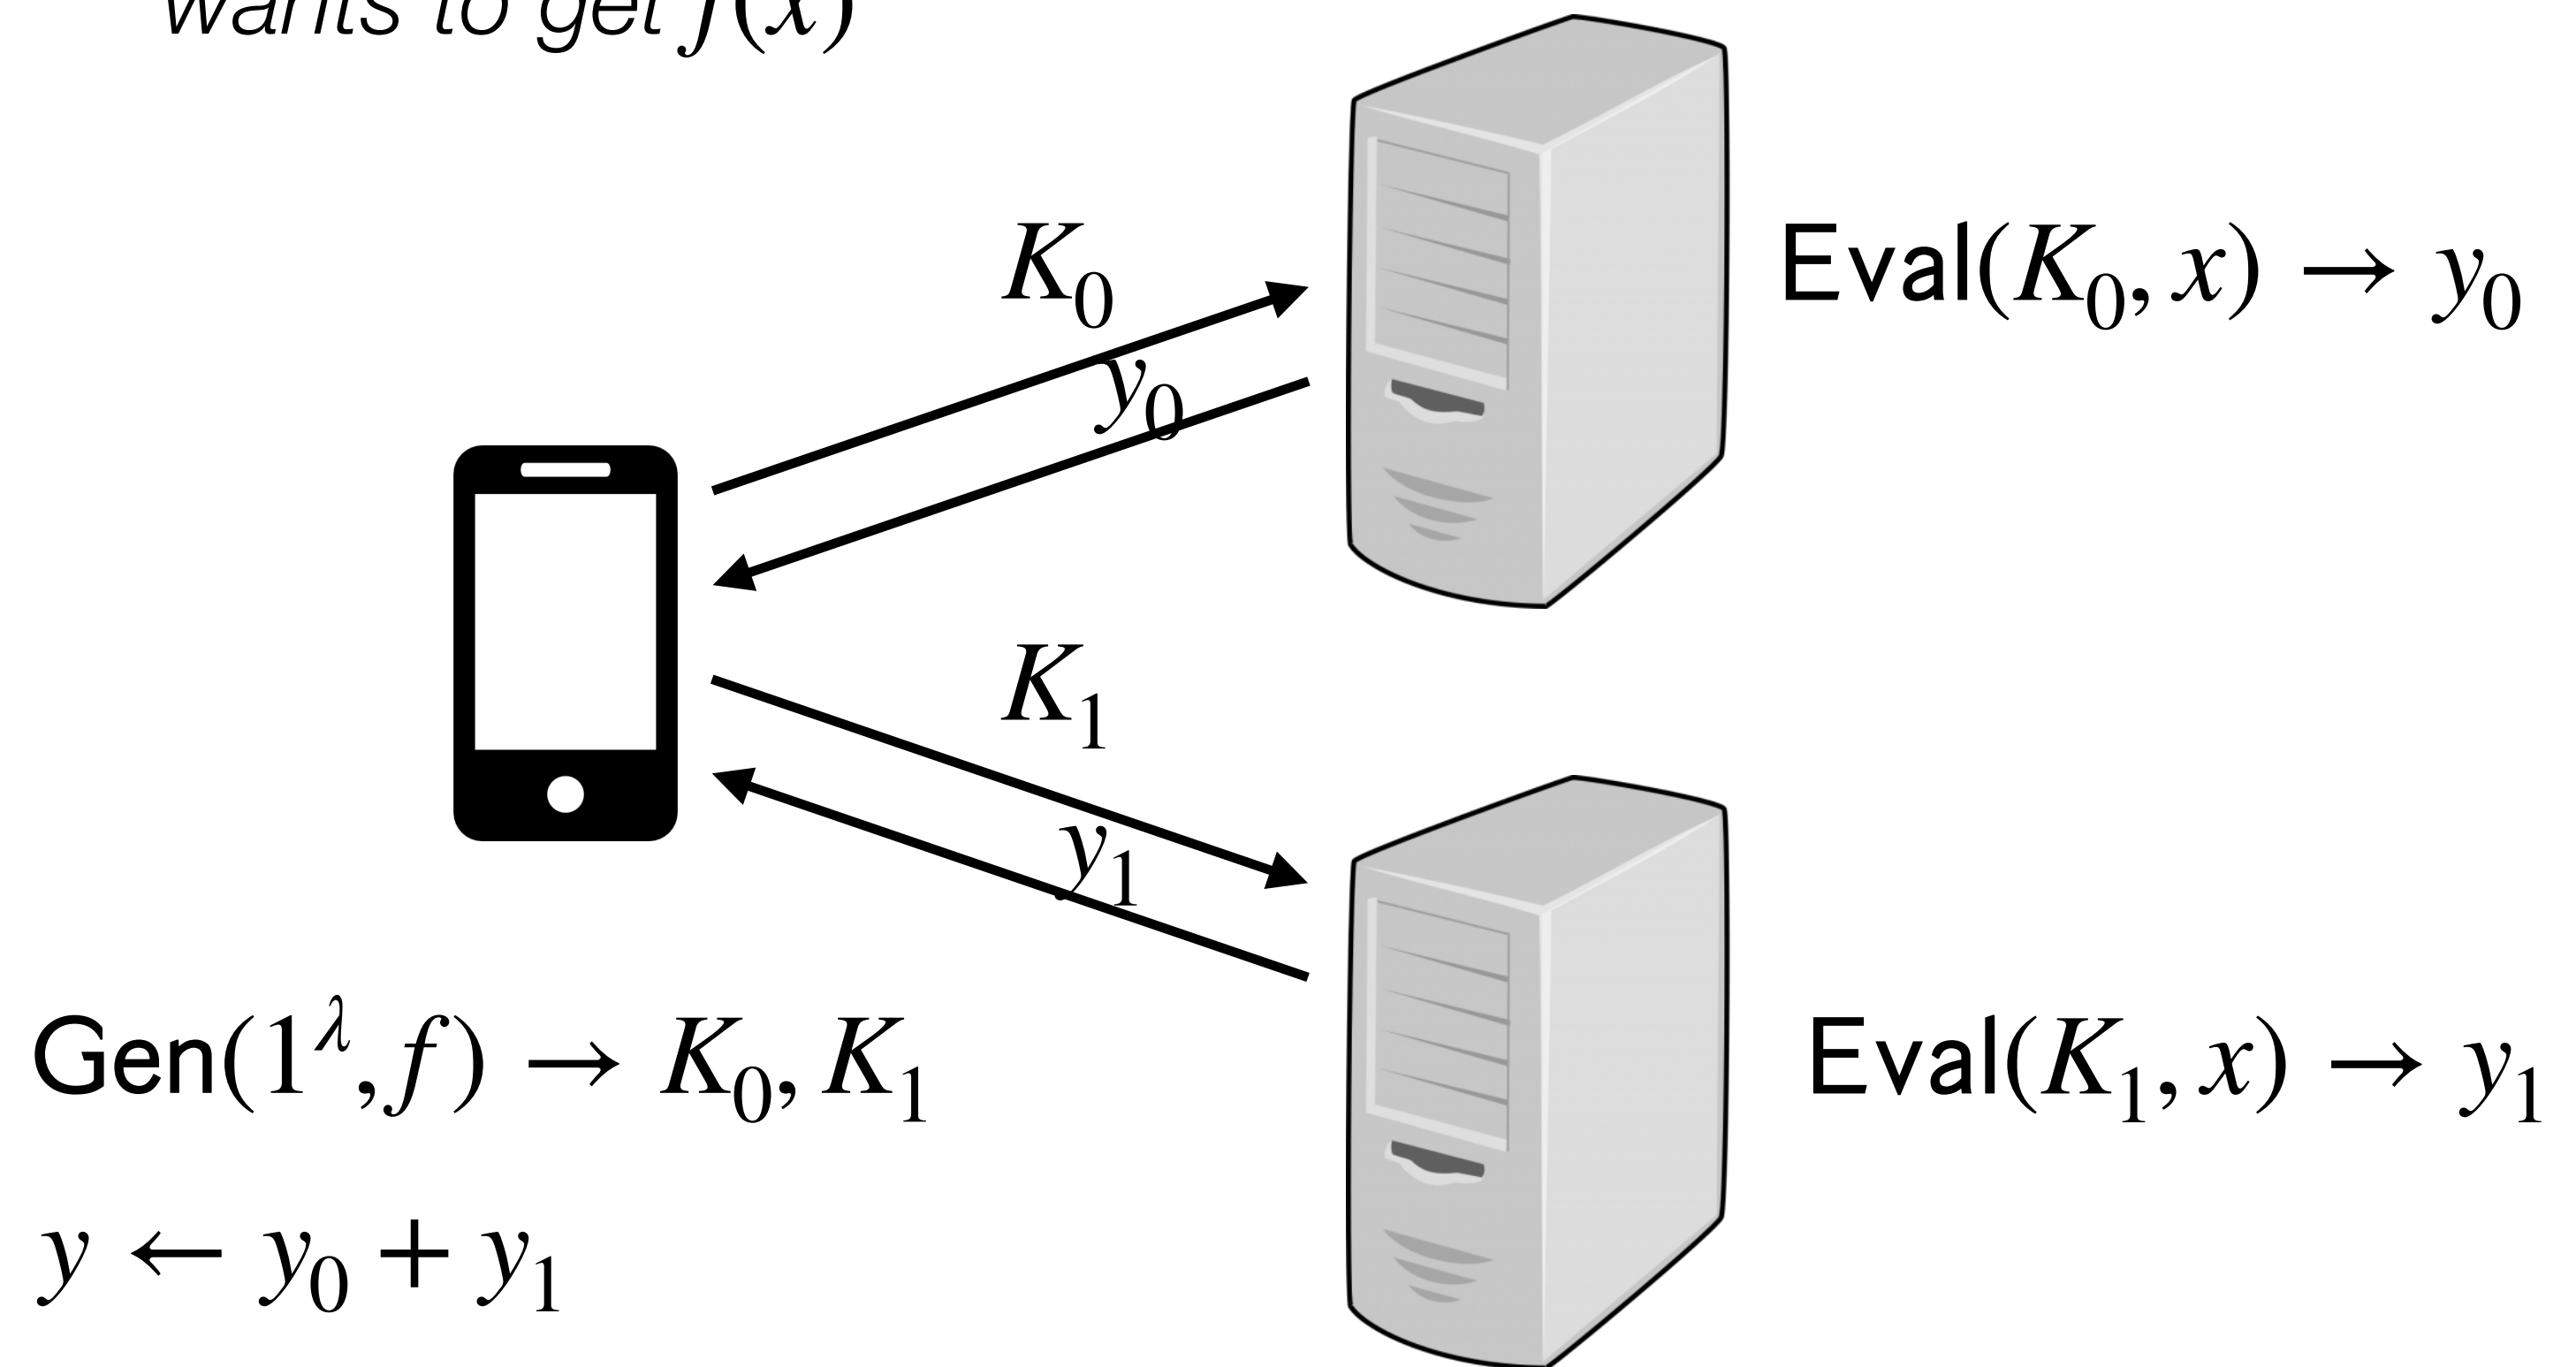
Considering just two parties

Two algorithms:

- $\text{Gen}(1^\lambda, f) \rightarrow K_0, K_1$
- $\text{Eval}(K_i, x) \rightarrow y_i$

*Client has secret  $f$ ,  
wants to get  $f(x)$*

*Servers have  $x$*



# FSS definitions

FSS is defined over a function class  $\mathcal{F}$  with security parameter  $\lambda$  has the following syntax:

- **Gen**( $1^\lambda, f$ )  $\rightarrow K_0, K_1$ : On input security parameter and  $f \in \mathcal{F}$ , output keys  $K_0, K_1$
- **Eval**( $K_i, x$ )  $\rightarrow y_i$ : On input key  $K_i$  (output from **Gen**) and input string  $x \in \mathcal{D}_f$ , where  $\mathcal{D}_f$  is the input domain of function  $f$ , output a share of  $f(x)$

**Correctness:** For all  $f \in \mathcal{F}, x \in \mathcal{D}_f$ ,

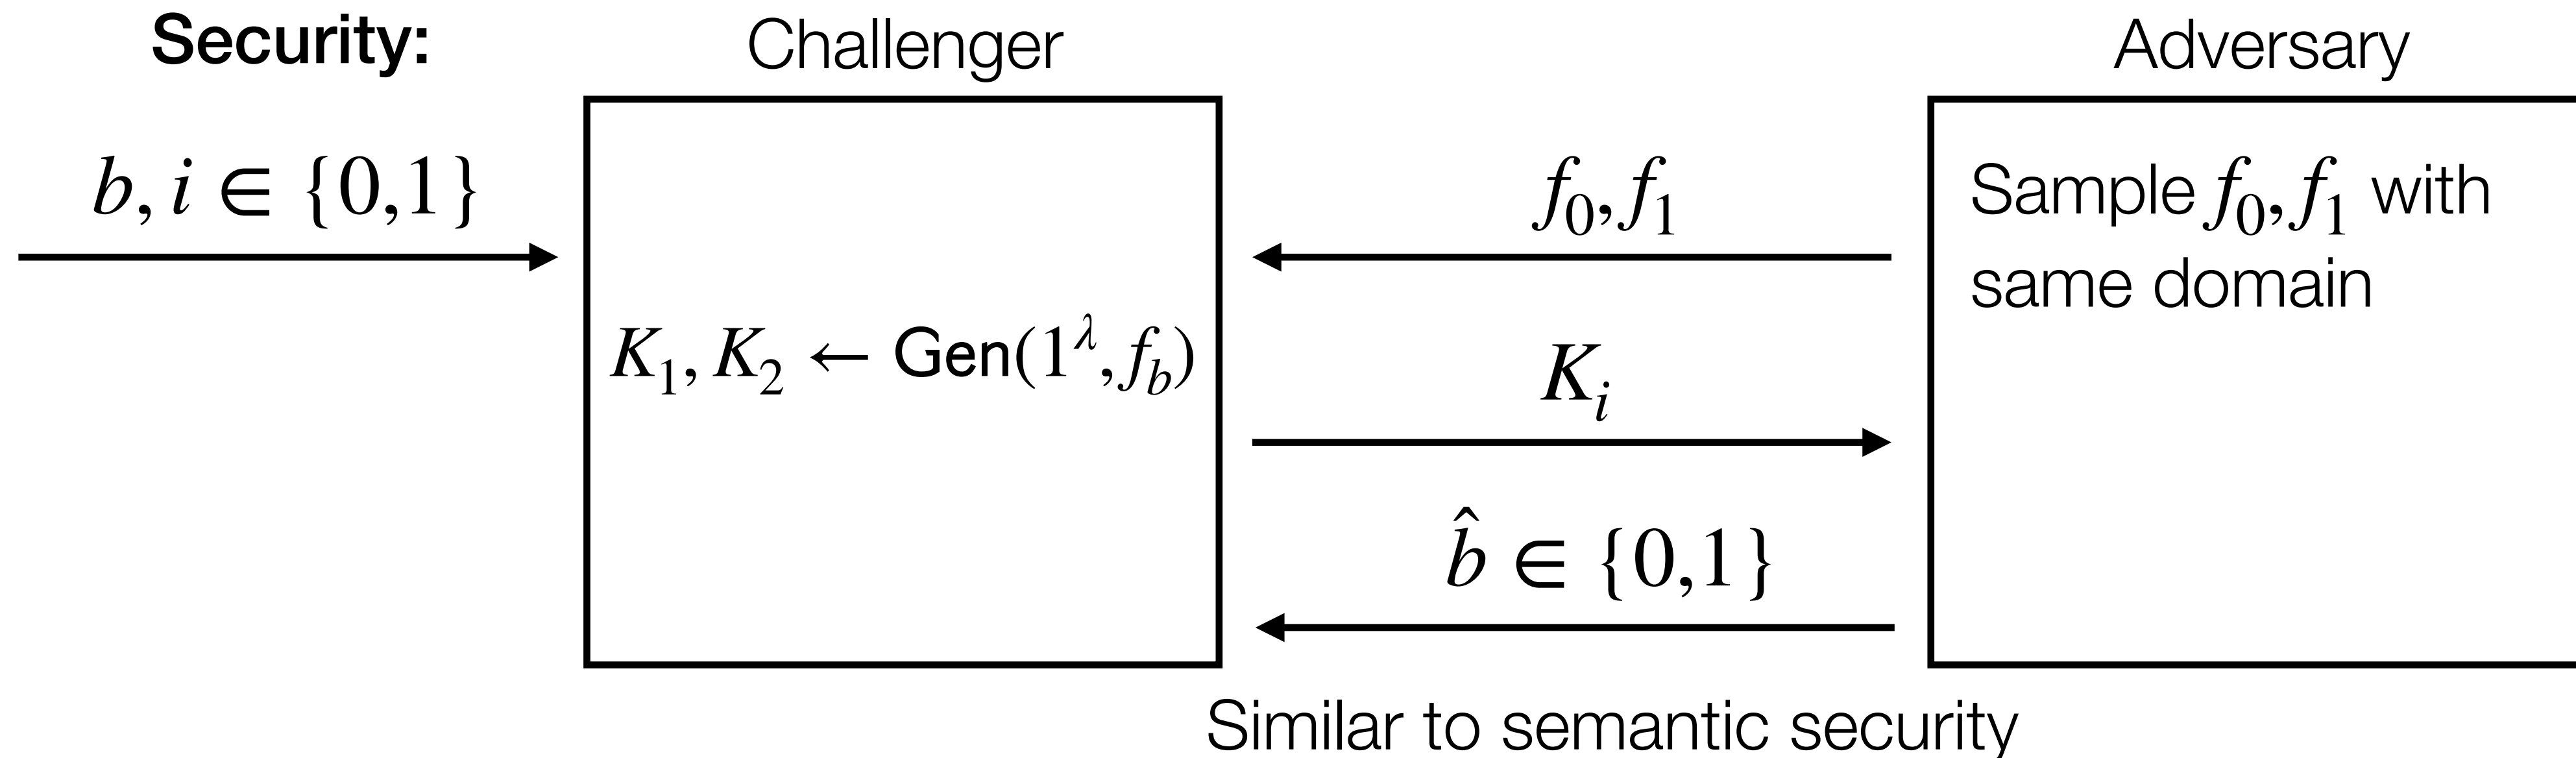
$$\Pr[K_1, K_2 \leftarrow \text{Gen}(1^\lambda, f) : \text{Eval}(K_0, x) + \text{Eval}(K_1, x) = f(x)] = 1$$



# FSS definitions

FSS is defined over a function class  $\mathcal{F}$  with security parameter  $\lambda$  has the following syntax:

- **Gen**( $1^\lambda, f$ )  $\rightarrow K_0, K_1$ : On input security parameter and  $f \in \mathcal{F}$ , output keys  $K_0, K_1$
- **Eval**( $K_i, x$ )  $\rightarrow y_i$ : On input key  $K_i$  (output from **Gen**) and input string  $x \in \mathcal{D}_f$ , where  $\mathcal{D}_f$  is the input domain of function  $f$ , output a share of  $f(x)$



# Functions with efficient FSS constructions


**Point functions:** For  $a, b$ , the corresponding point function  $f$  is defined as  $f(a) = b$  and, for all  $x \neq a, f(x) = 0$

- Corresponding FSS construction: distributed point function (DPF)


**Comparison functions:** For  $a, b$ , the corresponding comparison function  $f$  is defined as  $f(x) = b$  for  $x < a$  and  $f(x) = 0$  for  $x \geq a$

- Corresponding FSS construction: distributed comparison function (DCF)
- Can generalize to  $>$  and intervals

# PIR from distributed point functions



# PIR by keywords from DPFs



# Costs for DPFs and DCFs

FSS is defined over a function class  $\mathcal{F}$  with security parameter  $\lambda$  has the following syntax:

- **Gen**( $1^\lambda, f$ )  $\rightarrow K_0, K_1$ : On input security parameter and  $f \in \mathcal{F}$ , output keys  $K_0, K_1$
- **Eval**( $K_i, x$ )  $\rightarrow y_i$ : On input key  $K_i$  (output from **Gen**) and input string  $x \in \mathcal{D}_f$ , where  $\mathcal{D}_f$  is the input domain of function  $f$ , output a share of  $f(x)$

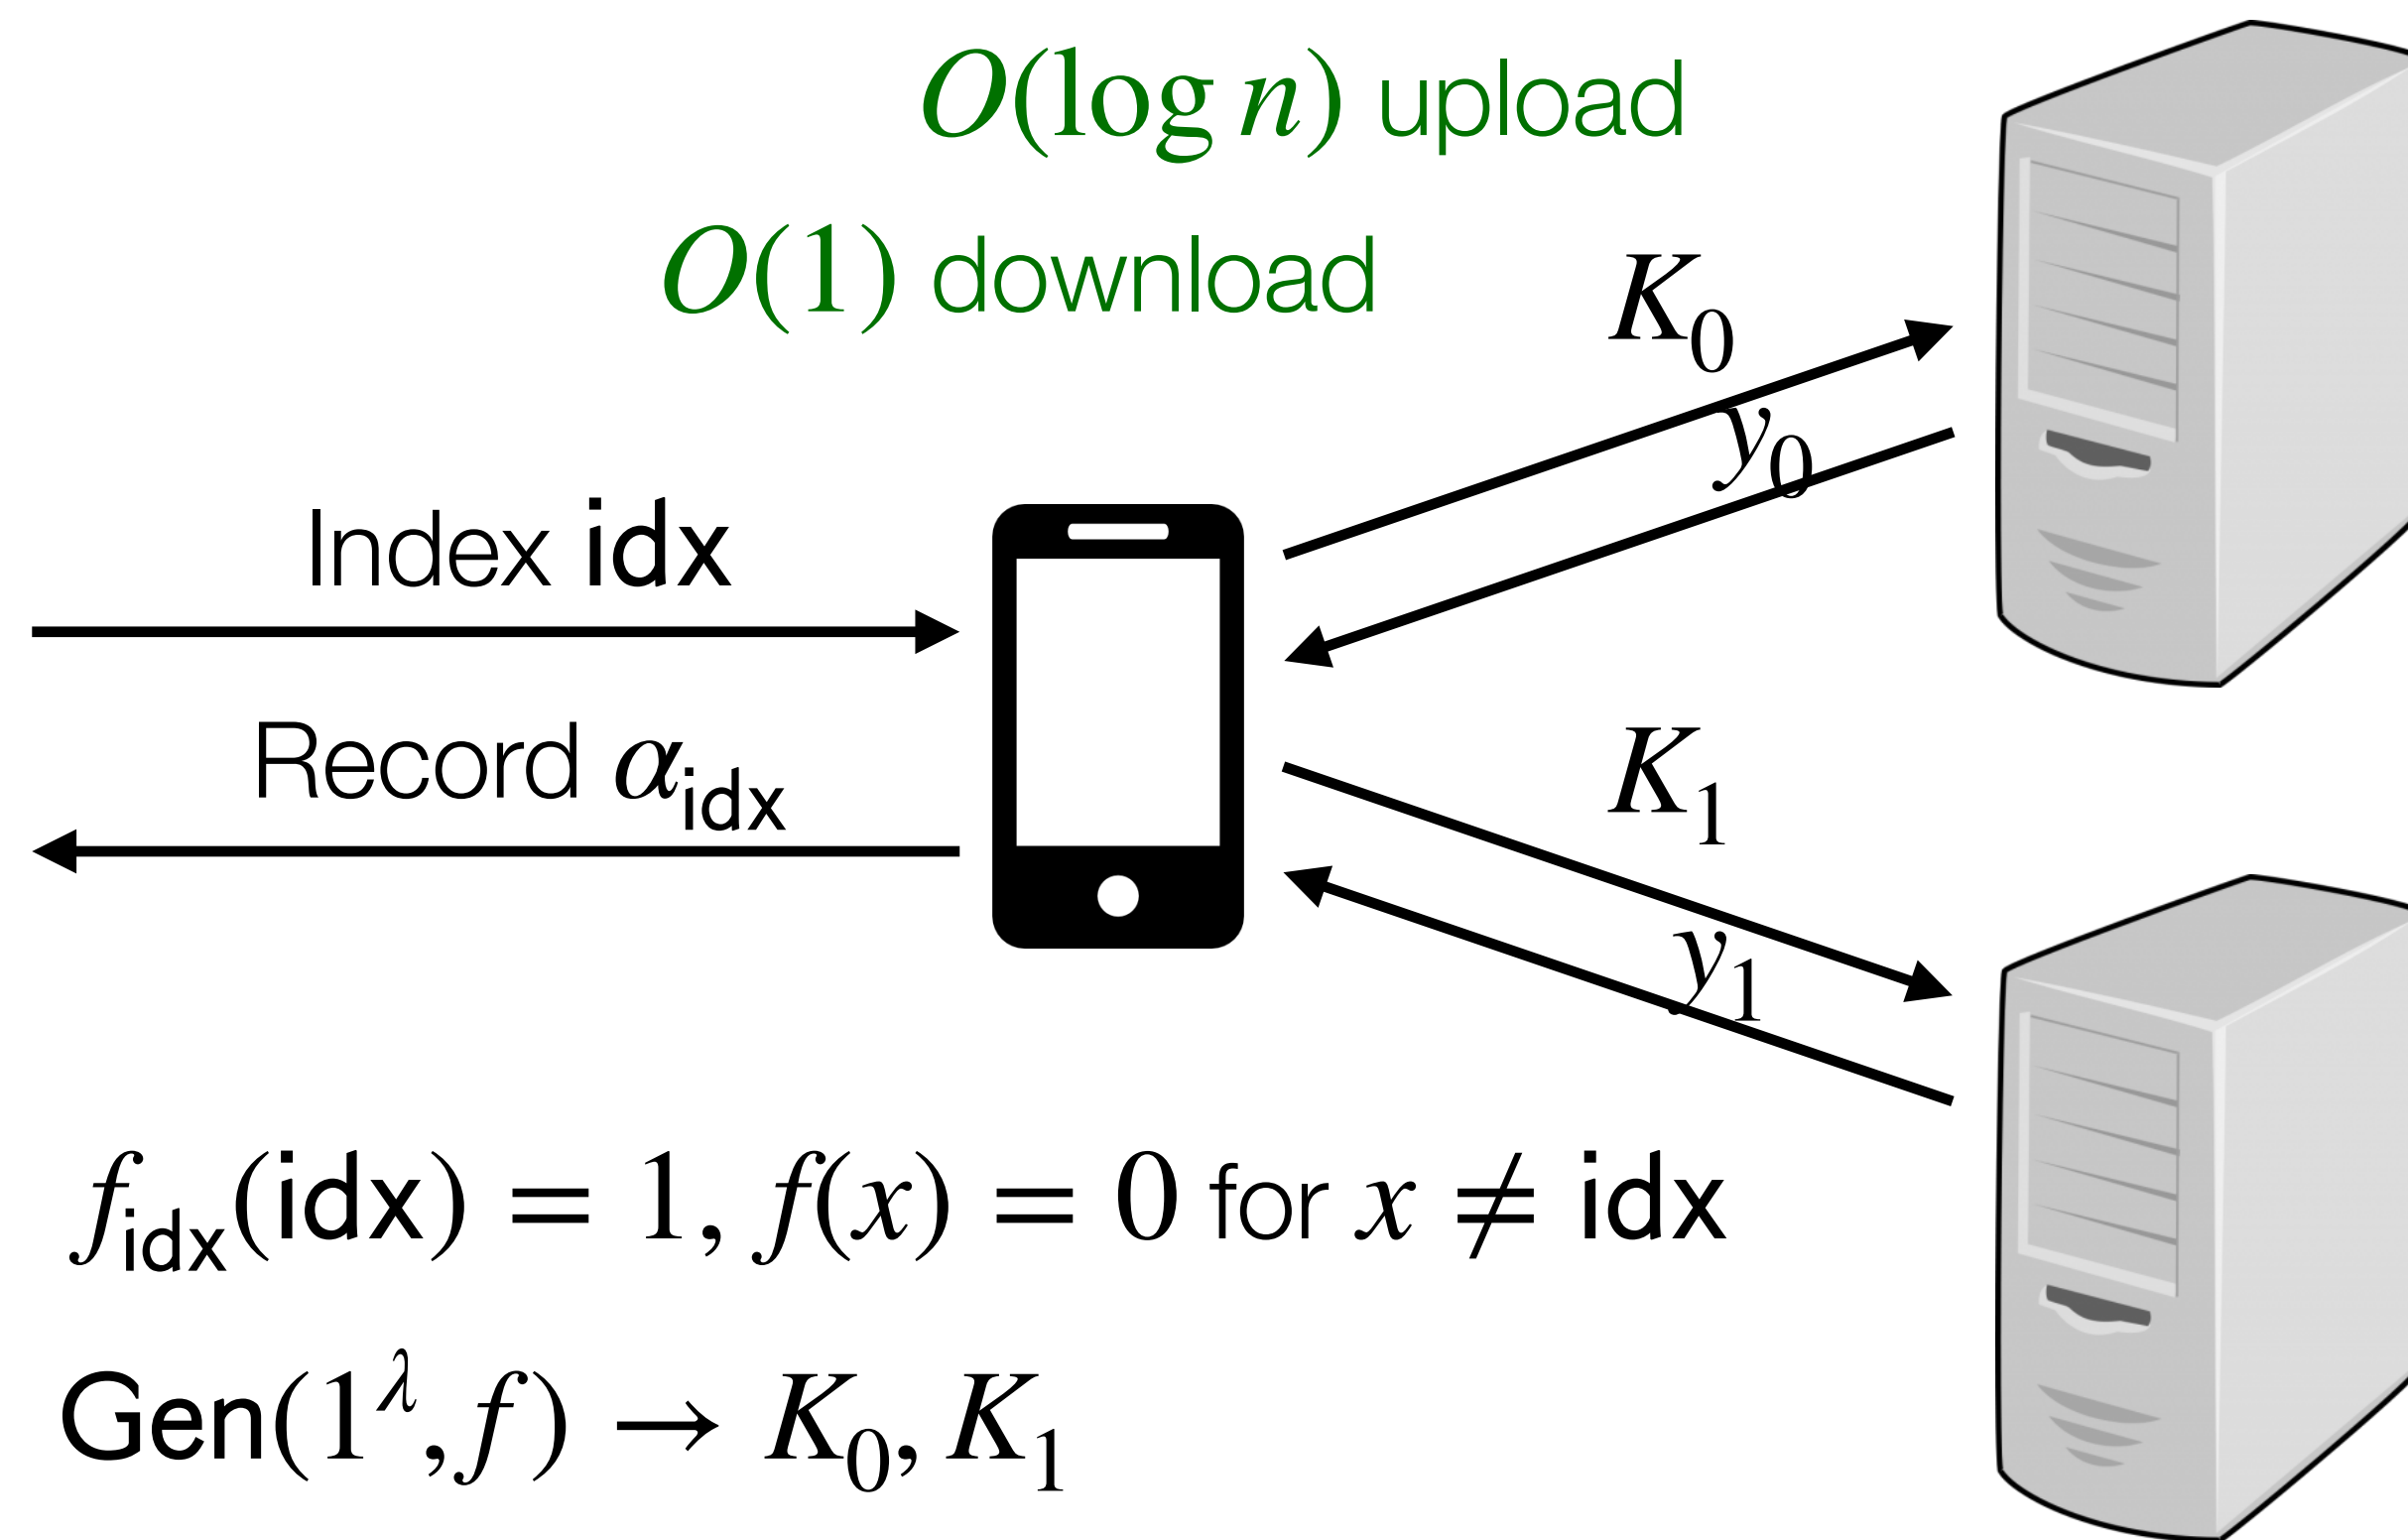
Only AES operations



For  $\mathcal{D} = \{0,1\}^m$ , key size is  $O(m)$



# Costs for PIR from DPFs



$$y \leftarrow y_0 + y_1$$

Linear scan over data


$$\begin{array}{l} 1 : \alpha_1 \\ 2 : \alpha_2 \\ \dots \\ n : \alpha_n \end{array} \quad \begin{array}{l} \text{Eval}(K_0, 1) \cdot \alpha_1 + \\ \text{Eval}(K_0, 2) \cdot \alpha_2 + \\ \dots \\ \text{Eval}(K_0, n) \cdot \alpha_n \end{array} = y_0$$

$$\begin{array}{l} 1 : \alpha_1 \\ 2 : \alpha_2 \\ \dots \\ n : \alpha_n \end{array} \quad \begin{array}{l} \text{Eval}(K_1, 1) \cdot \alpha_1 + \\ \text{Eval}(K_1, 2) \cdot \alpha_2 + \\ \dots \\ \text{Eval}(K_1, n) \cdot \alpha_n \end{array} = y_1$$

# DPF construction intuition

**Observation:** If  $K_0$  outputs  $y_0$ ,  $K_1$  outputs  $y_1$ , and  $y_0 = y_1$ , then if we assemble shares as  $y_0 - y_1$ ,  $(y_0, y_1)$  are shares of 0.

**Construction idea:** Each server can use its key to construct a tree that is identical to the other server's tree (secret shares of 0) *except* along a path from the root to the non-zero (“special”) point.




Each node is associated with a seed and a control bit

**Invariant:** Seed and control bit are identical across the two evaluations, except along the special path

# DPF construction intuition

**Construction idea:** Each server can use its key to construct a tree that is identical to the other server's tree (secret shares of 0) *except* along a path from the root to the non-zero (“special”) point.

**Invariant:** Seed and control bit are identical across the two evaluations, except along the special path



Easy to maintain invariant at root (always on special path)

What about when we leave the special path?

Idea: use a “correction word” to set the node values to the same value when leaving the special path

-  $[s] \oplus [t] \cdot \mathbf{CW}$  where  $s$  is the seed,  $t$  is the control bit, and  $\mathbf{CW}$  is the correction word — correction conditioned on  $t$

# Applications of FSS beyond PIR

Privately writing to a shared database

- Want to hide which element the client is writing to
- Applications to metadata-hiding messaging (see later in class) [Riposte]

Private aggregate statistics

- Privately collect histograms
- Compute most popular set of strings without revealing all strings [Poplar]

Multiparty computation

- Execute some computation across parties, but without each party revealing its inputs to the other parties [BGI19]

# Outline

1. Function secret sharing
- 2. Splinter**
3. Searching on encrypted data



# Splinter

[Wang, Yun, Goldwasser, Vaikuntanathan, Zaharia]

Gap between PIR and what applications need:

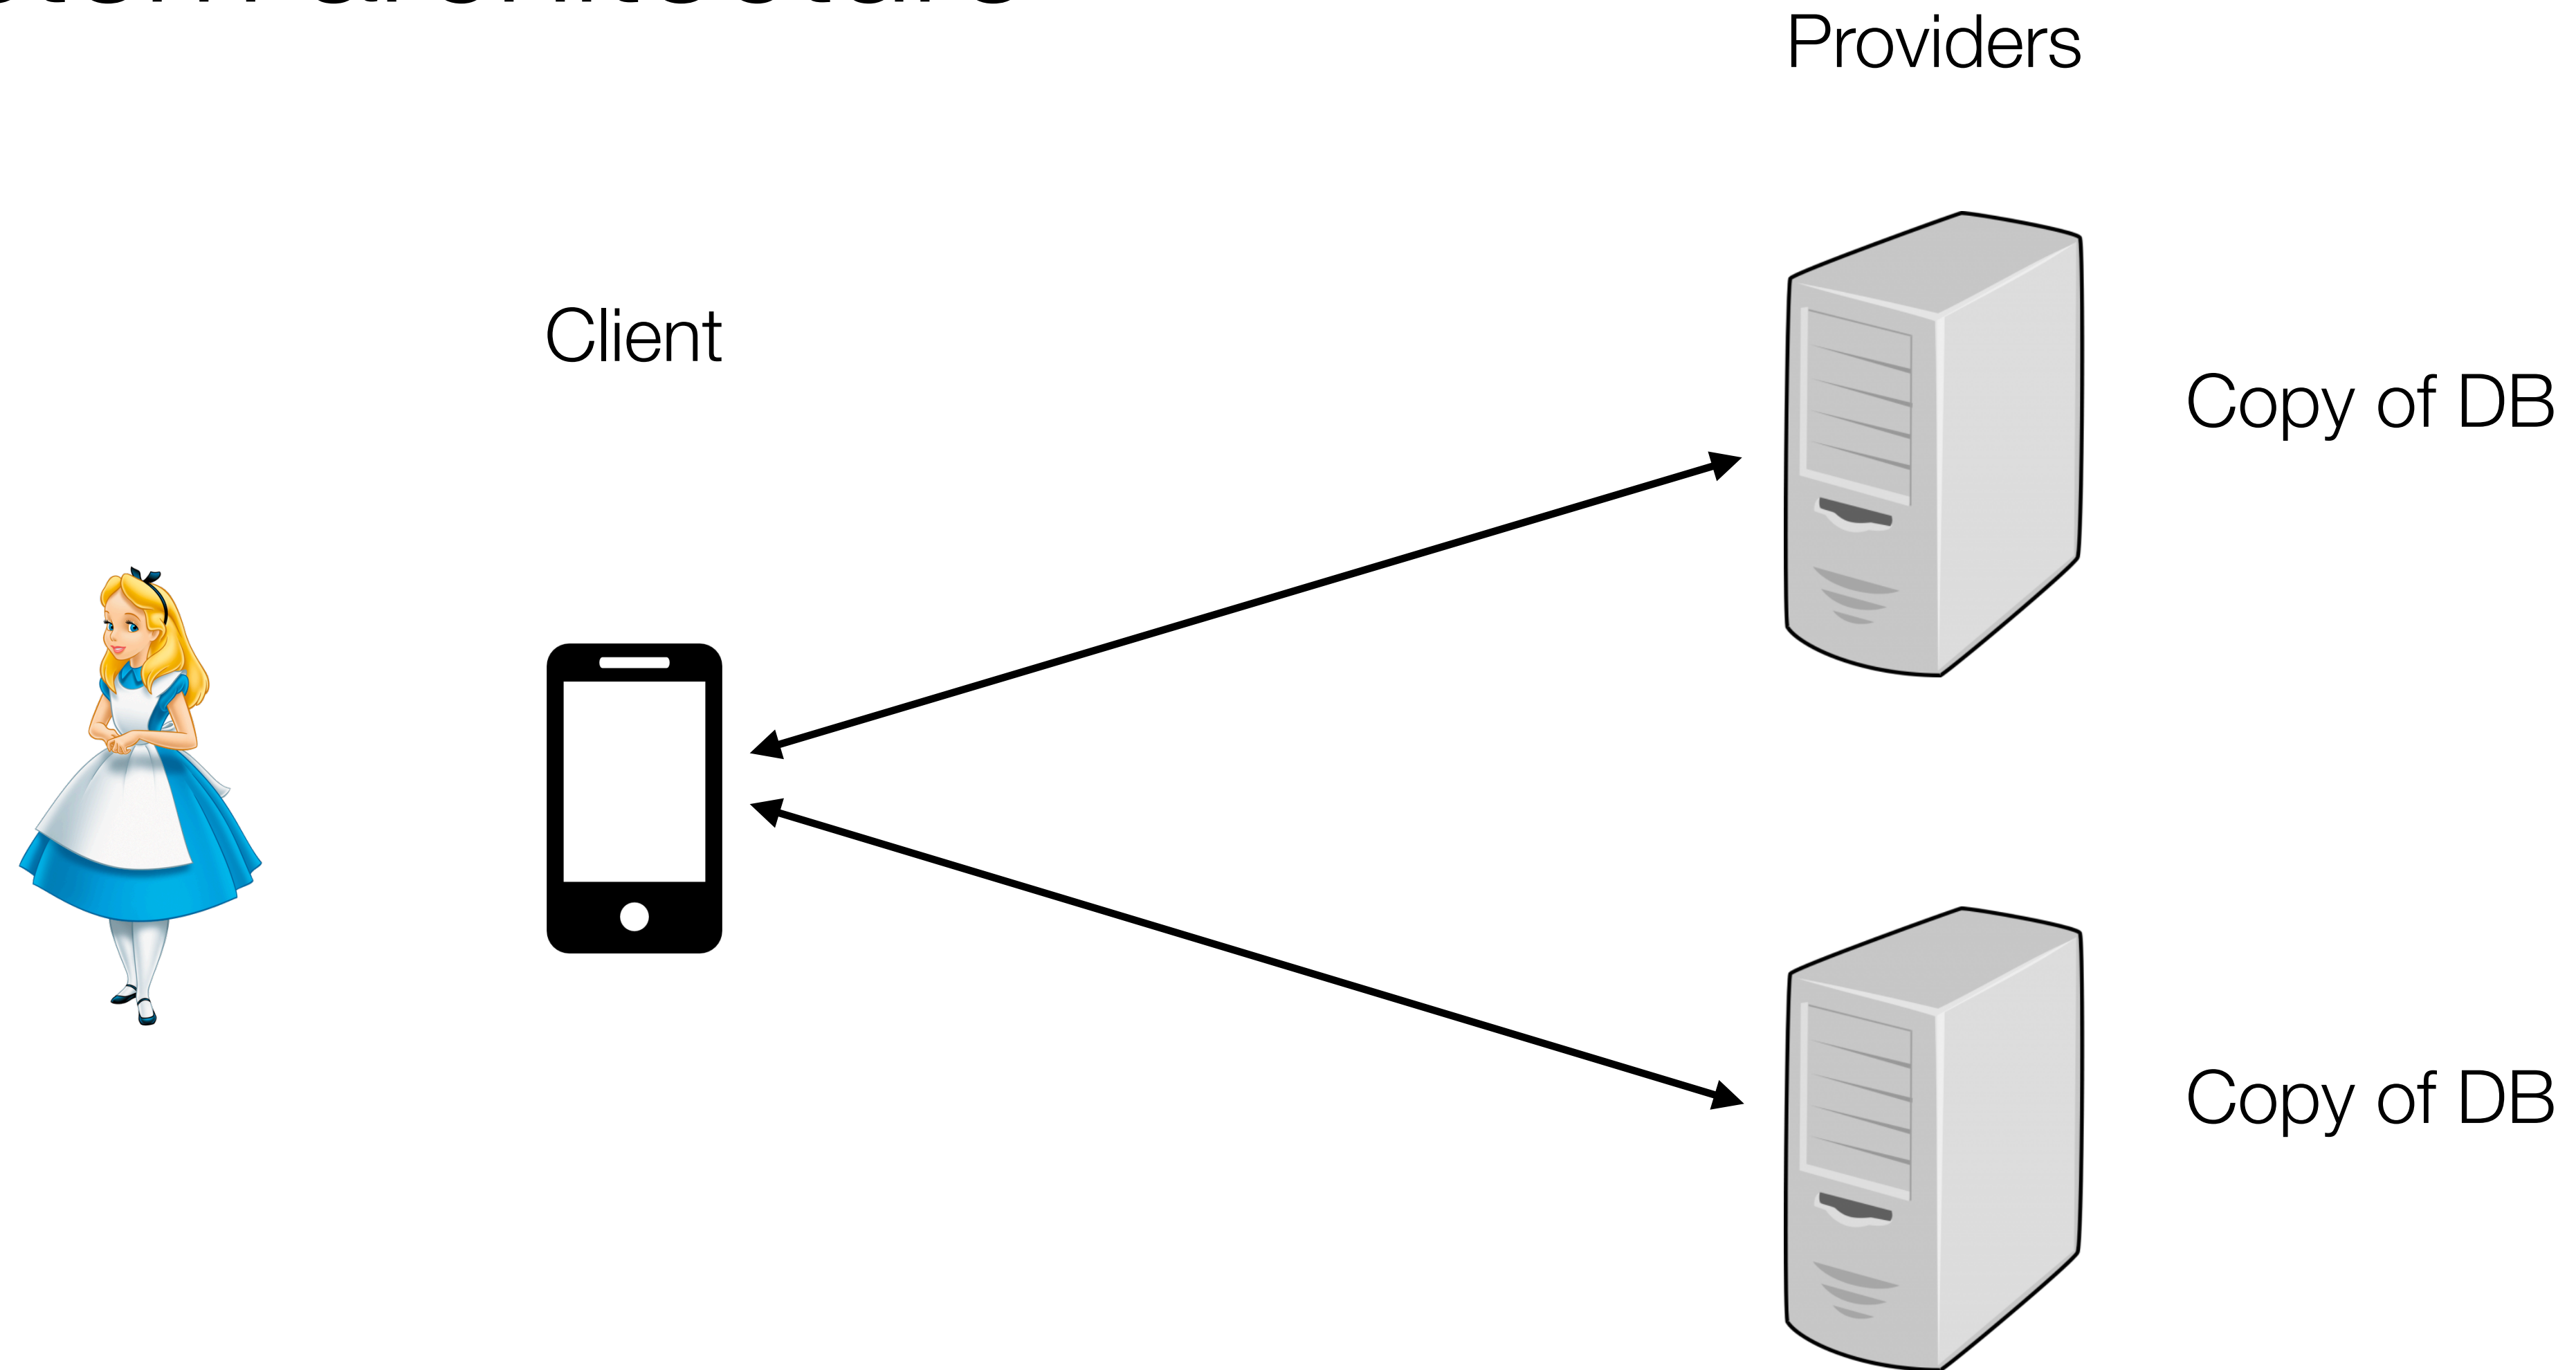
- Private information retrieval: Privately look up a record in a database by some index
- Applications need: Private SQL-like queries based on filters

Splinter helps to bridge this gap

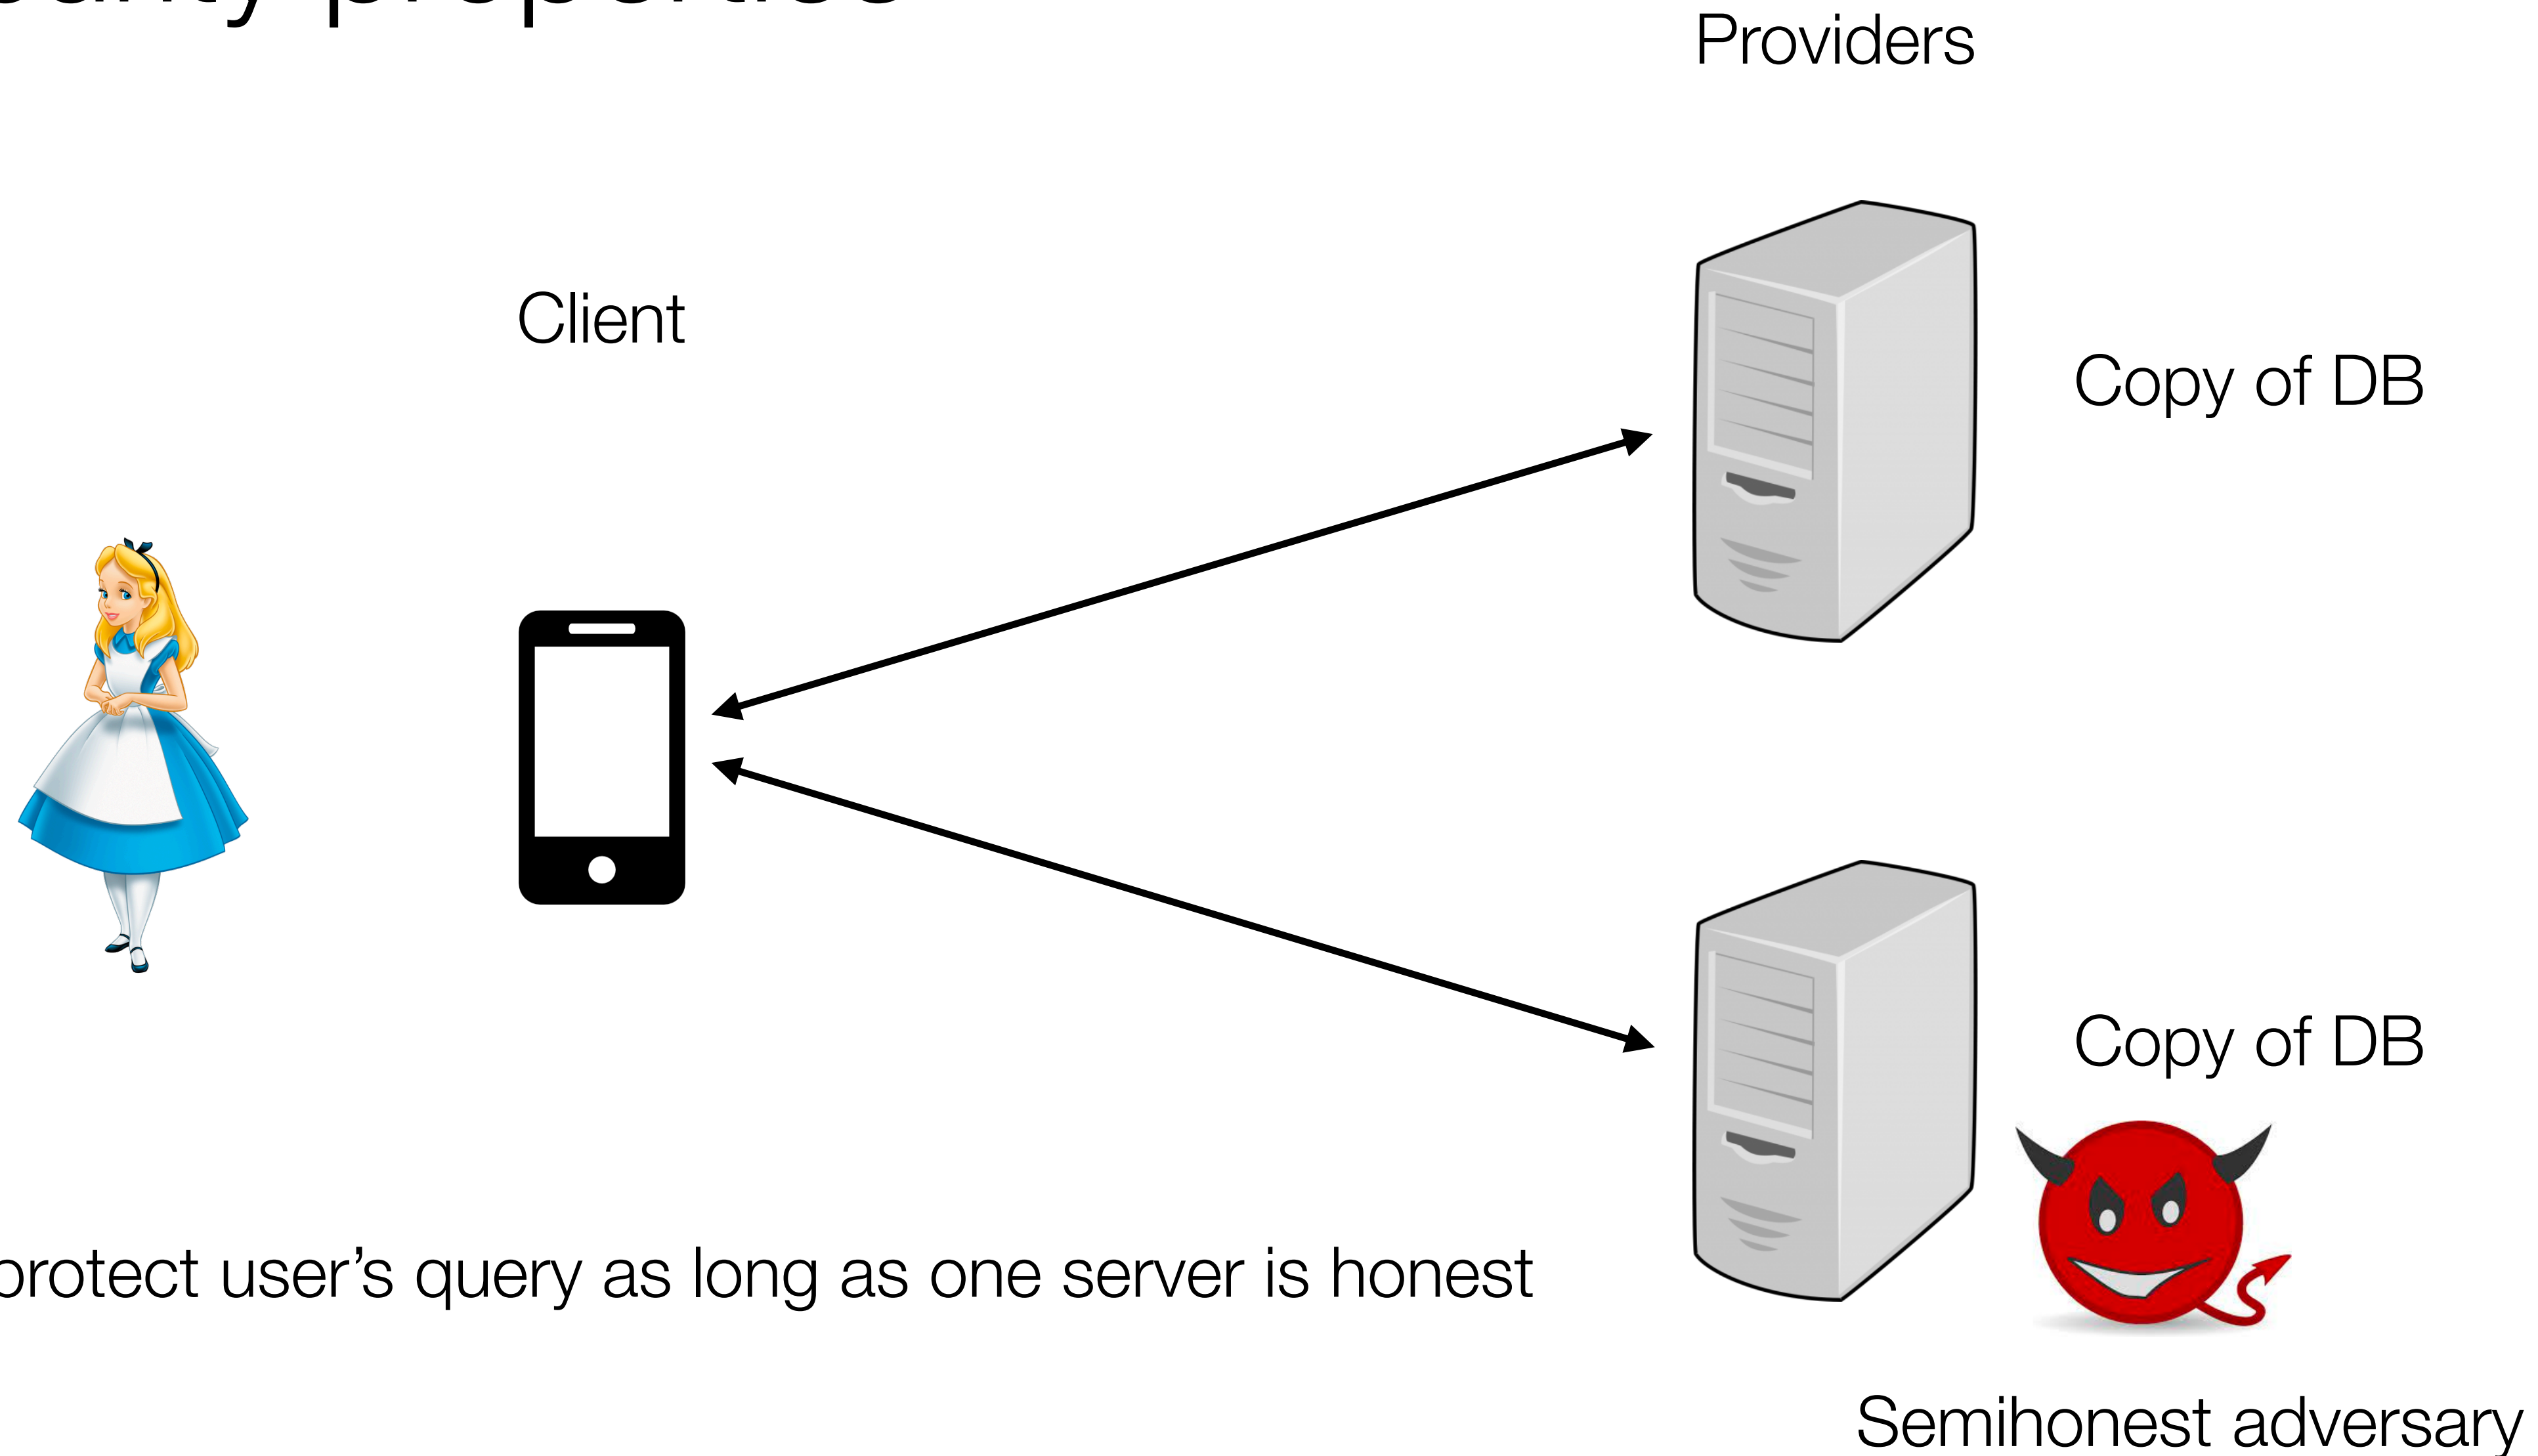
```
SELECT TOP 10 flights from flights  
WHERE source = SFO AND dest = EWR  
ORDER BY price
```

```
SELECT AVG(price) WHERE month=3  
AND origin=SFO AND dest=EWR
```

# System architecture



# Security properties



# Security properties

Splinter hides some query parameters, but not the query structure

Column names, aggregation function, and other features of query structure are not hidden from the server (only the “?” are hidden)

```
SELECT TOP 10 flights from flights  
WHERE source = ? AND dest = ?  
ORDER BY price
```

# Queries supported

Query format:

```
SELECT aggregate1, aggregate2, ...  
FROM table  
WHERE condition  
[GROUP BY expr1, expr2, ...]
```

*aggregate*:

- COUNT | SUM | AVG | STDEV (*expr*)
- MAX | MIN (*expr*)
- TOPK (*expr*, *k*, *sort\_expr*)
- HISTOGRAM (*expr*, *bins*)

*condition*:

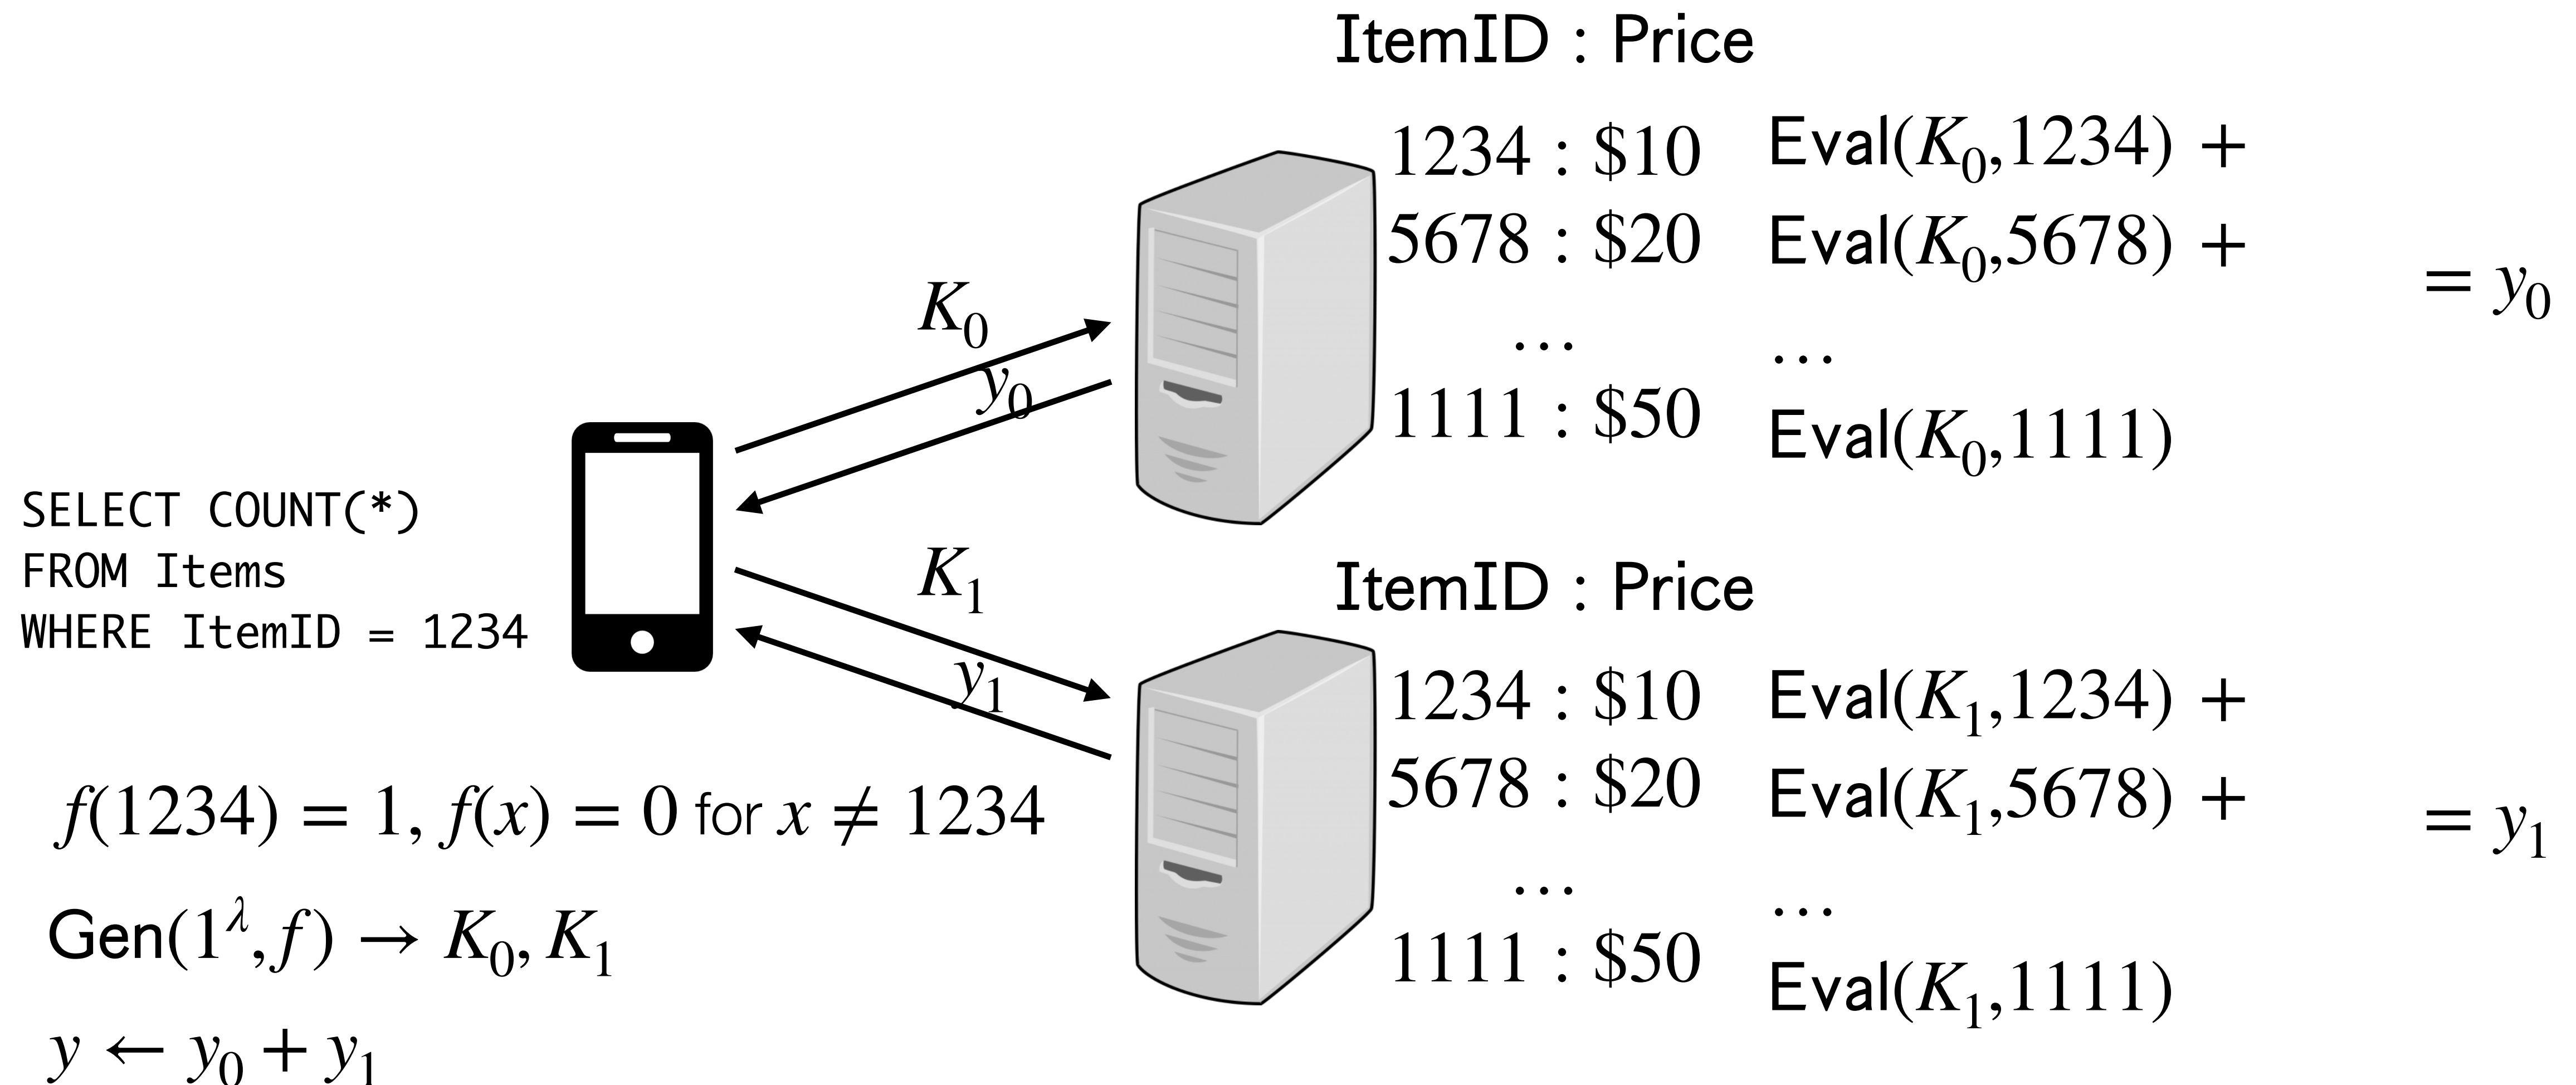
- *expr* = *secret*
- *secret*<sub>1</sub> ≤ *expr* ≤ *secret*<sub>2</sub>
- AND of '=' conditions and up to one interval
- OR of multiple disjoint conditions  
(e.g., country="UK" OR country="USA")

*expr*: any public function of the fields in a table row  
(e.g., ItemId + 1 or Price \* Tax)



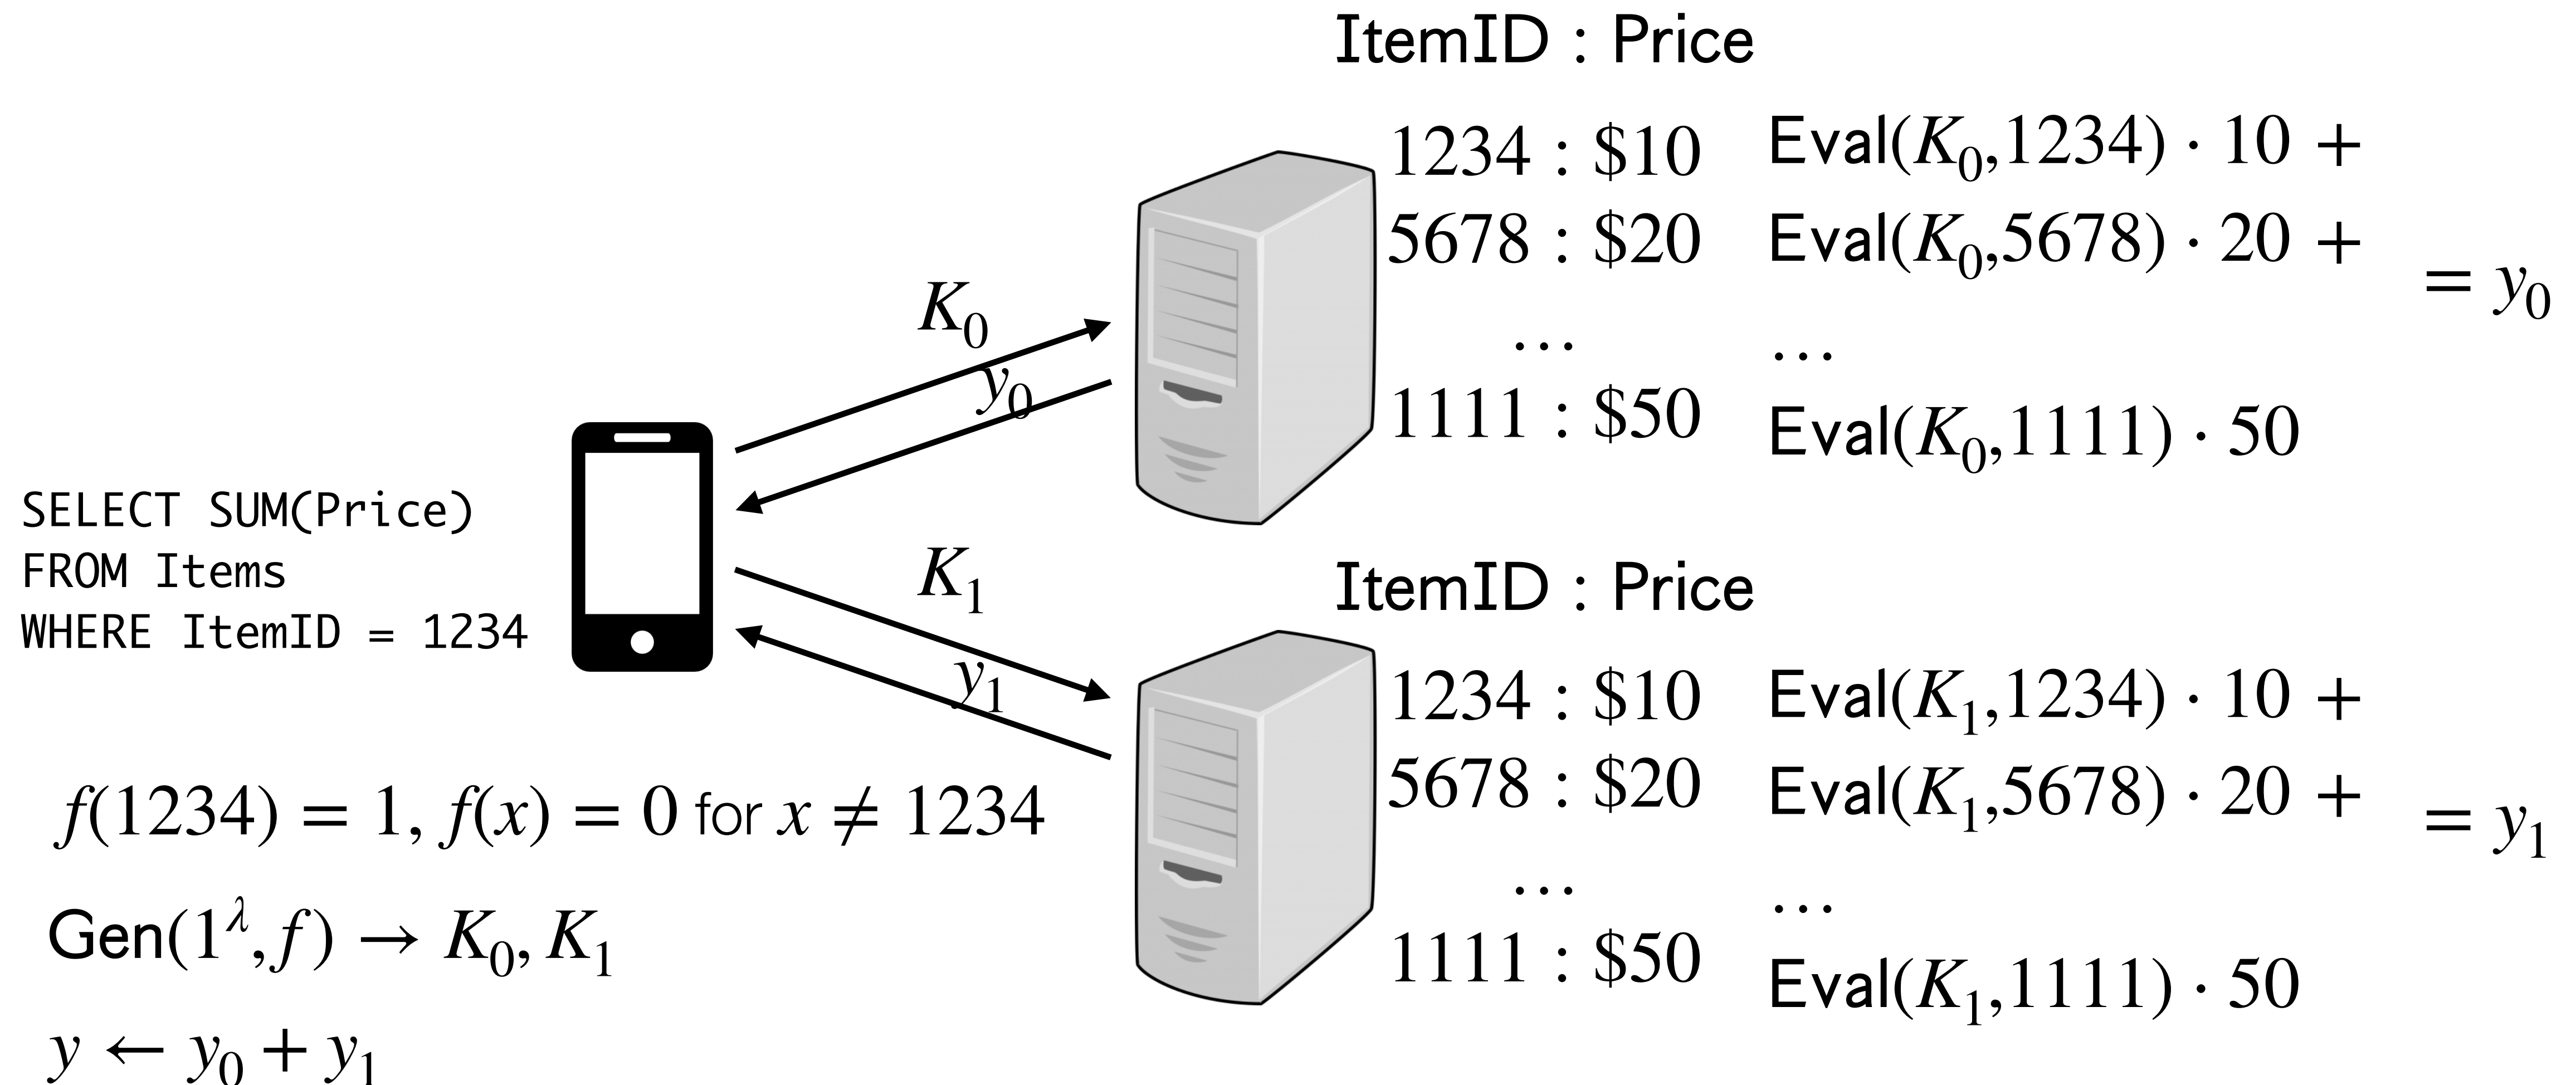
# Count queries

Can extend to interval conditions with interval FSS



# Sum queries

Can extend to interval conditions with interval FSS






# Extending to other linear aggregation functions

- $\text{AVG}(x)$ : Fetch  $\text{SUM}(x)$  and  $\text{COUNT}(x)$  and use to compute the mean
- $\text{STDEV}(x)$ : Fetch  $\text{SUM}(x)$  and  $\text{COUNT}(x)$ , along with  $\text{SUM}(x^2)$  and  $\text{COUNT}(x^2)$

# More complex conditions: AND

```
SELECT COUNT(*)
FROM Items
WHERE Price=50 AND
Color=Blue AND Size=Small
```



Price, Color, Size

\$10, Green, M  
\$20, Blue, S  
...

\$50, Yellow, L

$\text{Eval}(K_0, 10 || \text{Green} || \text{M}) +$   
 $\text{Eval}(K_0, 20 || \text{Blue} || \text{S}) +$   
 ...  
 $\text{Eval}(K_0, 50 || \text{Yellow} || \text{L})$

$= y_0$



Price, Color, Size

\$10, Green, M  
\$20, Blue, S  
...

\$50, Yellow, L

$\text{Eval}(K_0, 10 || \text{Green} || \text{M}) +$   
 $\text{Eval}(K_0, 20 || \text{Blue} || \text{S}) +$   
 ...  
 $\text{Eval}(K_0, 50 || \text{Yellow} || \text{L})$

$= y_1$

$f(50 || \text{Blue} || \text{S}) = 1, f(x) = 0$  for  
 $x \neq (50 || \text{Blue} || \text{S})$

$\text{Gen}(1^\lambda, f) \rightarrow K_0, K_1$

$y \leftarrow y_0 + y_1$

# More complex conditions: Disjoint OR

```
SELECT COUNT(*)  
FROM Trips  
WHERE Dest=SFO OR Dest=EWR
```



$f(\text{SFO}) = 1, f(x) = 0$  for  $x \neq \text{SFO}$   
 $f'(\text{EWR}) = 1, f'(x) = 0$  for  $x \neq \text{EWR}$

$\text{Gen}(1^\lambda, f) \rightarrow K_0, K_1, \text{Gen}(1^\lambda, f') \rightarrow K'_0, K'_1$   
 $y \leftarrow y_0 + y_1$

|      |  |         |
|------|--|---------|
| Dest |  |         |
| SFO  | $\text{Eval}(K_0, \text{SFO}) + \text{Eval}(K'_0, \text{SFO})$ | $= y_0$ |
| JFK  | $\text{Eval}(K_0, \text{JFK}) + \text{Eval}(K'_0, \text{JFK})$ |         |
| ...  | ...  |         |
| LAX  | $\text{Eval}(K_0, \text{LAX}) + \text{Eval}(K'_0, \text{LAX})$ |         |
| Dest |  |         |
| SFO  | $\text{Eval}(K_1, \text{SFO}) + \text{Eval}(K'_1, \text{SFO})$ | $= y_1$ |
| JFK  | $\text{Eval}(K_1, \text{JFK}) + \text{Eval}(K'_1, \text{JFK})$ |         |
| ...  | ...  |         |
| LAX  | $\text{Eval}(K_1, \text{LAX}) + \text{Eval}(K'_1, \text{LAX})$ |         |

# Max/min for interval conditions

```
SELECT MAX(Price) FROM Items WHERE 9 <= Rating <= 10
```

Setup: build an array  $A$  ordered by ratings

Round 1: Find indices  $i, j$  in  $A$  where all **rating**  $\in [9, 10]$  using keys for interval function for **rating**  $\in [0, 8]$ , **rating**  $\in [0, 9]$

Round 2: Select maxes from power-of-2 sized intervals that cover exactly records  $i$  through  $j$  in  $A$

|                  |           |   |   |   |   |   |   |   |
|------------------|-----------|---|---|---|---|---|---|---|
|                  | $A[3..6]$ |   |   |   |   |   |   |   |
| $A$              | 3         | 5 | 1 | 2 | 4 | 1 | 0 | 1 |
| Size-2 intervals | 5         |   | 2 |   | 4 |   | 1 |   |
| Size-4 intervals | 5         |   |   |   | 4 |   |   |   |
| Size-8 intervals | 5         |   |   |   |   |   |   |   |

See paper for max for single conditions

# Max/min for disjoint OR

```
SELECT MIN(Price) FROM Flights WHERE Dest=SFO OR Dest=EWR
```

Setup: build an array  $A$  ordered by ratings

Binary search over  $A$  in decreasing power-of-2 sizes to find the smallest index where  $\text{Dest}=\text{SFO}$  or  $\text{Dest}=\text{EWR}$

- `SELECT COUNT(*) FROM A WHERE (Dest=SFO OR Dest=EWR) AND index  $\in$  [start, end]`
- To search over interval, take advantage of fact that intervals are power-of-2 aligned (match on first  $x$  bits)

Use a DPF to retrieve the price for the entry with the smallest index satisfying  $\text{Dest}=\text{SFO}$  OR  $\text{Dest}=\text{EWR}$

[For small number of conditions, can also run one query for each clause and combine at the client]




# Outline

1. Function secret sharing
2. Splinter
- 3. Searching on encrypted data**

# End-to-end encrypted filesystems

Provide strong security guarantees if attacker compromises server.




# Users expect the ability to search

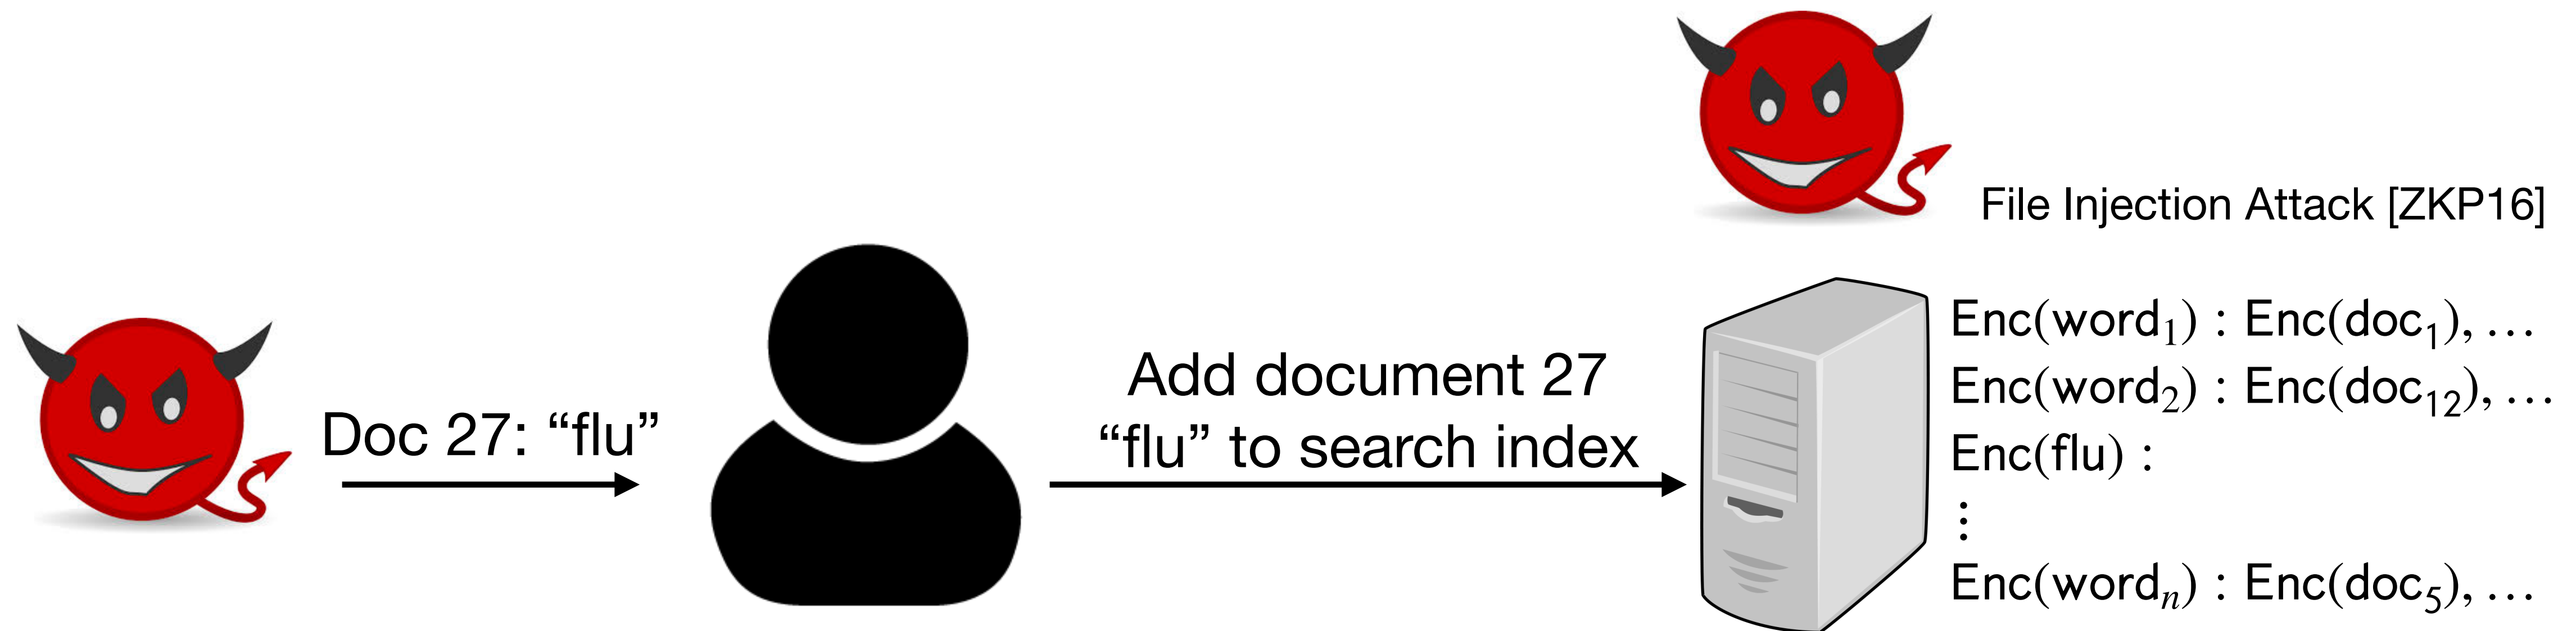
Doc 1  
Doc 7  
Doc 21  
Doc 53

# Search for end-to-end encrypted filesystems

**Challenge:** server cannot decrypt data to search.

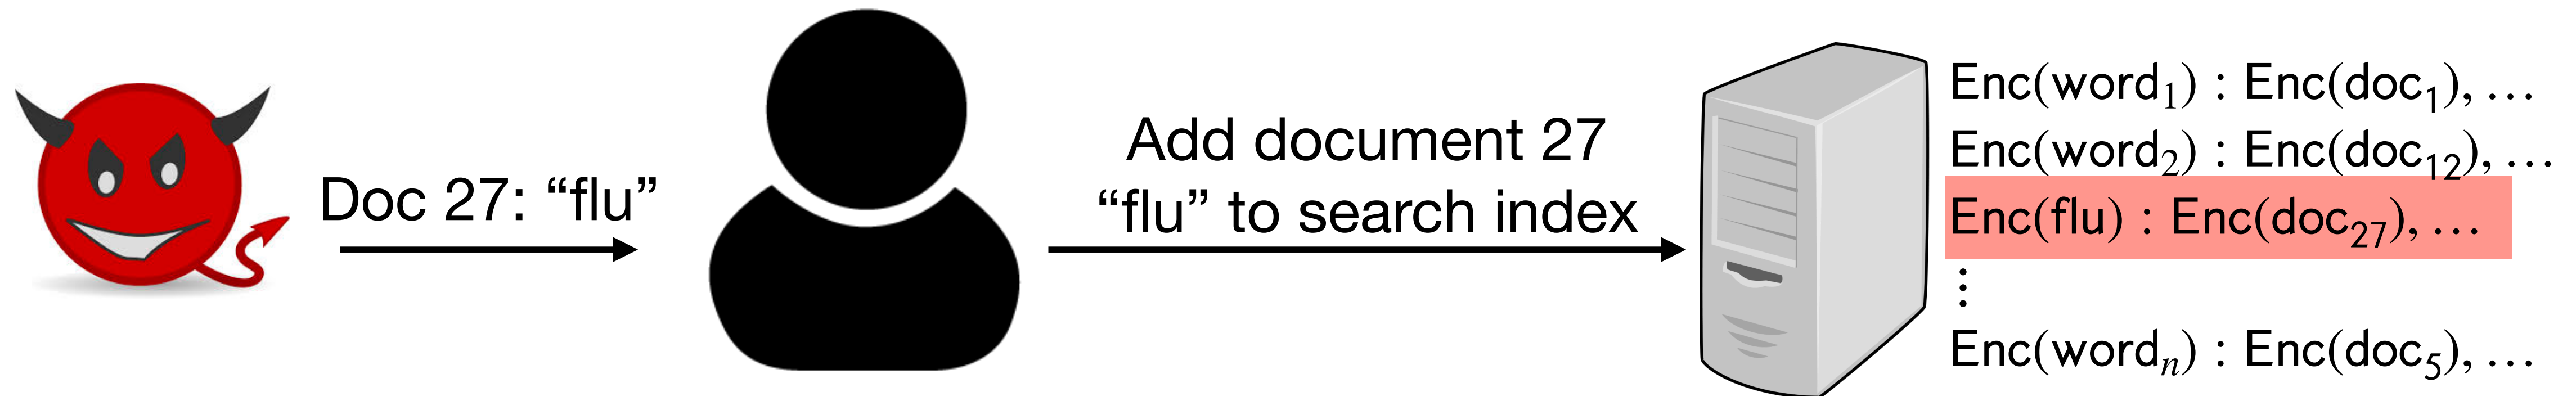


# How the server accesses an encrypted index can leak information




# How the server accesses an encrypted index can leak information


Repeat for all words in English dictionary.



# How the server accesses an encrypted index can leak information



# How the server accesses an encrypted index can leak information




... and many more attacks




# Oblivious RAM solutions are slow for encrypted search

Client can read/write server data without revealing data location [GO96, PathORAM]




# Security properties



One malicious server:

- Attacker can't learn anything about queries or document contents from memory accesses

# Security properties



One malicious server:


- Attacker can't learn anything about queries or document contents from memory accesses

Two malicious servers:

- Attacker can't immediately learn document contents (but can infer info from memory accesses)

# DORY: Efficient search without search-access-pattern leakage

[Dauterman, Feng, Luo, Popa, Stoica]



# High-level approach

## Searchable encryption schemes

[Song, Wagner, Perrig], ...

Queries require accessing  
part of search index

Leak info about query and data

## DORY

Queries require accessing  
entire search index  
 $\implies$  ok if concretely fast

Leak nothing about query or data

# Building the search index

Bloom filter containing keywords in doc 1

|         |           |           |           |         |           |
|---------|-----------|-----------|-----------|---------|-----------|
| doc 1   | $v_{1,1}$ | $v_{1,2}$ | $v_{1,3}$ | $\dots$ | $v_{1,m}$ |
| doc 2   | $v_{2,1}$ | $v_{2,2}$ | $v_{2,3}$ | $\dots$ | $v_{2,m}$ |
| doc 3   | $v_{3,1}$ | $v_{3,2}$ | $v_{3,3}$ | $\dots$ | $v_{3,m}$ |
|         | $\dots$   | $\dots$   | $\dots$   | $\dots$ | $\dots$   |
| doc $n$ | $v_{n,1}$ | $v_{n,2}$ | $v_{n,3}$ | $\dots$ | $v_{n,m}$ |

# Building the search index

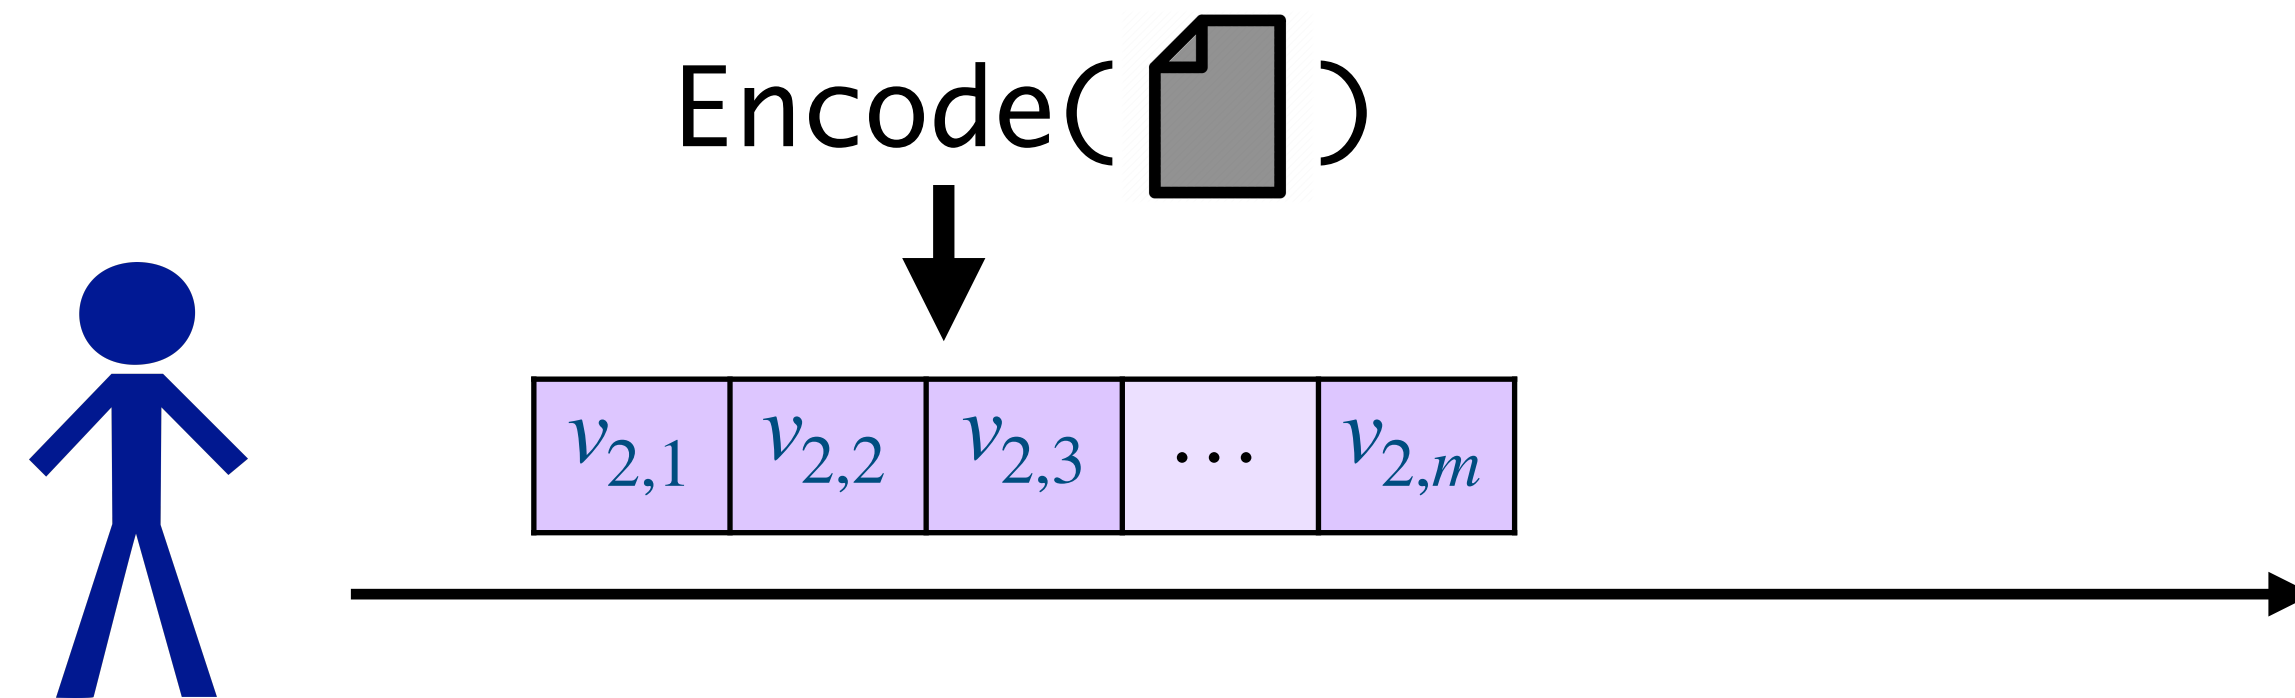
Bloom filter containing keywords in doc 1

|         |           |           |           |     |           |
|---------|-----------|-----------|-----------|-----|-----------|
| doc 1   | $v_{1,1}$ | $v_{1,2}$ | $v_{1,3}$ | ... | $v_{1,m}$ |
| doc 2   | $v_{2,1}$ | $v_{2,2}$ | $v_{2,3}$ | ... | $v_{2,m}$ |
| doc 3   | $v_{3,1}$ | $v_{3,2}$ | $v_{3,3}$ | ... | $v_{3,m}$ |
|         | ...       | ...       | ...       | ... | ...       |
| doc $n$ | $v_{n,1}$ | $v_{n,2}$ | $v_{n,3}$ | ... | $v_{n,m}$ |

Can check for a keyword  
with few point queries

# High-level approach

Update: write a row

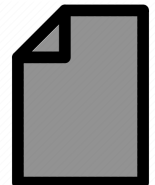



|         |           |           |           |         |           |
|---------|-----------|-----------|-----------|---------|-----------|
| doc 1   | $v_{1,1}$ | $v_{1,2}$ | $v_{1,3}$ | $\dots$ | $v_{1,m}$ |
| doc 2   | $v_{2,1}$ | $v_{2,2}$ | $v_{2,3}$ | $\dots$ | $v_{2,m}$ |
| doc 3   | $v_{3,1}$ | $v_{3,2}$ | $v_{3,3}$ | $\dots$ | $v_{3,m}$ |
|         | $\dots$   | $\dots$   | $\dots$   | $\dots$ | $\dots$   |
| doc $n$ | $v_{n,1}$ | $v_{n,2}$ | $v_{n,3}$ | $\dots$ | $v_{n,m}$ |



# High-level approach

Update: write a row

Encode(  )



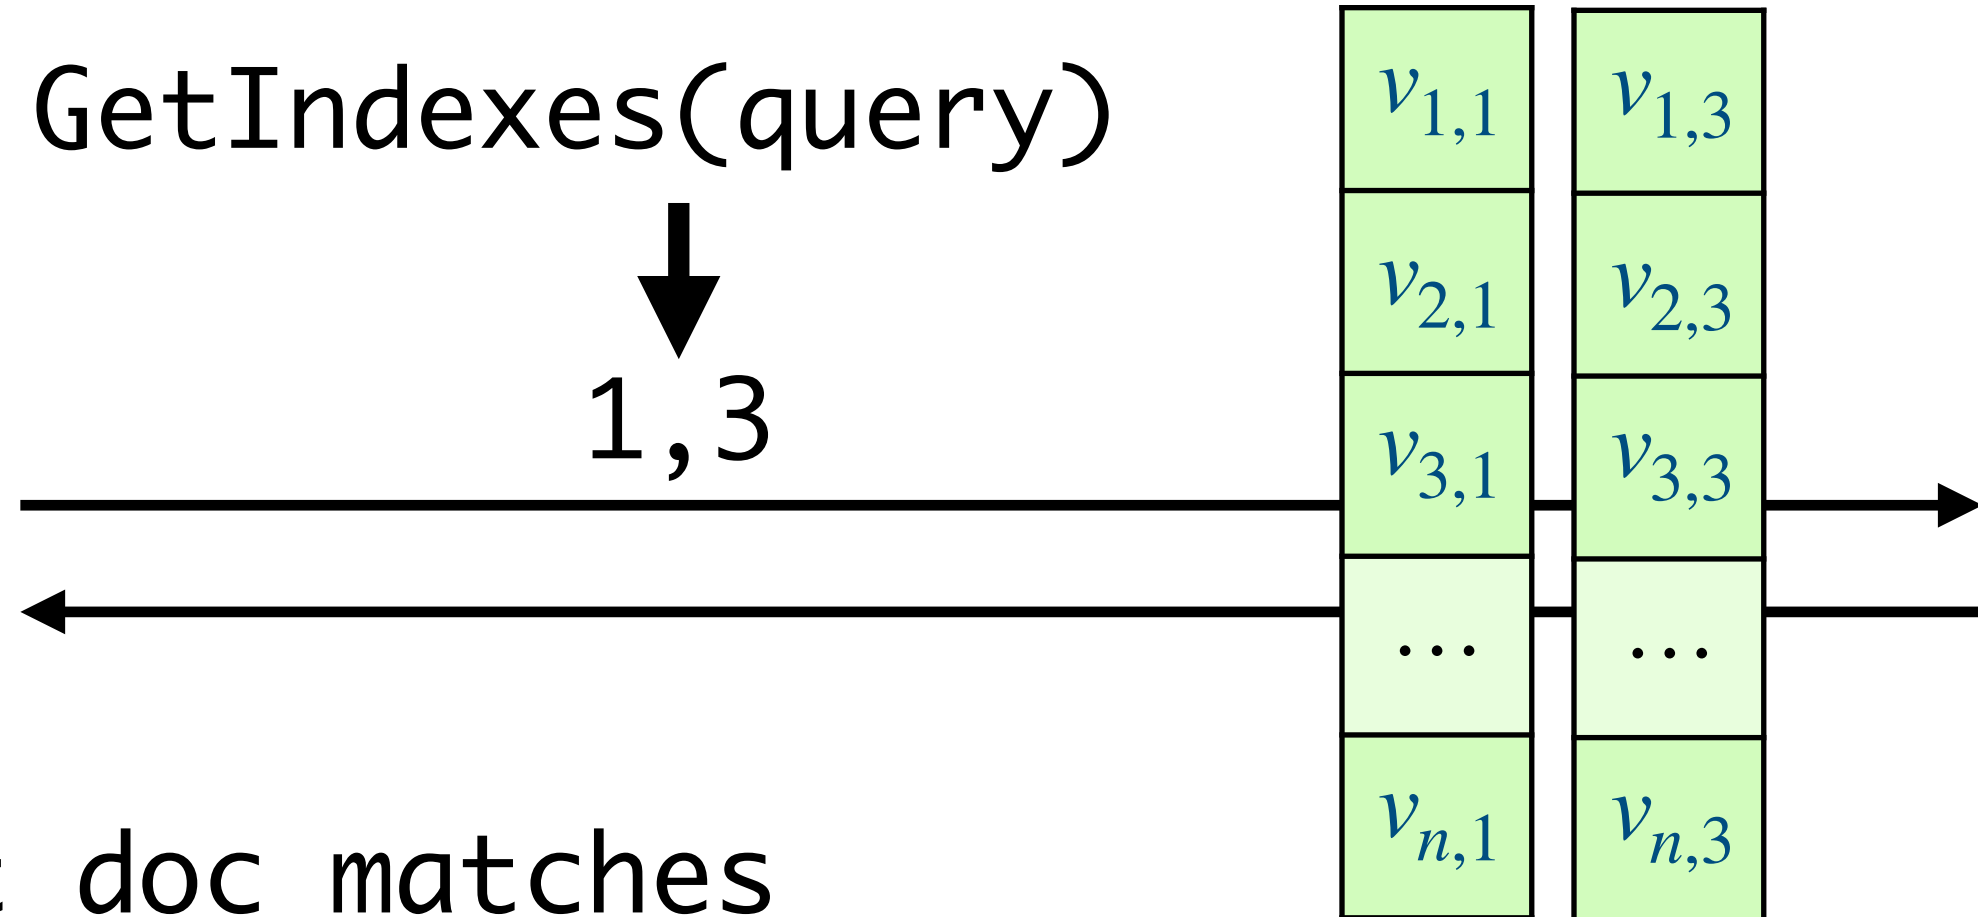

|           |           |           |     |           |
|-----------|-----------|-----------|-----|-----------|
| $v_{2,1}$ | $v_{2,2}$ | $v_{2,3}$ | ... | $v_{2,m}$ |
|-----------|-----------|-----------|-----|-----------|



Search: *privately* read columns

GetIndexes(query)

↓  
1, 3





Output doc matches

|         |           |           |           |     |           |
|---------|-----------|-----------|-----------|-----|-----------|
| doc 1   | $v_{1,1}$ | $v_{1,2}$ | $v_{1,3}$ | ... | $v_{1,m}$ |
| doc 2   | $v_{2,1}$ | $v_{2,2}$ | $v_{2,3}$ | ... | $v_{2,m}$ |
| doc 3   | $v_{3,1}$ | $v_{3,2}$ | $v_{3,3}$ | ... | $v_{3,m}$ |
|         | ...       | ...       | ...       | ... | ...       |
| doc $n$ | $v_{n,1}$ | $v_{n,2}$ | $v_{n,3}$ | ... | $v_{n,m}$ |

# High-level approach

Update: write a row

Encode(  )



Privately  
retrieve column


$v_{2,m}$

Check search  
result integrity

Search: *privately* read column

GetIndexes(query)

1, 3



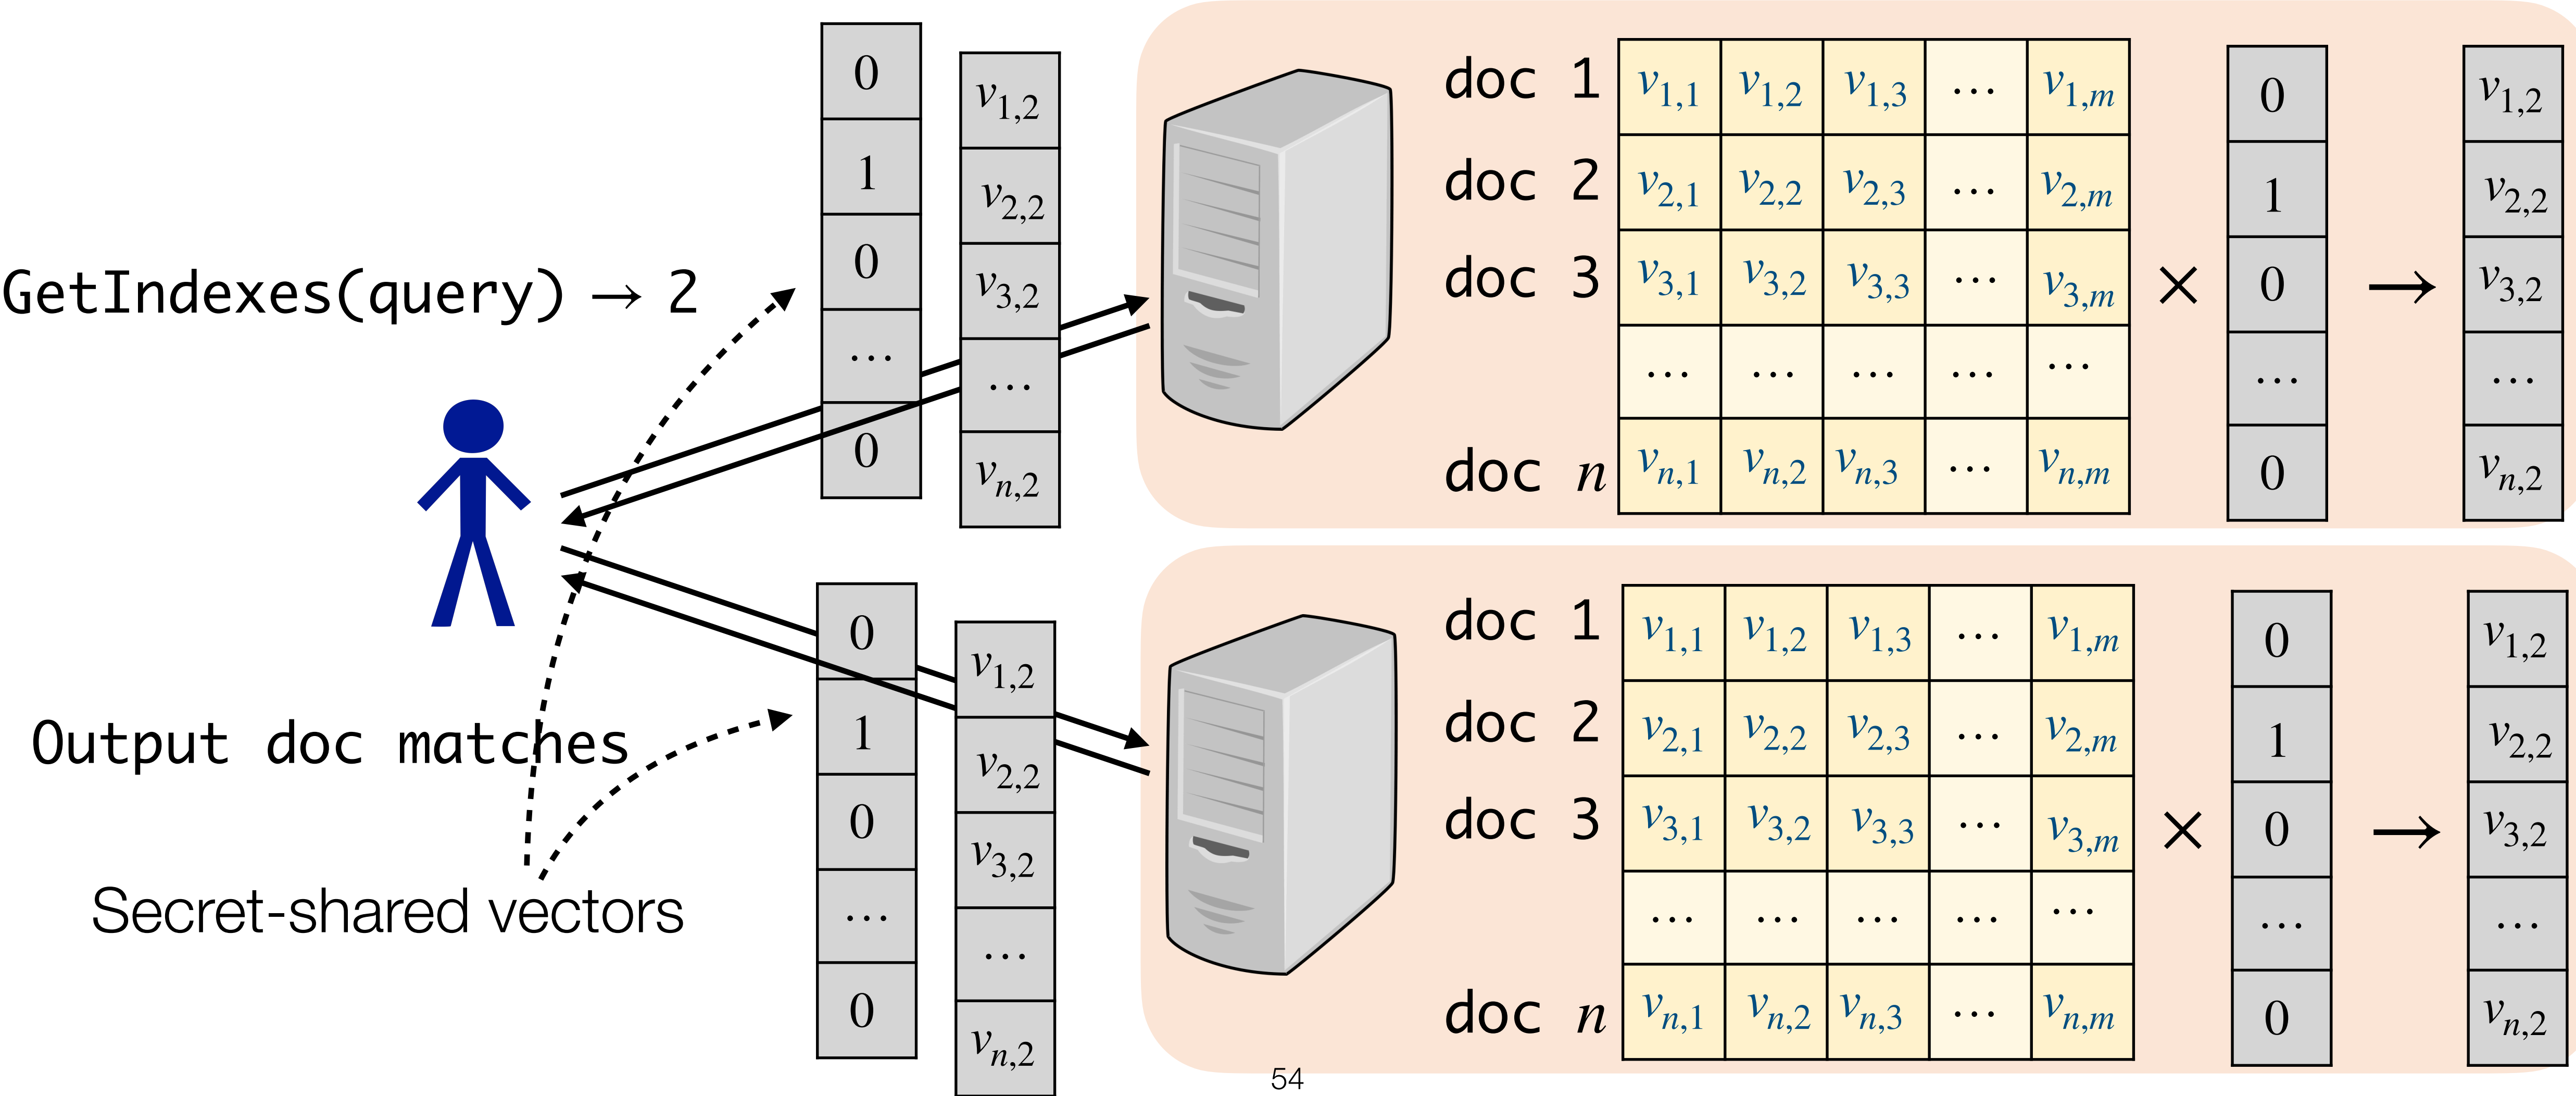
Output doc matches

|           |           |
|-----------|-----------|
| $v_{1,1}$ | $v_{1,3}$ |
| $v_{2,1}$ | $v_{2,3}$ |
| $v_{3,1}$ | $v_{3,3}$ |
| ...       | ...       |
| $v_{n,1}$ | $v_{n,3}$ |

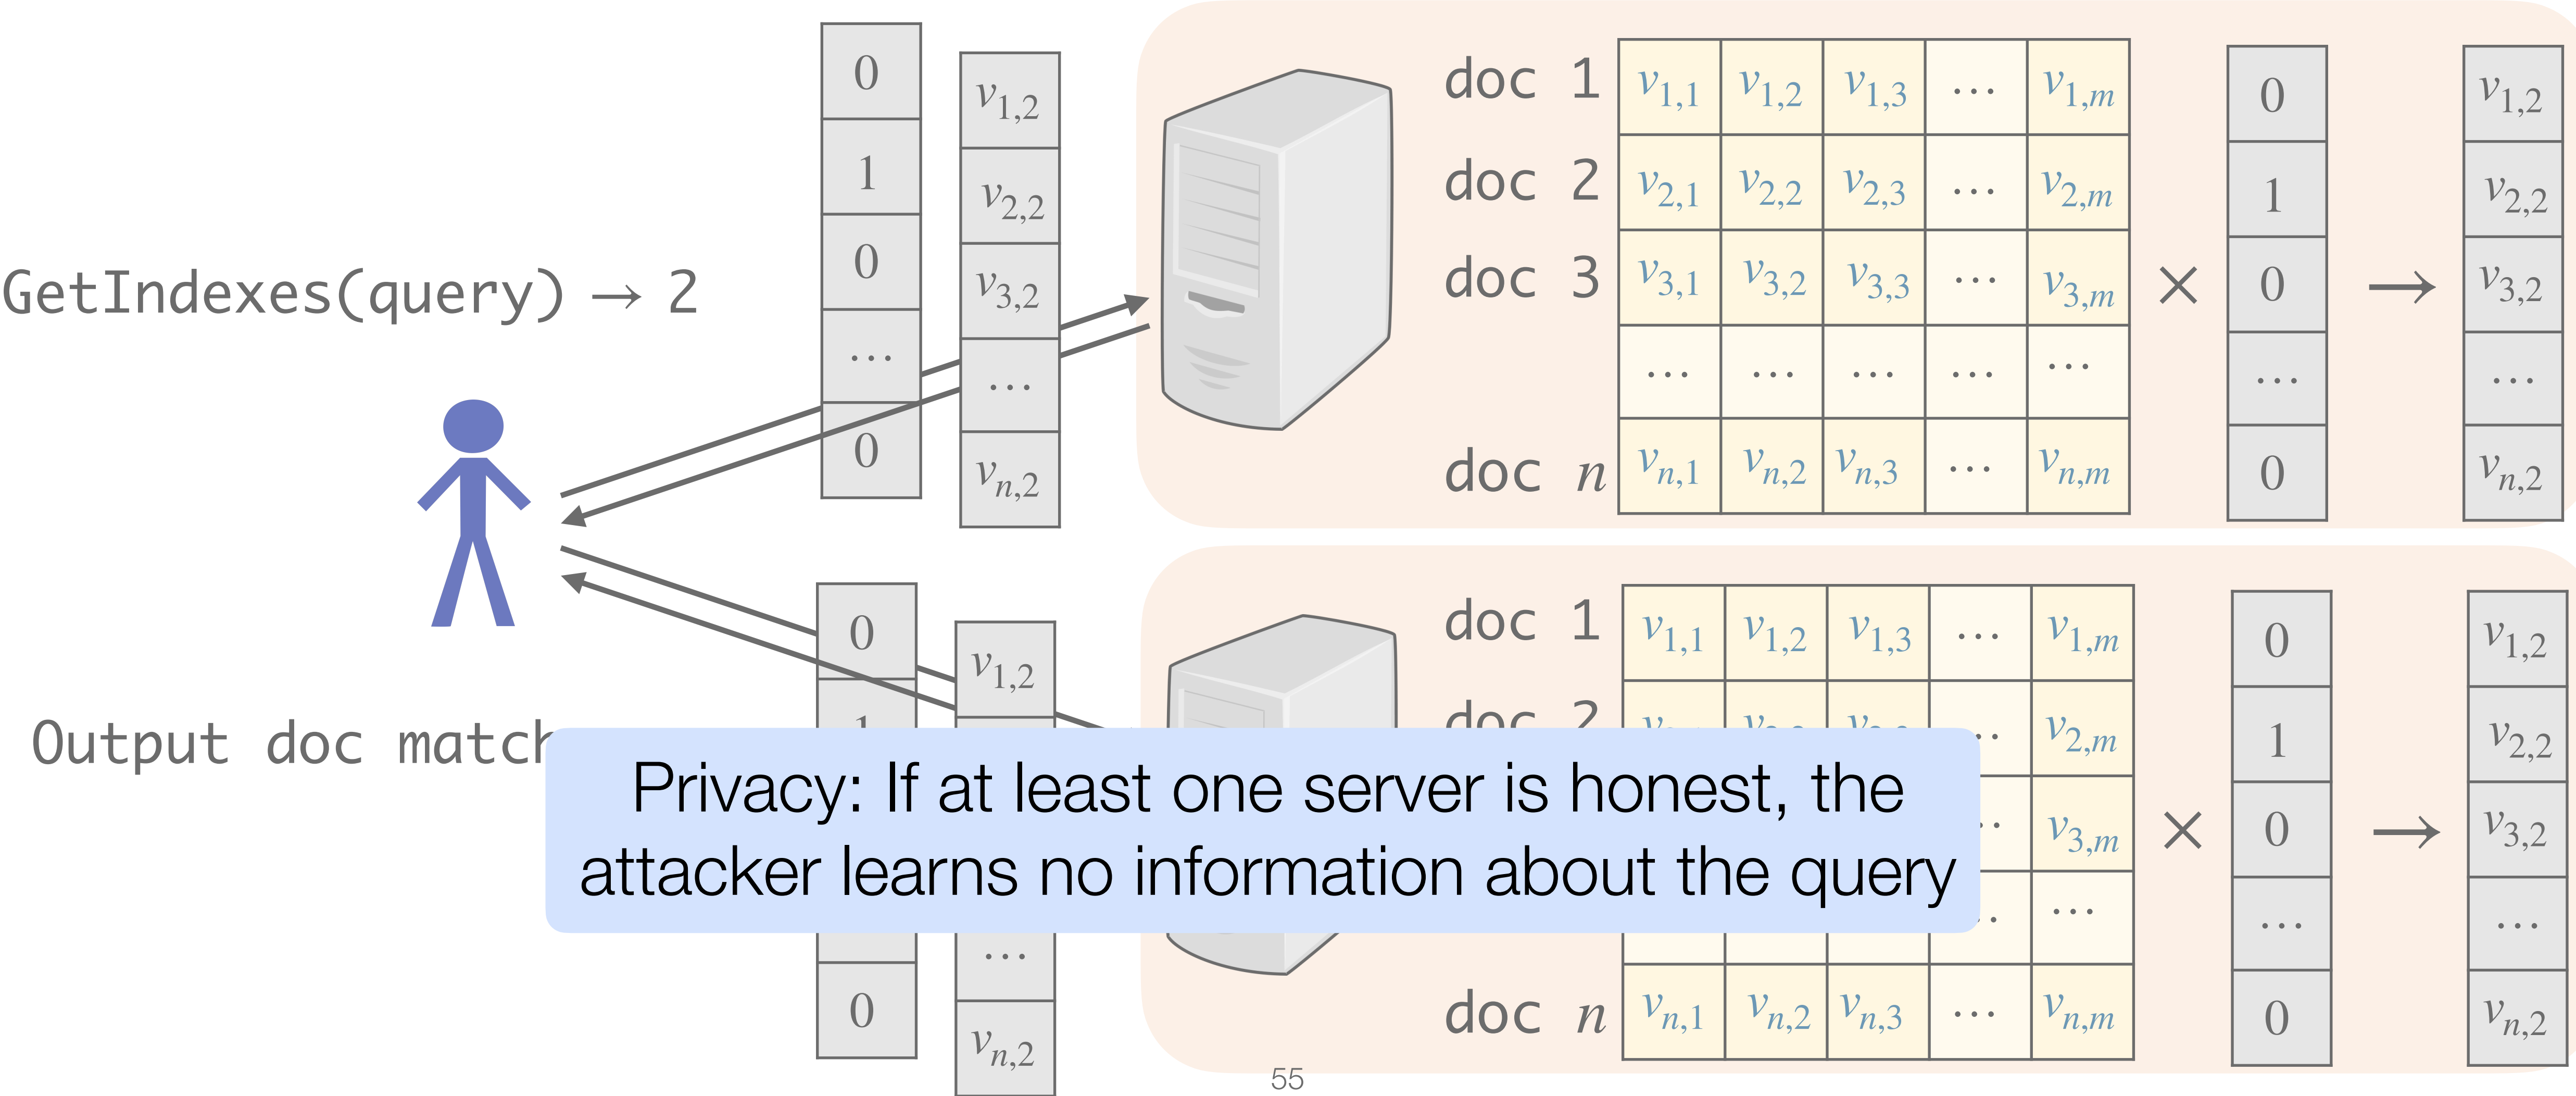
Encrypting the  
search index

|         |           |           |           |     |           |
|---------|-----------|-----------|-----------|-----|-----------|
| doc 1   | $v_{1,1}$ | $v_{1,2}$ | $v_{1,3}$ | ... | $v_{1,m}$ |
| doc 2   | $v_{2,1}$ | $v_{2,2}$ | $v_{2,3}$ | ... | $v_{2,m}$ |
| doc 3   | $v_{3,1}$ | $v_{3,2}$ | $v_{3,3}$ | ... | $v_{3,m}$ |
|         | ...       | ...       | ...       | ... | ...       |
| doc $n$ | $v_{n,1}$ | $v_{n,2}$ | $v_{n,3}$ | ... | $v_{n,m}$ |

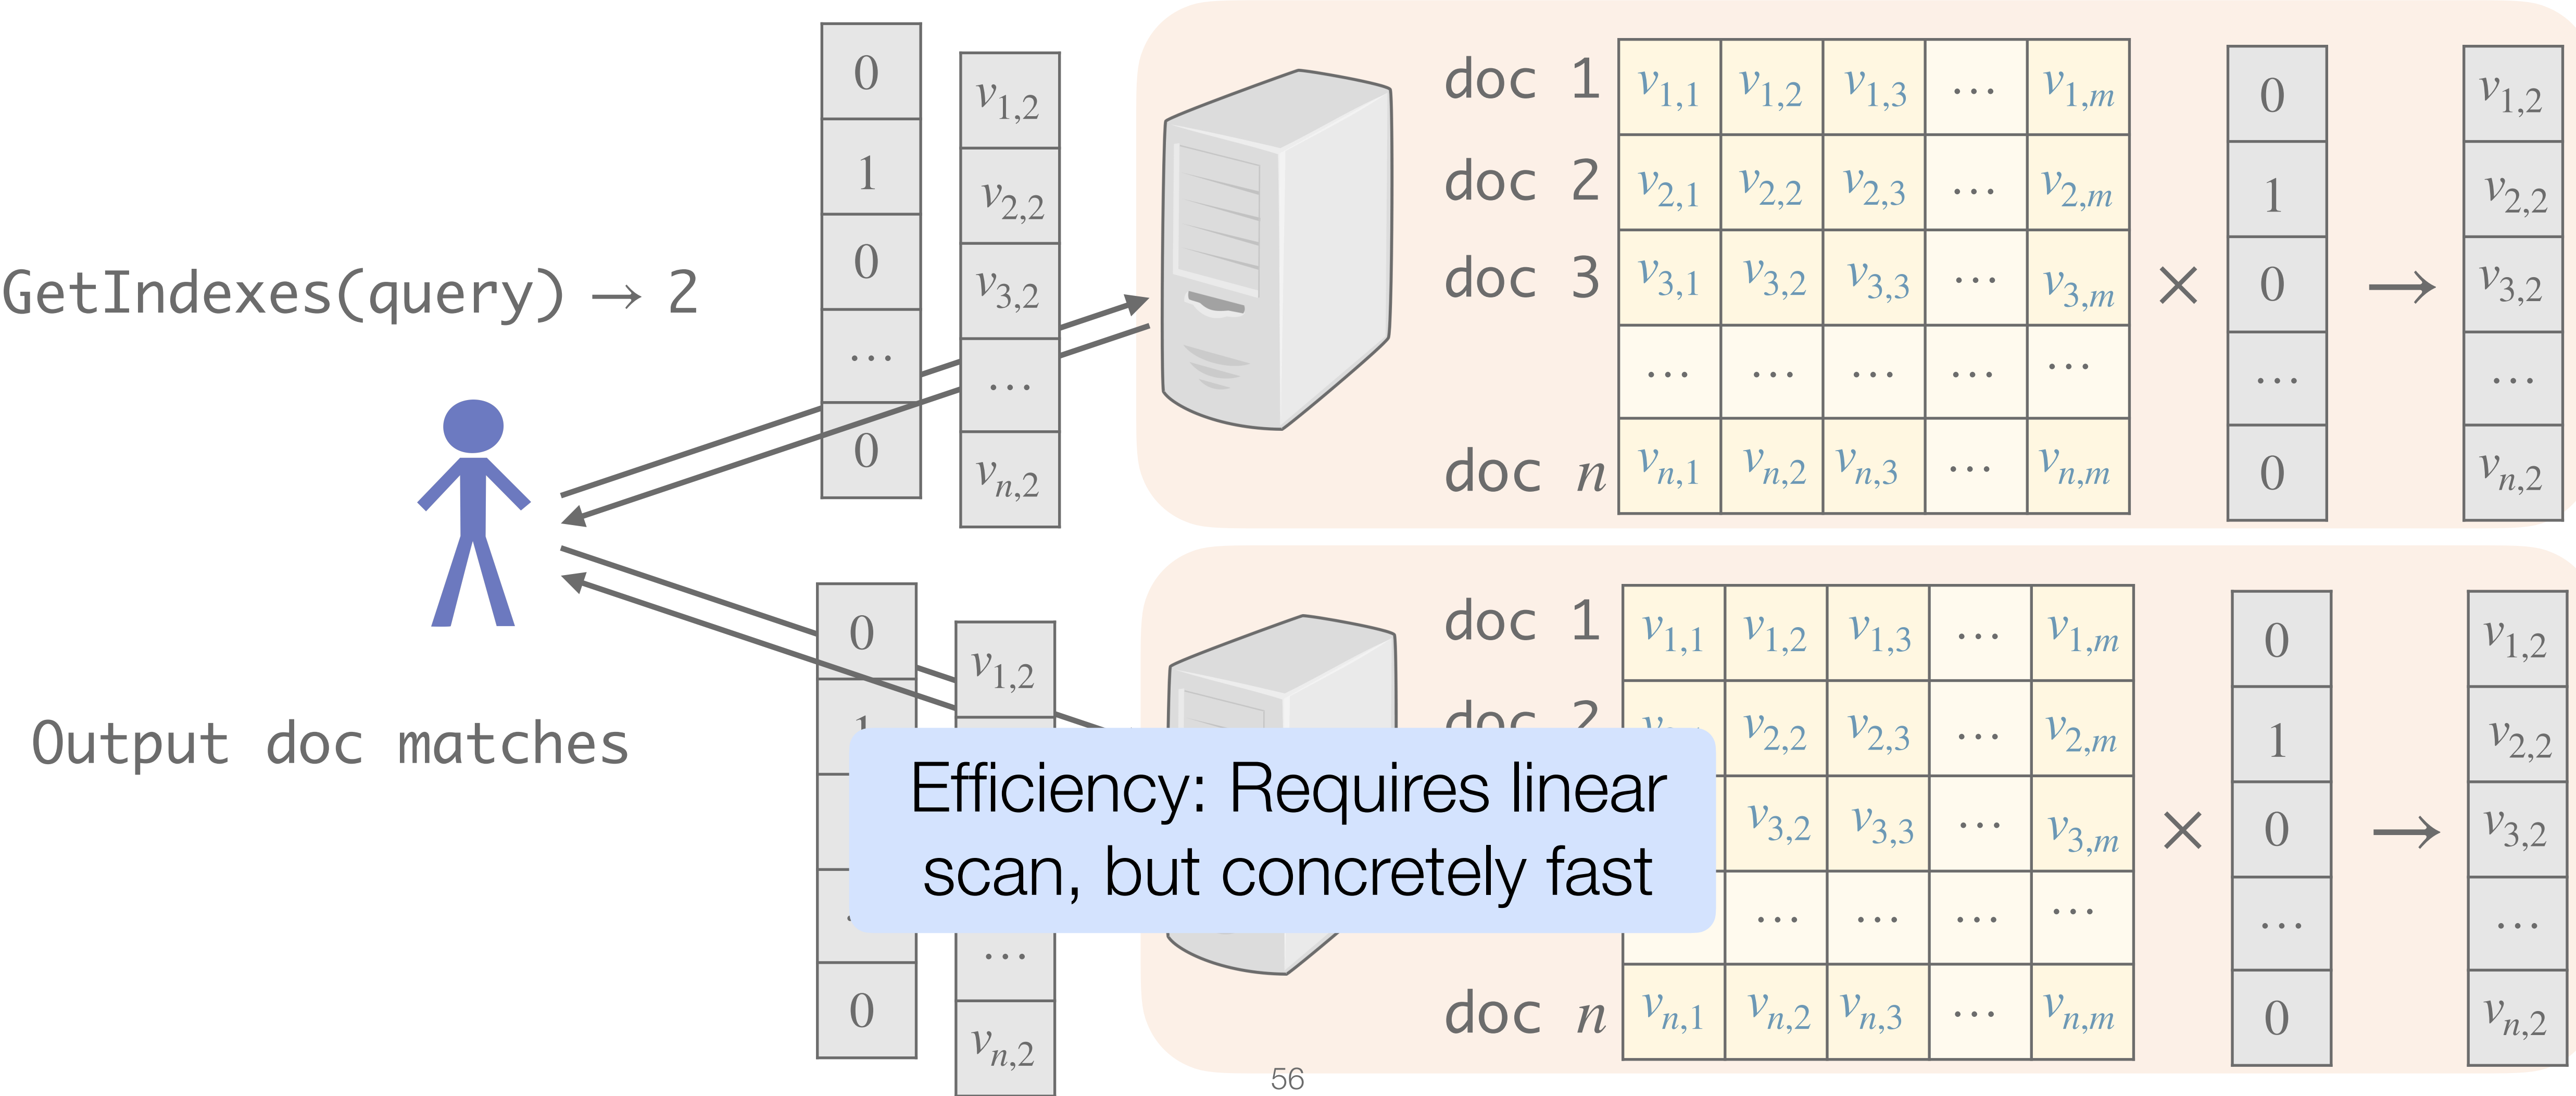
# Privately reading a column to search



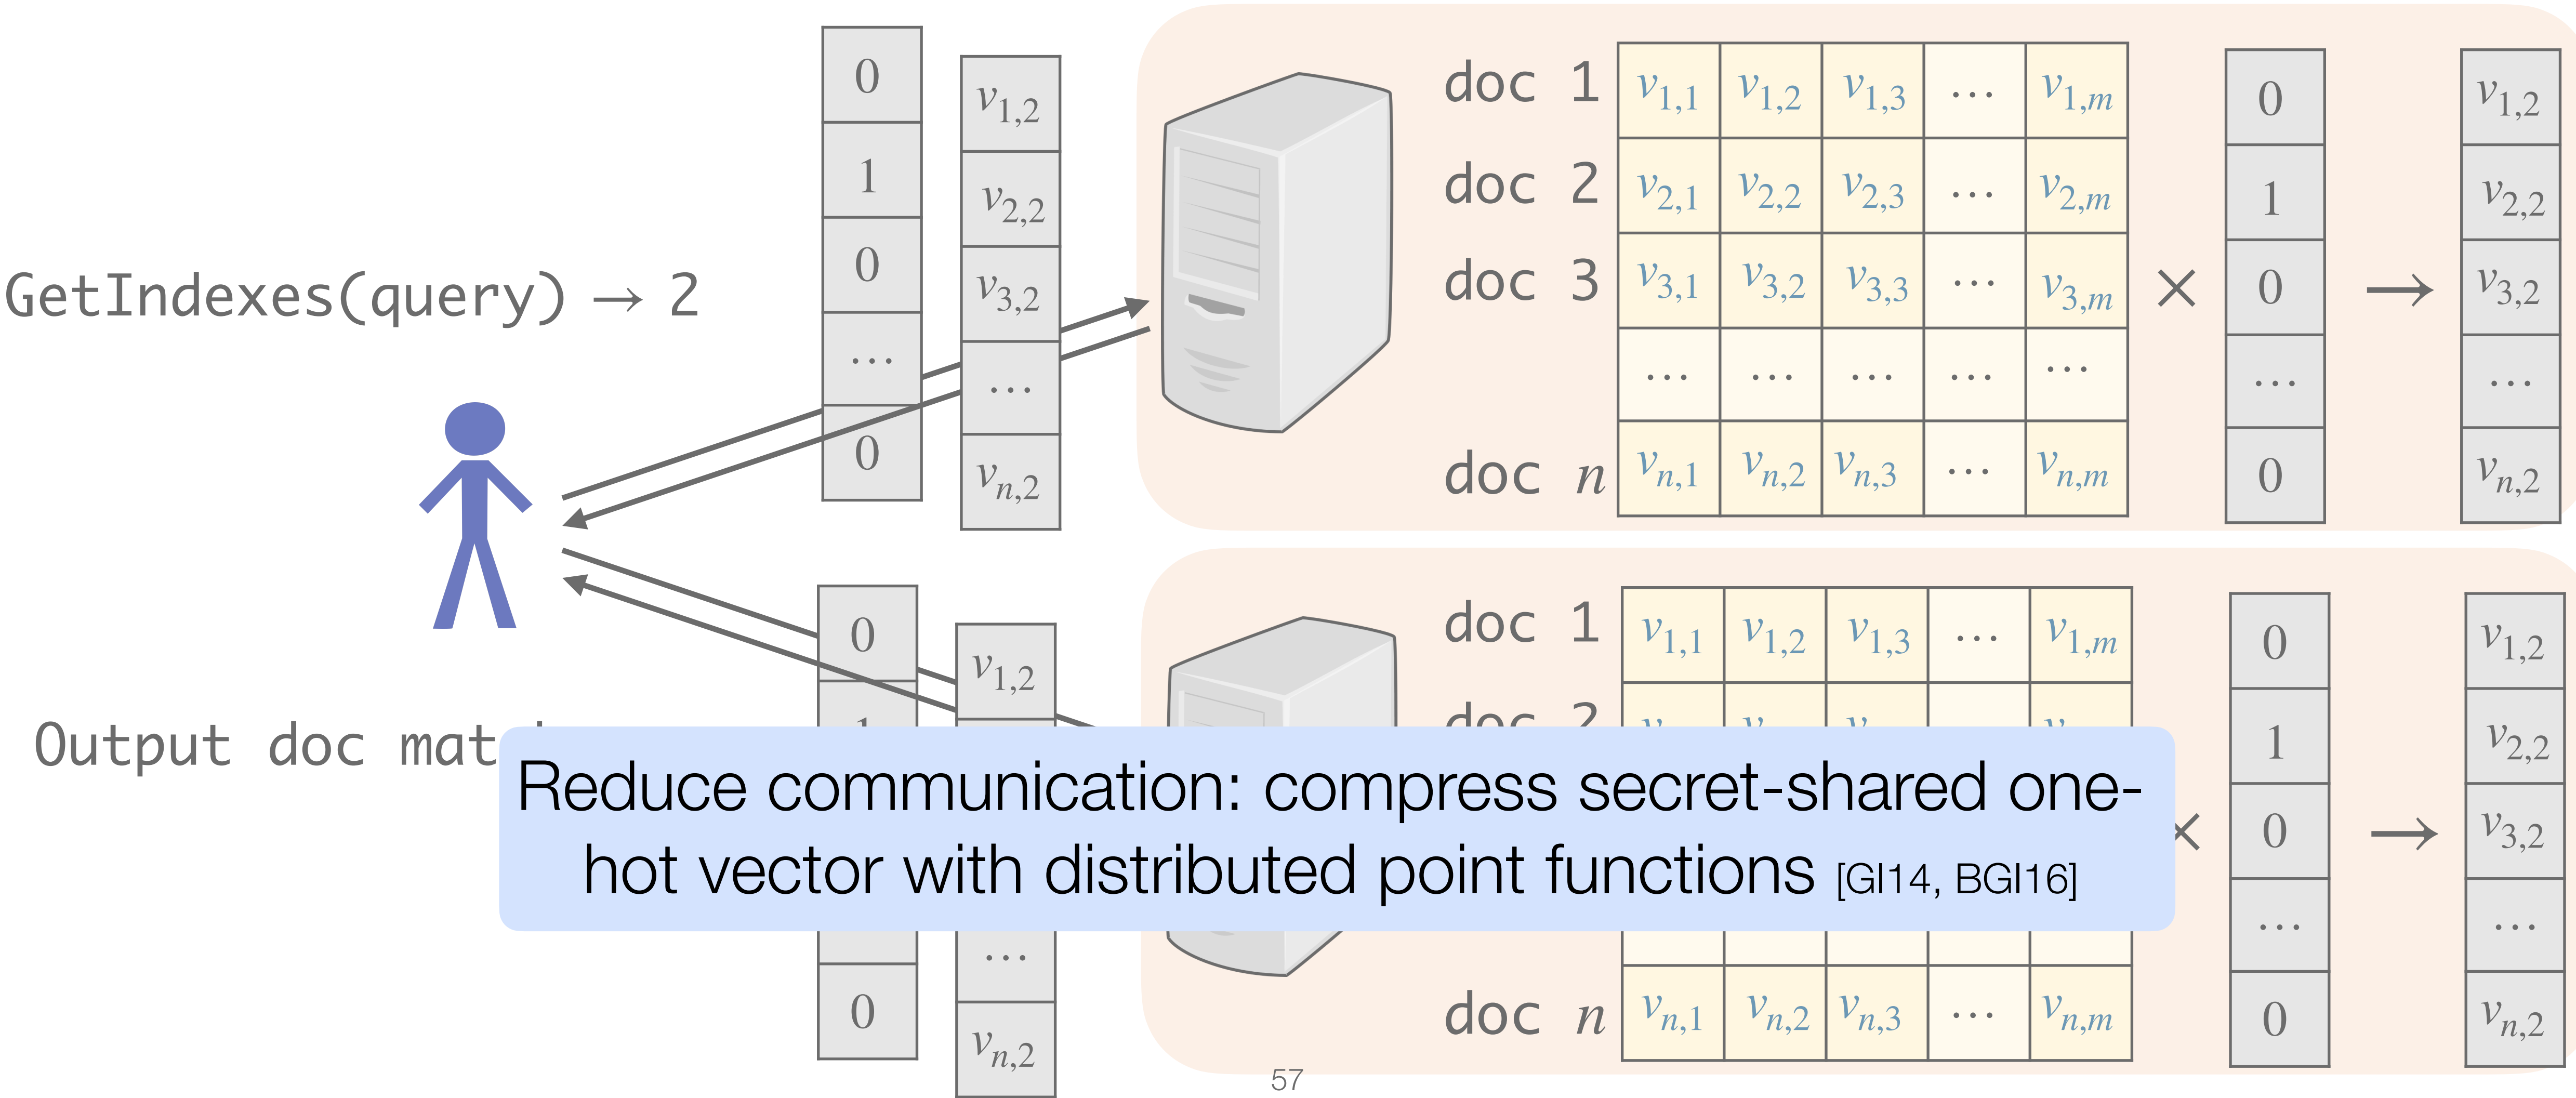
# Privately reading a column to search



# Privately reading a column to search



# Privately reading a column to search



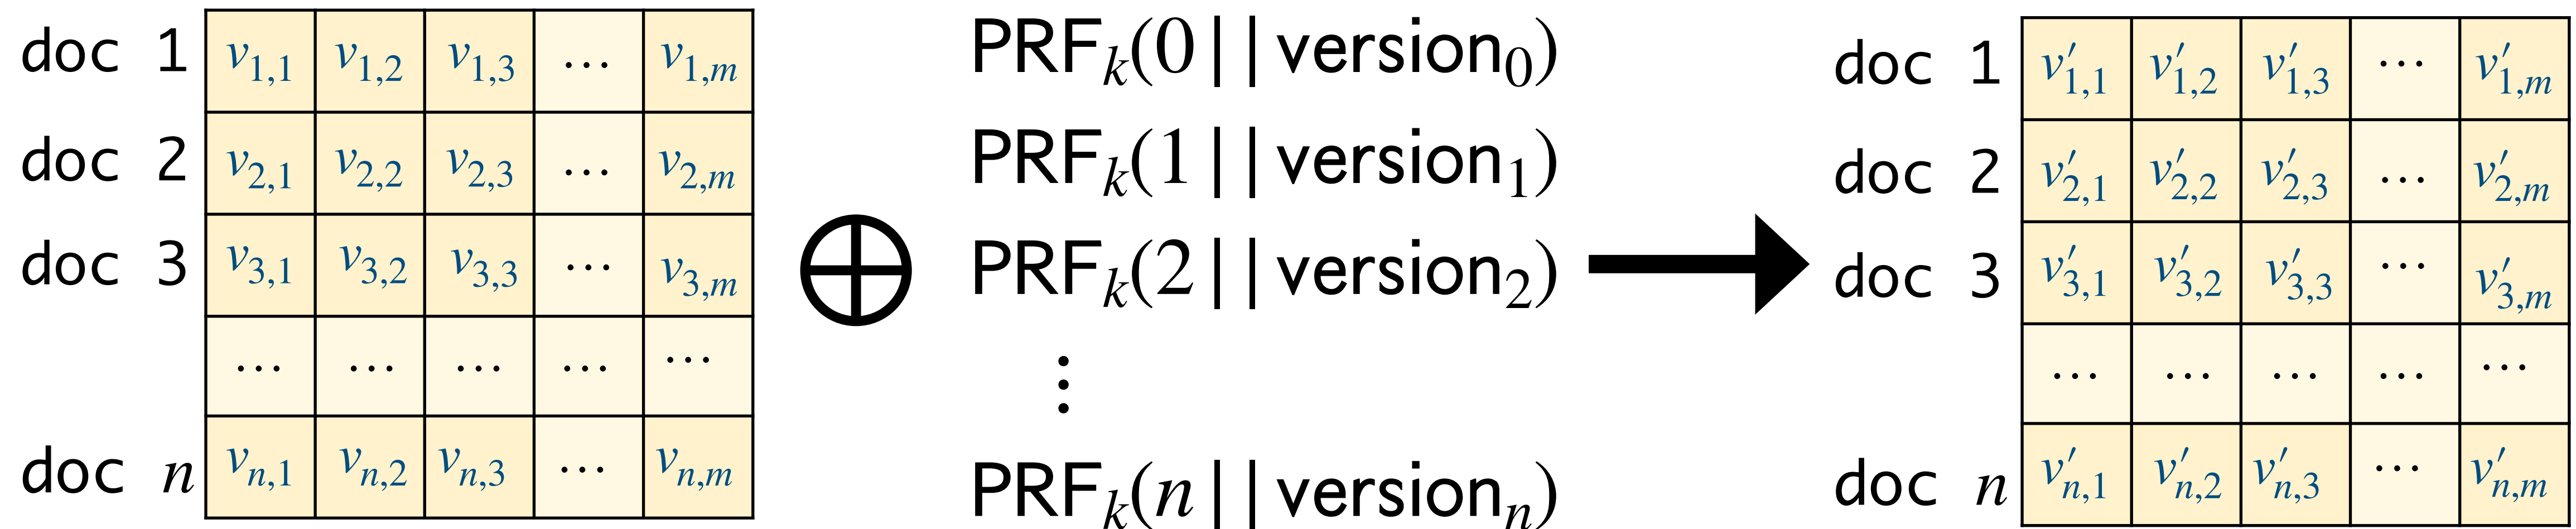


# Encrypting the search index

Challenge: Client uploads *rows* of bits, but retrieves *columns* of bits

- Don't want to increase communication by a factor of  $\lambda \approx 128$ .

Idea: Generate a unique mask using document version number



# Search result integrity

MAC tags allow client to verify result corresponds to valid update

**Aggregate MACs** compress all MAC tags in a column into a single tag [KL08]

- XOR individual MAC tags into single aggregate tag

| Aggregate MACs |  | $t_1$     | $t_2$     | $t_3$     | $\dots$ | $t_m$     |
|----------------|--|-----------|-----------|-----------|---------|-----------|
| doc 1          |  | $v_{1,1}$ | $v_{1,2}$ | $v_{1,3}$ | $\dots$ | $v_{1,m}$ |
| doc 2          |  | $v_{2,1}$ | $v_{2,2}$ | $v_{2,3}$ | $\dots$ | $v_{2,m}$ |
| doc 3          |  | $v_{3,1}$ | $v_{3,2}$ | $v_{3,3}$ | $\dots$ | $v_{3,m}$ |
|                |  | $\dots$   | $\dots$   | $\dots$   | $\dots$ | $\dots$   |
| doc $n$        |  | $v_{n,1}$ | $v_{n,2}$ | $v_{n,3}$ | $\dots$ | $v_{n,m}$ |



# Logistics

Signups for meetings for feedback on project proposals

- Required for all groups

Course feedback form on Ed

# References

- Beimel, Amos, Yuval Ishai, and Tal Malkin. "Reducing the servers computation in private information retrieval: PIR with preprocessing." In *Annual International Cryptology Conference*, pp. 55-73. Berlin, Heidelberg: Springer Berlin Heidelberg, 2000.
- Boneh, Dan, Elette Boyle, Henry Corrigan-Gibbs, Niv Gilboa, and Yuval Ishai. "Lightweight techniques for private heavy hitters." In *2021 IEEE Symposium on Security and Privacy (SP)*, pp. 762-776. IEEE, 2021.
- Boyle, Elette, Niv Gilboa, and Yuval Ishai. "Function secret sharing." In *Annual international conference on the theory and applications of cryptographic techniques*, pp. 337-367. Berlin, Heidelberg: Springer Berlin Heidelberg, 2015.
- Boyle, Elette, Niv Gilboa, and Yuval Ishai. "Function secret sharing: Improvements and extensions." In *Proceedings of the 2016 ACM SIGSAC conference on computer and communications security*, pp. 1292-1303. 2016.
- Boyle, Elette, Niv Gilboa, and Yuval Ishai. "Secure computation with preprocessing via function secret sharing." In *Theory of Cryptography Conference*, pp. 341-371. Cham: Springer International Publishing, 2019.
- Chor, Benny, Eyal Kushilevitz, Oded Goldreich, and Madhu Sudan. "Private information retrieval." *Journal of the ACM (JACM)* 45, no. 6 (1998): 965-981.
- Corrigan-Gibbs, Henry, Dan Boneh, and David Mazières. "Riposte: An anonymous messaging system handling millions of users." In *2015 IEEE Symposium on Security and Privacy*, pp. 321-338. IEEE, 2015.
- Goldreich, Oded, and Rafail Ostrovsky. "Software protection and simulation on oblivious RAMs." *Journal of the ACM (JACM)* 43, no. 3 (1996): 431-473.
- Katz, Jonathan, and Andrew Y. Lindell. "Aggregate message authentication codes." In *Cryptographers' Track at the RSA Conference*, pp. 155-169. Berlin, Heidelberg: Springer Berlin Heidelberg, 2008.
- Stefanov, Emil, Marten van Dijk, Elaine Shi, T-H. Hubert Chan, Christopher Fletcher, Ling Ren, Xiangyao Yu, and Srinivas Devadas. "Path ORAM: an extremely simple oblivious RAM protocol." *Journal of the ACM (JACM)* 65, no. 4 (2018): 1-26.
- Song, Dawn Xiaoding, David Wagner, and Adrian Perrig. "Practical techniques for searches on encrypted data." In *Proceeding 2000 IEEE symposium on security and privacy. S&P 2000*, pp. 44-55. IEEE, 2000.
- Zhang, Yupeng, Jonathan Katz, and Charalampos Papamanthou. "All your queries are belong to us: the power of {File-Injection} attacks on searchable encryption." In *25th USENIX Security Symposium (USENIX Security 16)*, pp. 707-720. 2016.